

Data Structures in Java

Lecture 17: Introduction to Graphs.

11/13/2019

Daniel Bauer

Graphs

- A **Graph** is a pair of two sets $G=(V,E)$:
 - V : the set of **vertices** (or **nodes**)
 - E : the set of **edges**.
 - each edge is a pair (v,w) where $v,w \in V$

Graphs

- A **Graph** is a pair of two sets $G=(V,E)$:

v_1

v_2

v_3

v_4

- V : the set of **vertices** (or **nodes**)

- E : the set of **edges**.

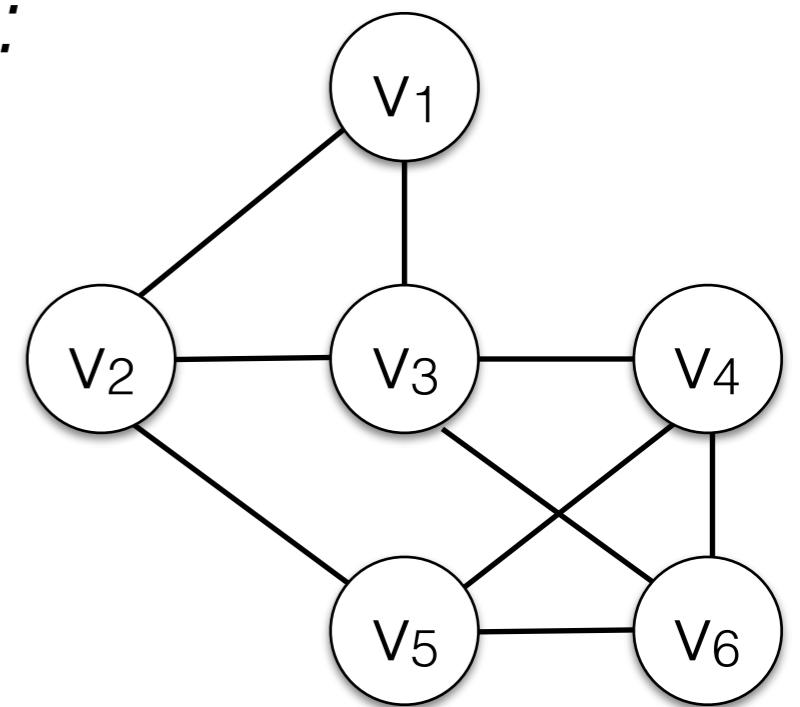
- each edge is a pair (v,w) where
 $v,w \in V$

v_5

v_6

Graphs

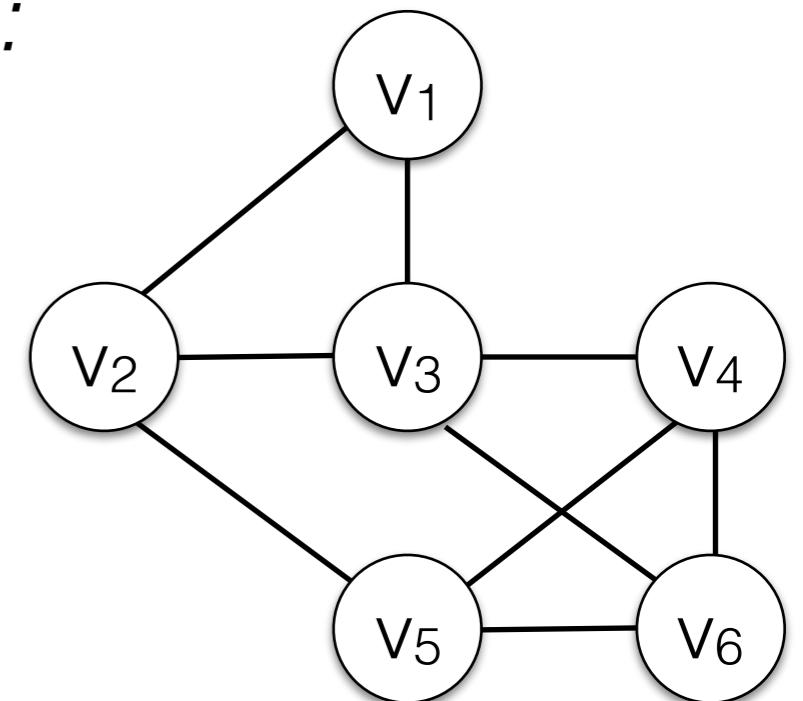
- A **Graph** is a pair of two sets $G=(V,E)$:
 - V : the set of **vertices** (or **nodes**)
 - E : the set of **edges**.
 - each edge is a pair (v,w) where $v,w \in V$



Graphs

- A **Graph** is a pair of two sets $G=(V,E)$:

- V : the set of **vertices** (or **nodes**)
- E : the set of **edges**.
 - each edge is a pair (v,w) where $v,w \in V$



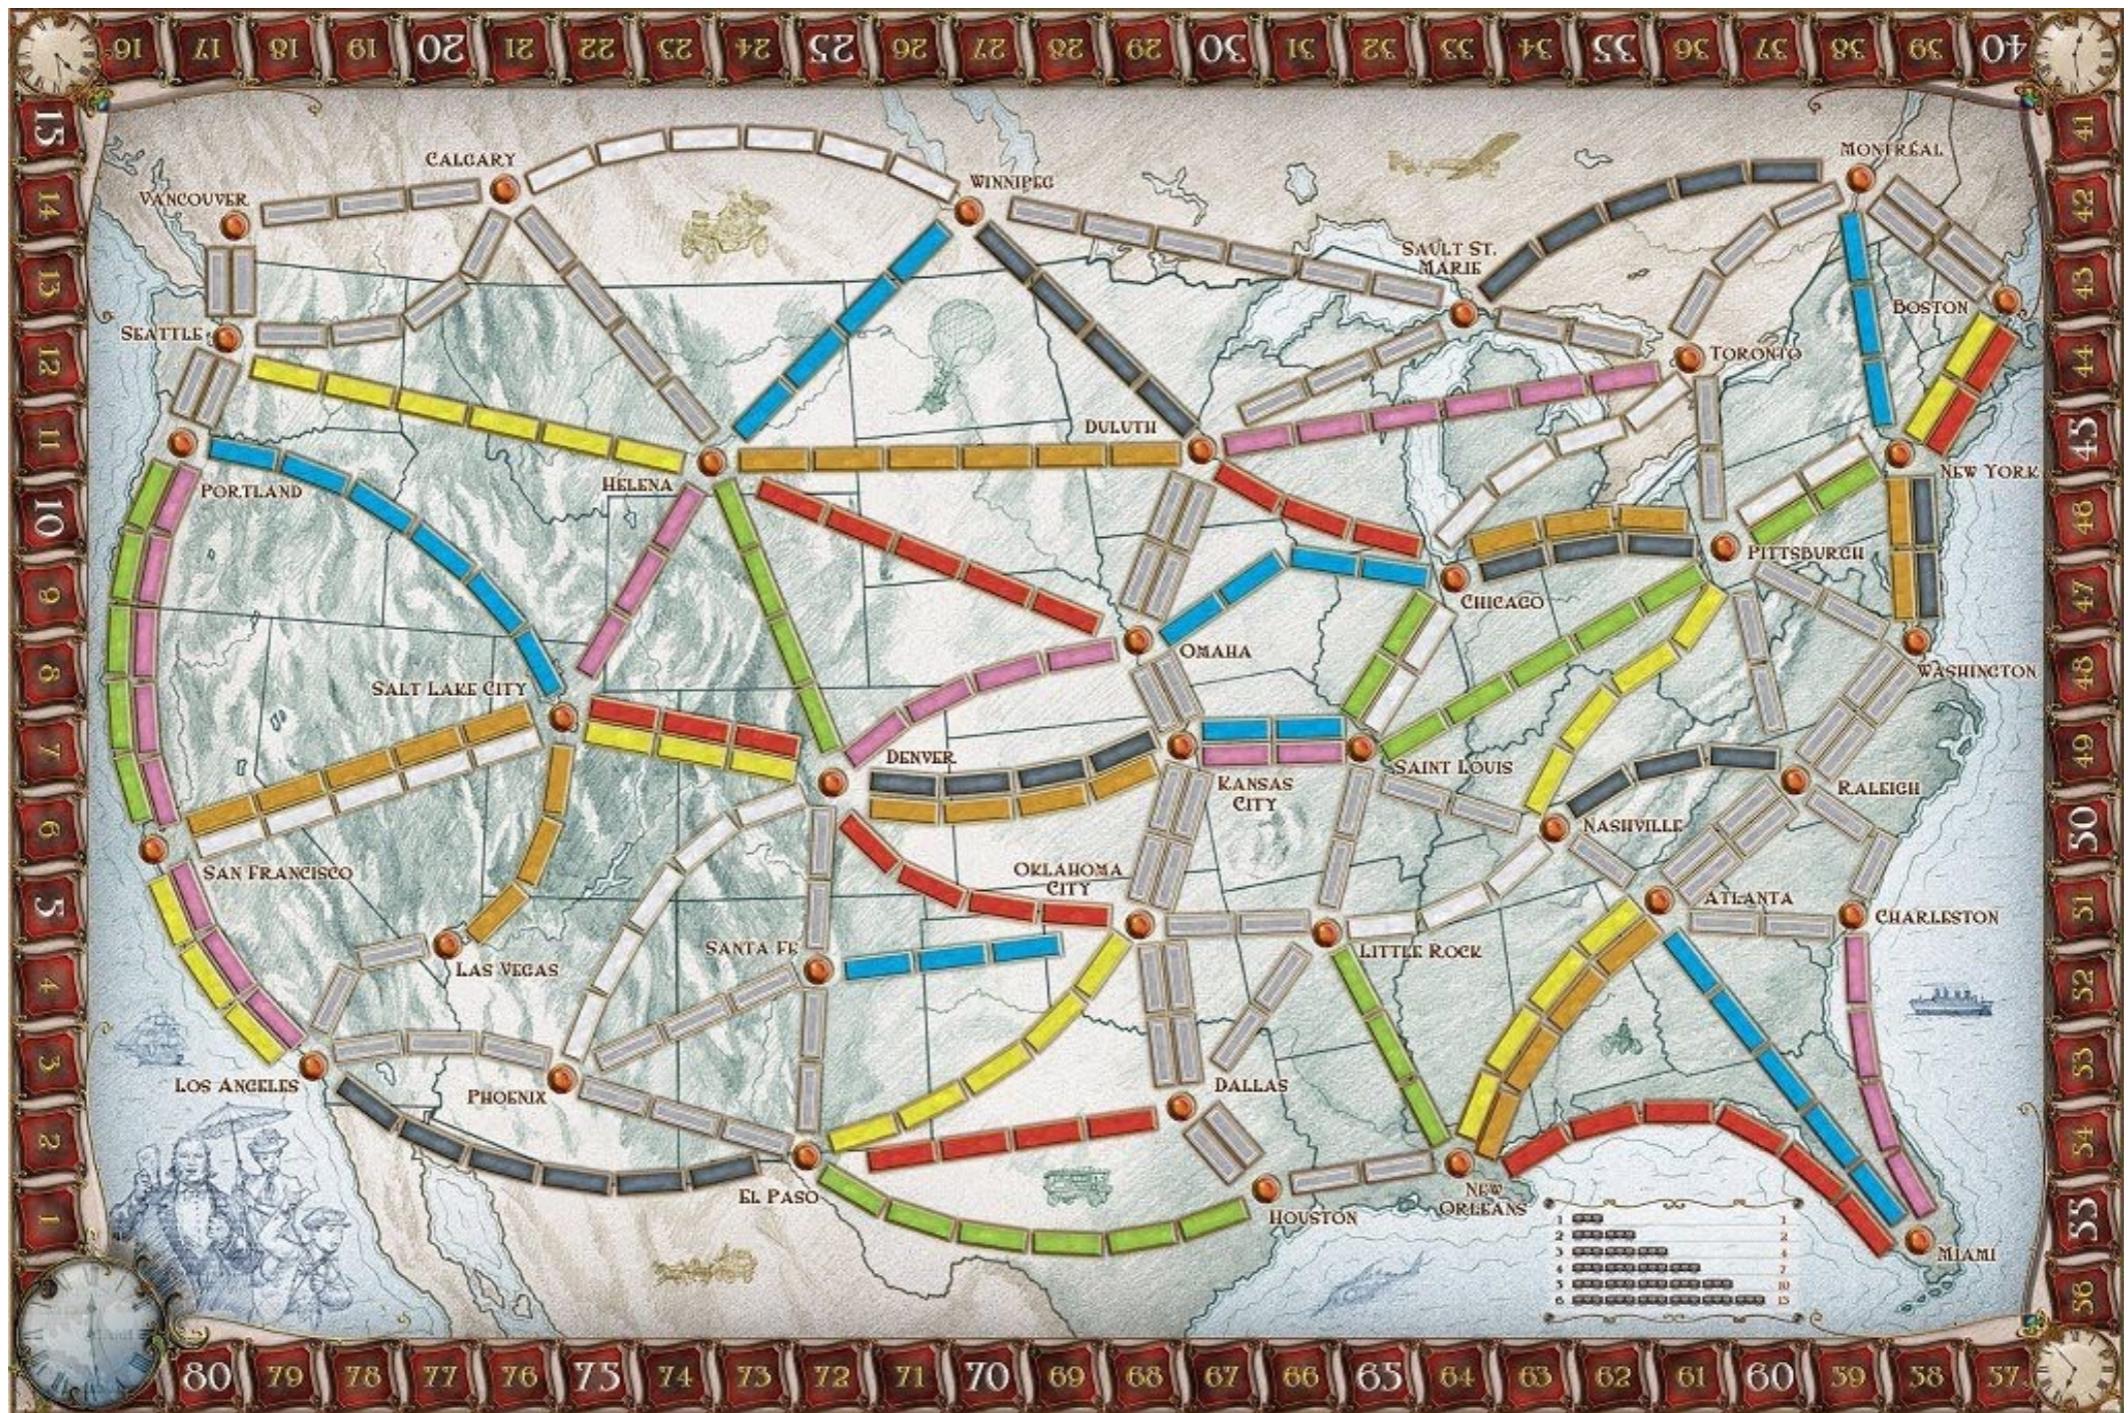
$$V = \{v_1, v_2, v_3, v_4, v_5, v_6\}$$

$$E = \{(v_1, v_2), (v_1, v_3), (v_2, v_3), (v_2, v_5), (v_3, v_4), (v_3, v_6), (v_4, v_5), (v_4, v_6), (v_5, v_6)\}$$

Graphs in Computer Science

- Graphs are used to model all kinds of relational data.
- General purpose algorithms make it possible to solve problems on these models.
 - Shortest Paths, Spanning Tree, Finding Cliques, Strongly Connected Components, Network Flow, Graph Coloring, Minimum Edge/Vertex Cover, Graph Partitioning, ...

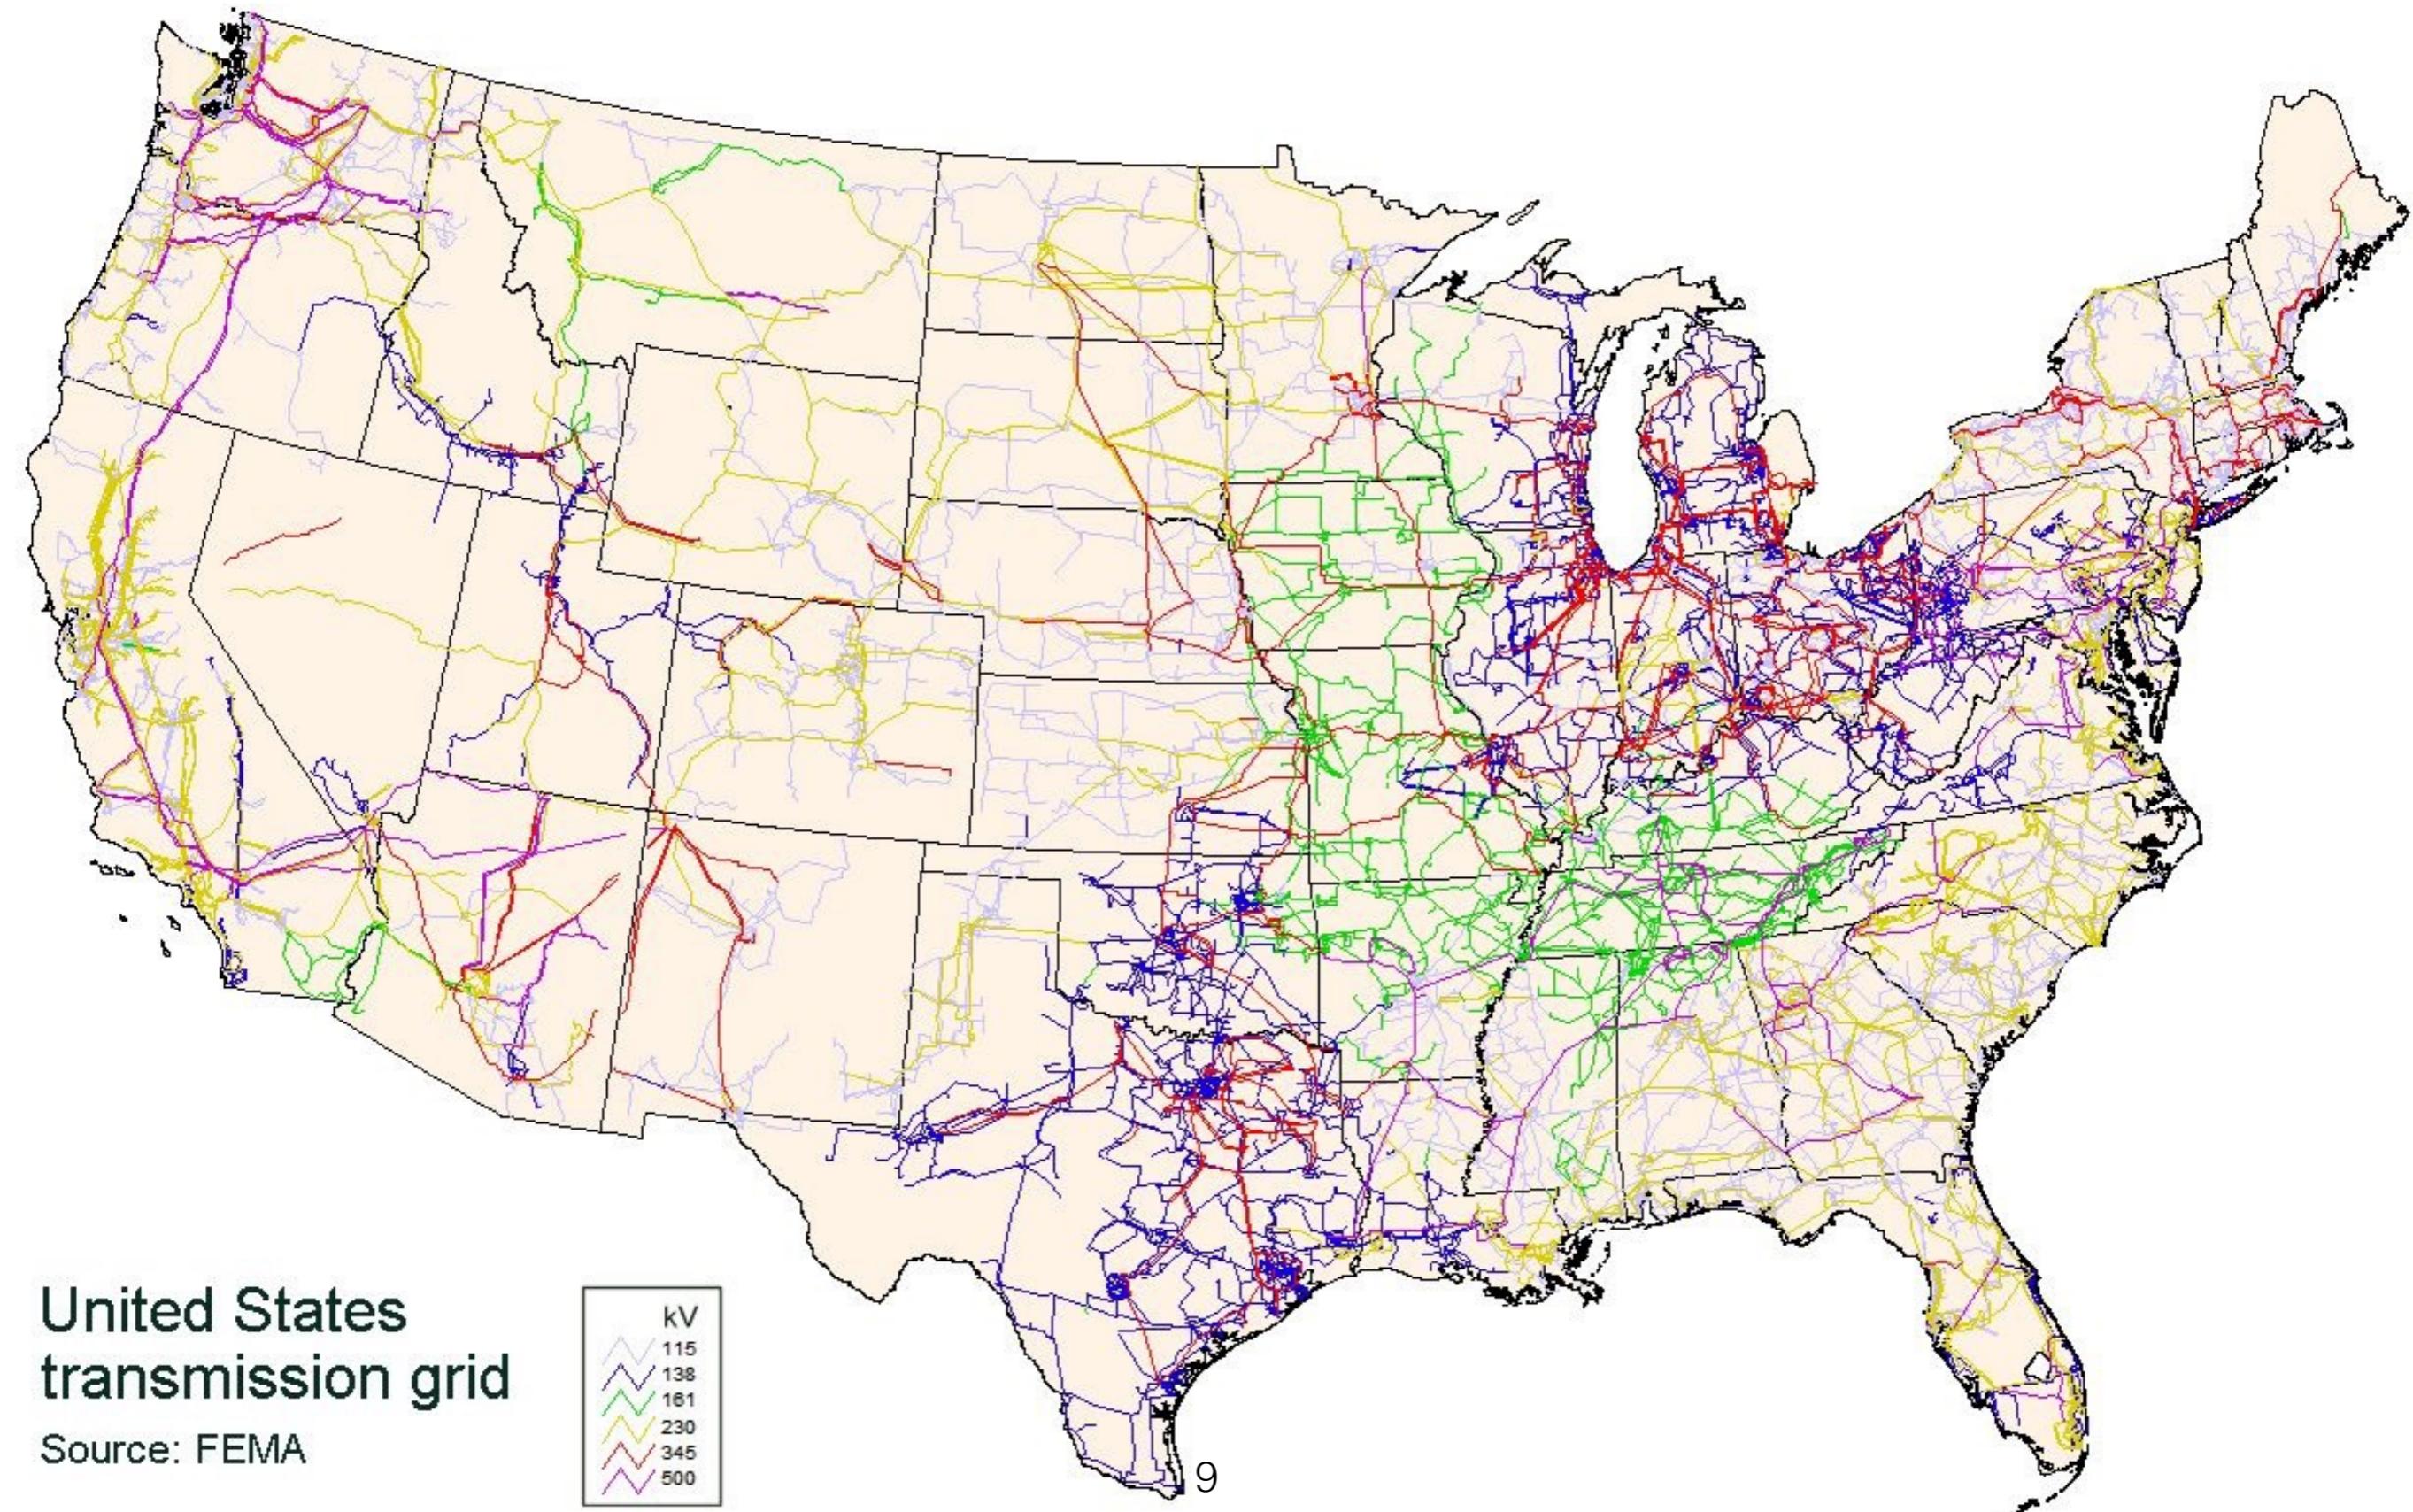
Rail Network



Social Networks

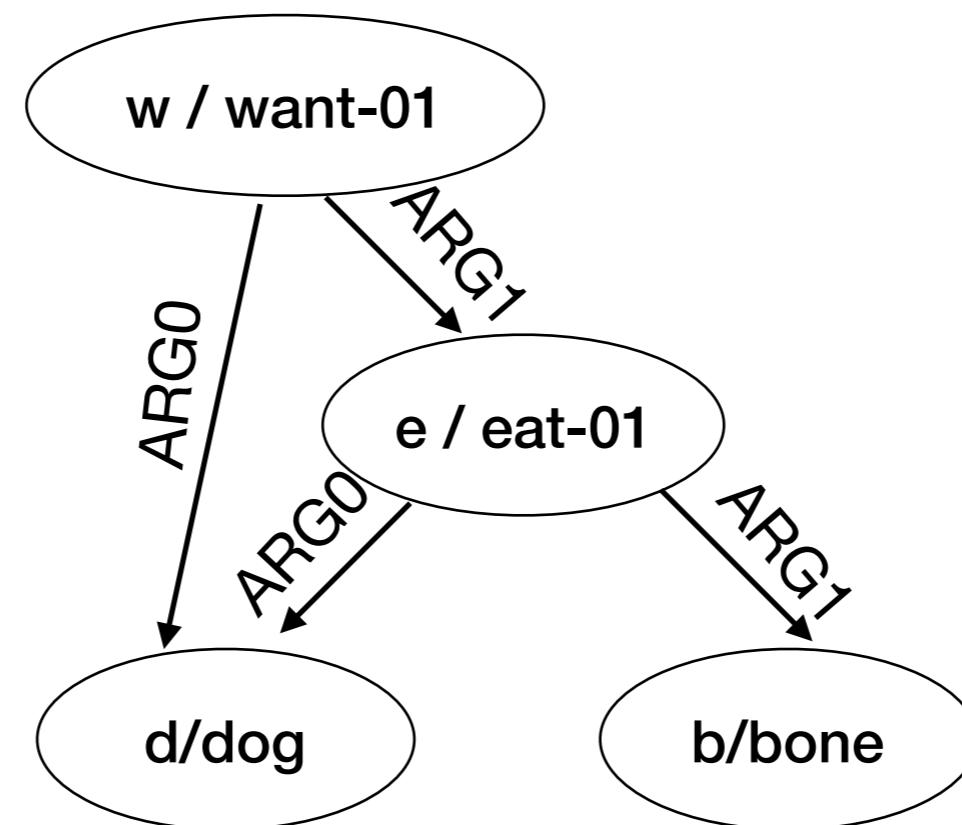


US Power Grid



Graph-Based Representation of Sentence Meaning

The dog wants to eat a bone.



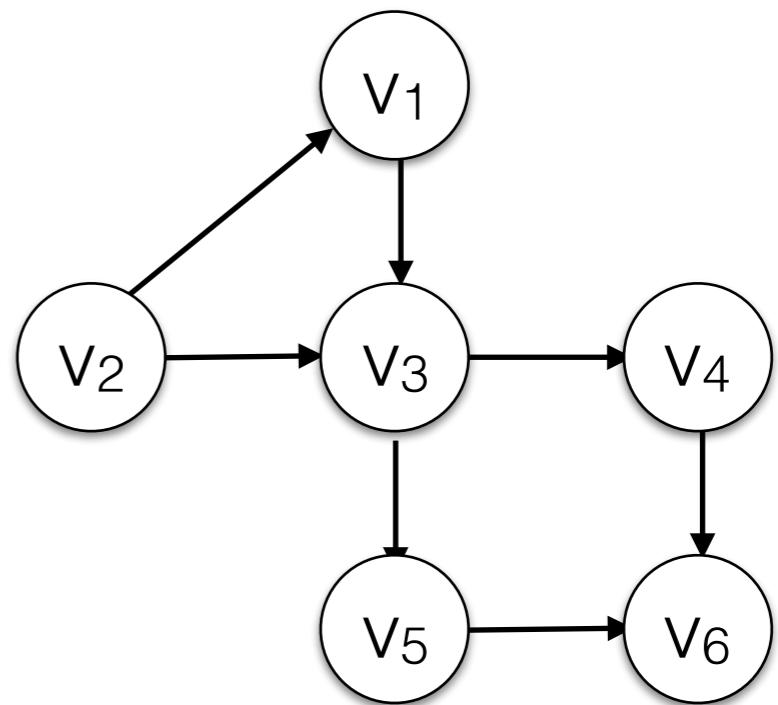
- E.g. "Abstract Meaning Representation"

Edges

- Graphs may be **directed** or **undirected**.
 - In directed graphs, the edge pairs are ordered.

$$V = \{v_1, v_2, v_3, v_4, v_5, v_6\}$$

$$E = \{(v_1, v_3), (v_2, v_1), (v_2, v_3), (v_3, v_4), (v_3, v_5), (v_4, v_6), (v_5, v_6)\}$$



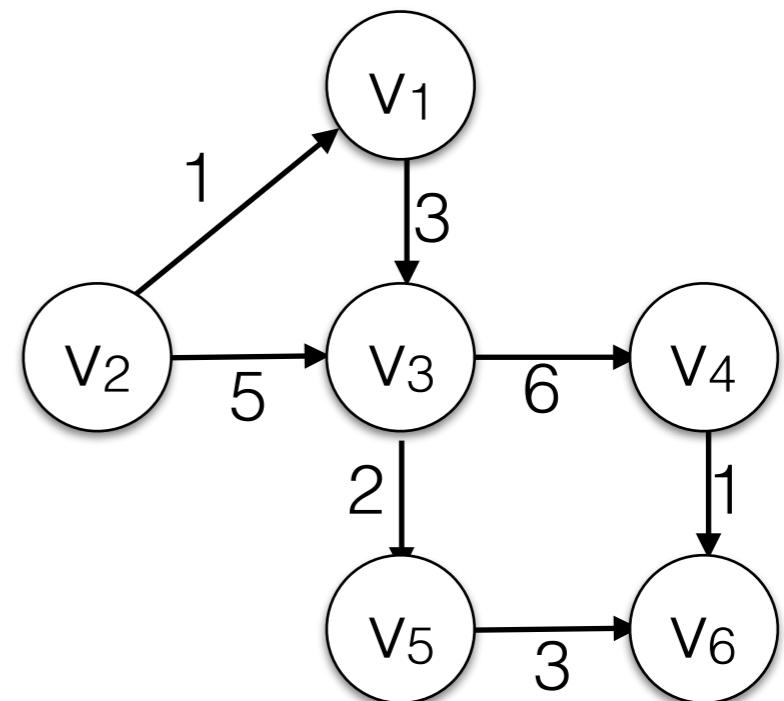
directed graph

Edges

- Graphs may be **directed** or **undirected**.
 - In directed graphs, the edge pairs are ordered.
 - Edges often have some weight or cost associated with them (**weighted** graphs).

$$V = \{v_1, v_2, v_3, v_4, v_5, v_6\}$$

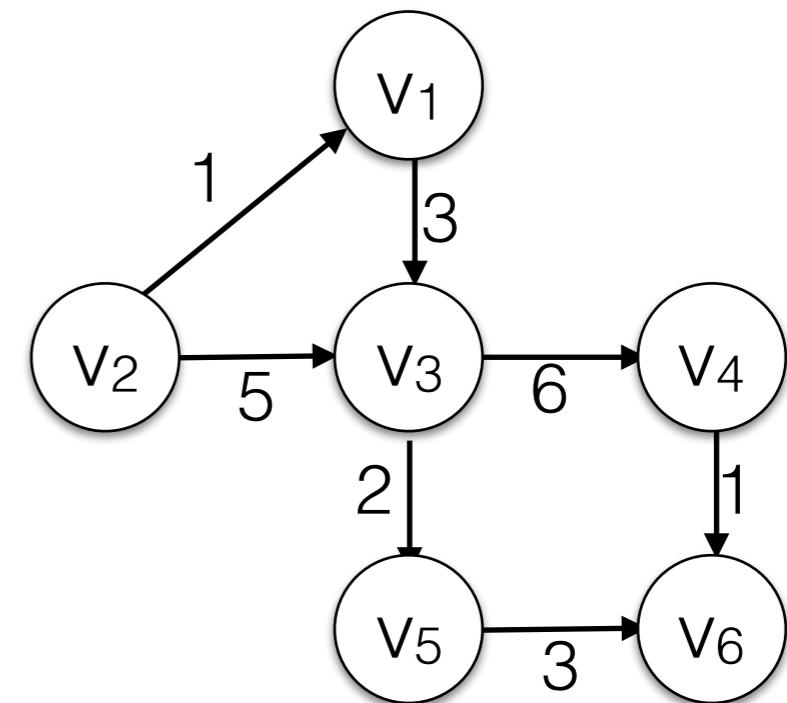
$$E = \{(v_1, v_3), (v_2, v_1), (v_2, v_3), (v_3, v_4), (v_3, v_5), (v_4, v_6), (v_5, v_6)\}$$



12 directed and weighted graph

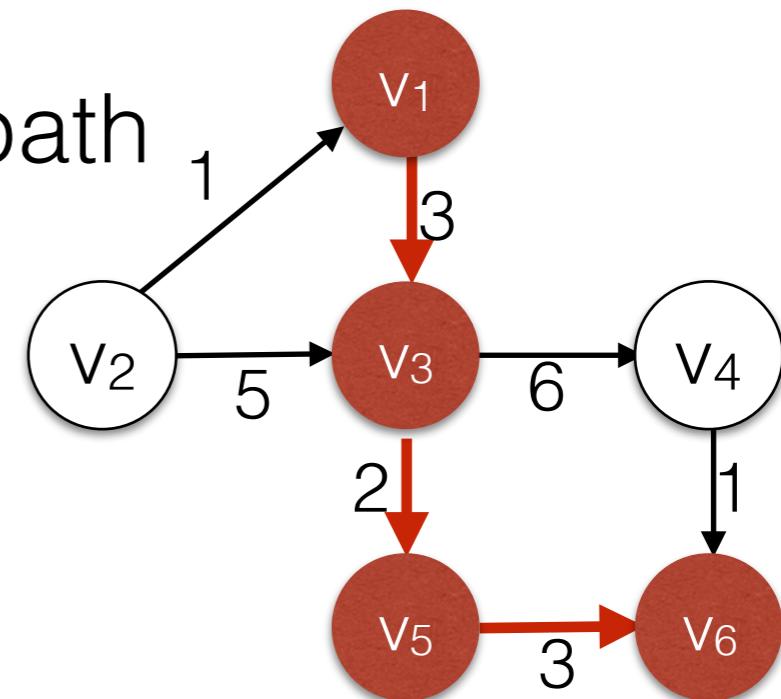
Paths

- Vertex w is **adjacent** to vertex v iff $(w,v) \in E$.
- A **path** is a sequence of vertices w_1, w_2, \dots, w_k such that $(w_i, w_{i+1}) \in E$.



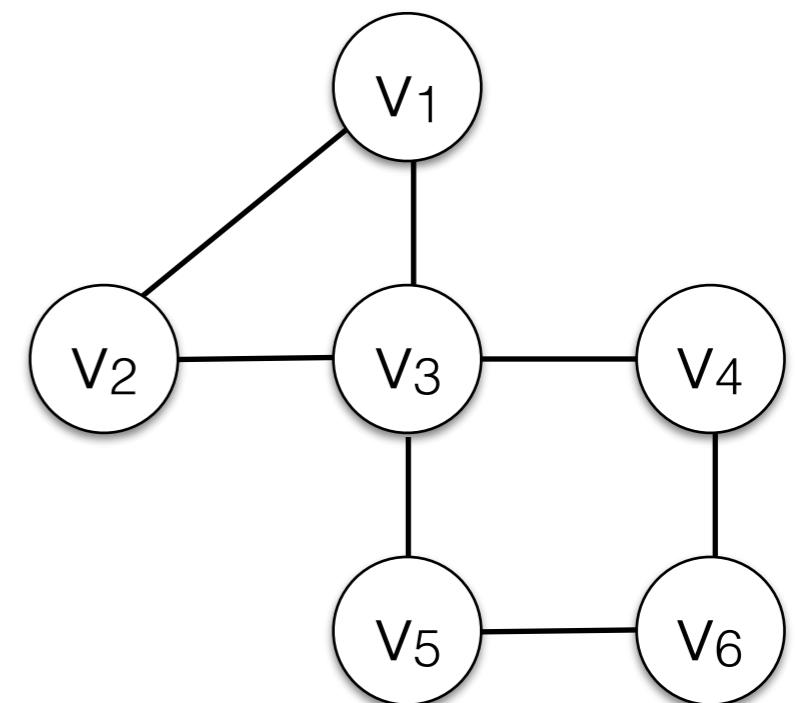
Paths

- Vertex w is **adjacent** to vertex v iff $(w,v) \in E$.
- A **path** is a sequence of vertices w_1, w_2, \dots, w_k such that $(w_i, w_{i+1}) \in E$.
- **length** of a path:
 $k-1$ = number of edges on path
- **cost** of a path:
Sum of all edge costs.



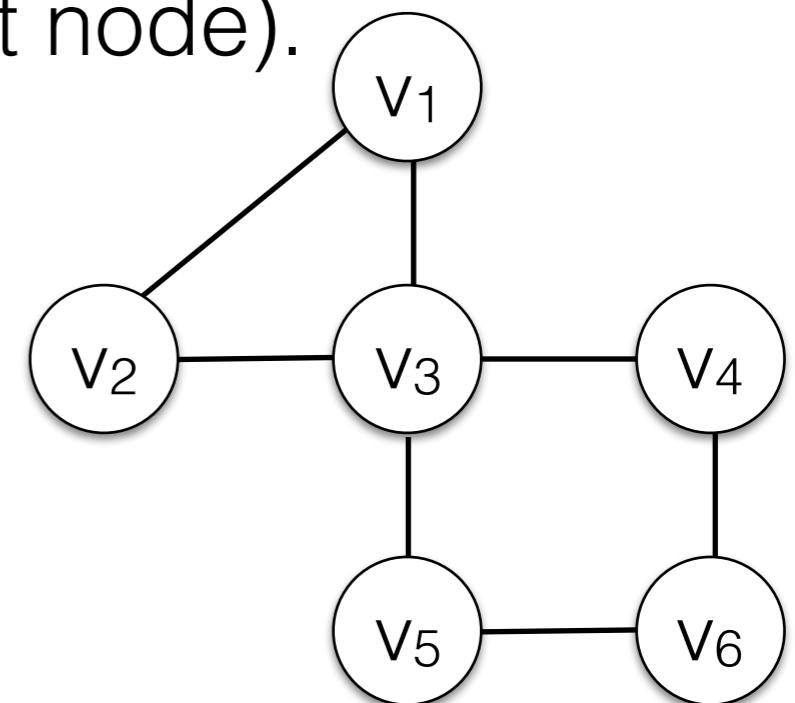
Path from v_1 to v_6 , length 3, cost 8
 $(v_1, v_3), (v_3, v_5), (v_5, v_6)$

Simple Paths



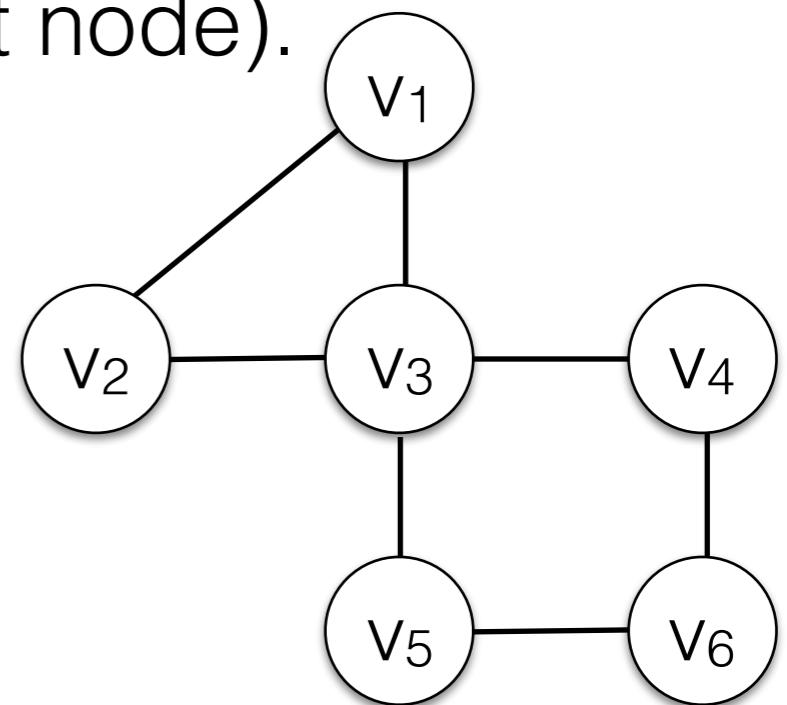
Simple Paths

- A **simple path** is a path that contains every node only once (except possibly the first and last node).



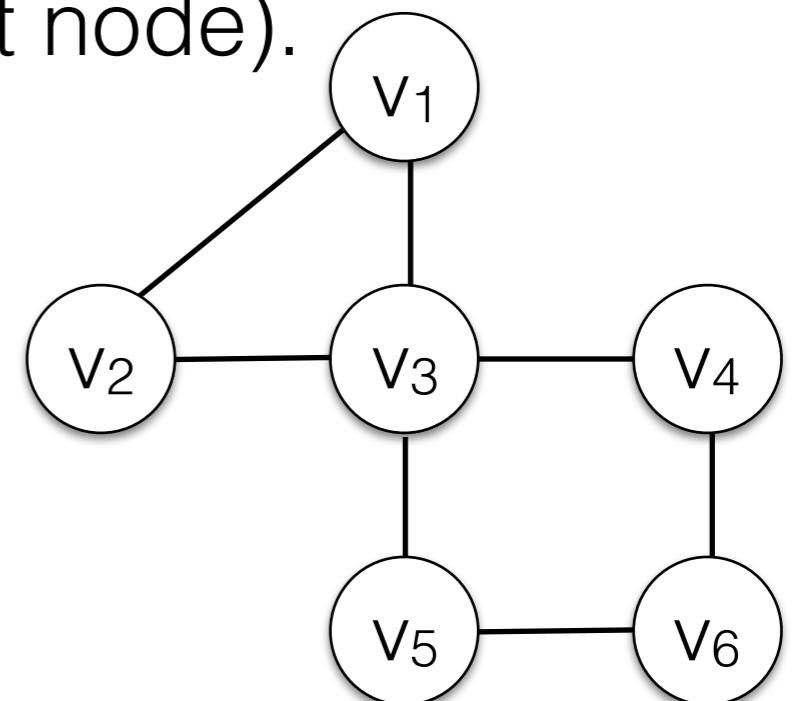
Simple Paths

- A **simple path** is a path that contains every node only once (except possibly the first and last node).
- $(v_2, v_3, v_4, v_6, v_5, v_3, v_1)$ is a path but not a simple path.



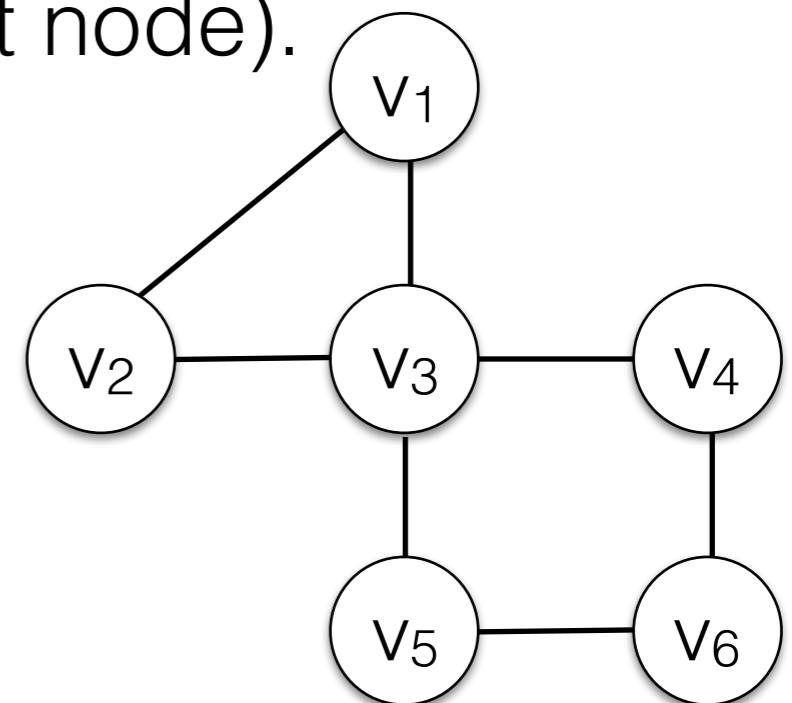
Simple Paths

- A **simple path** is a path that contains every node only once (except possibly the first and last node).
- $(v_2, v_3, v_4, v_6, v_5, v_3, v_1)$ is a path but not a simple path.
- There are only two simple paths between v_2 and v_1 : (v_2, v_1) and (v_2, v_3, v_1)



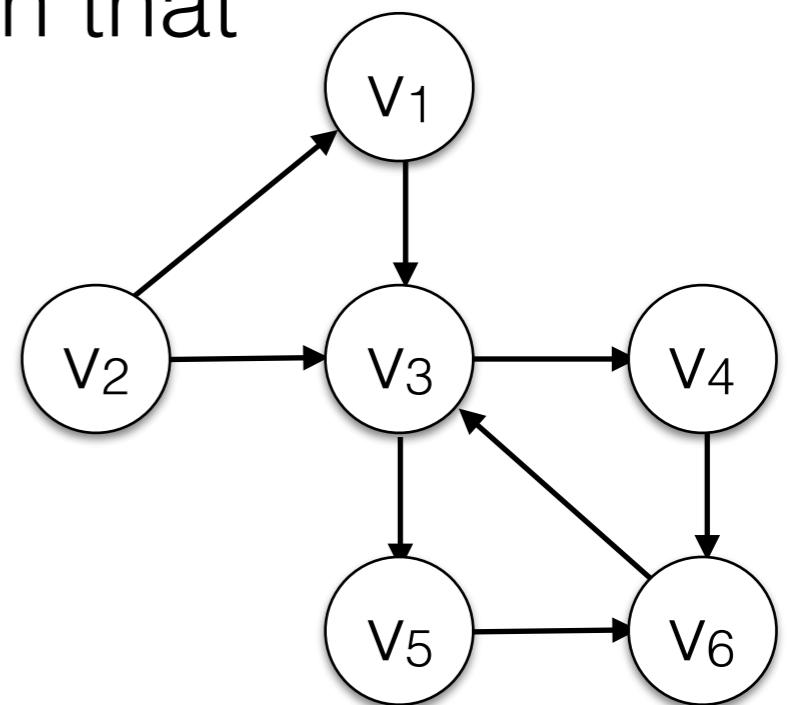
Simple Paths

- A **simple path** is a path that contains every node only once (except possibly the first and last node).
- $(v_2, v_3, v_4, v_6, v_5, v_3, v_1)$ is a path but not a simple path.
- There are only two simple paths between v_2 and v_1 : (v_2, v_1) and (v_2, v_3, v_1)
- (v_1, v_3, v_2, v_1) is a simple path.



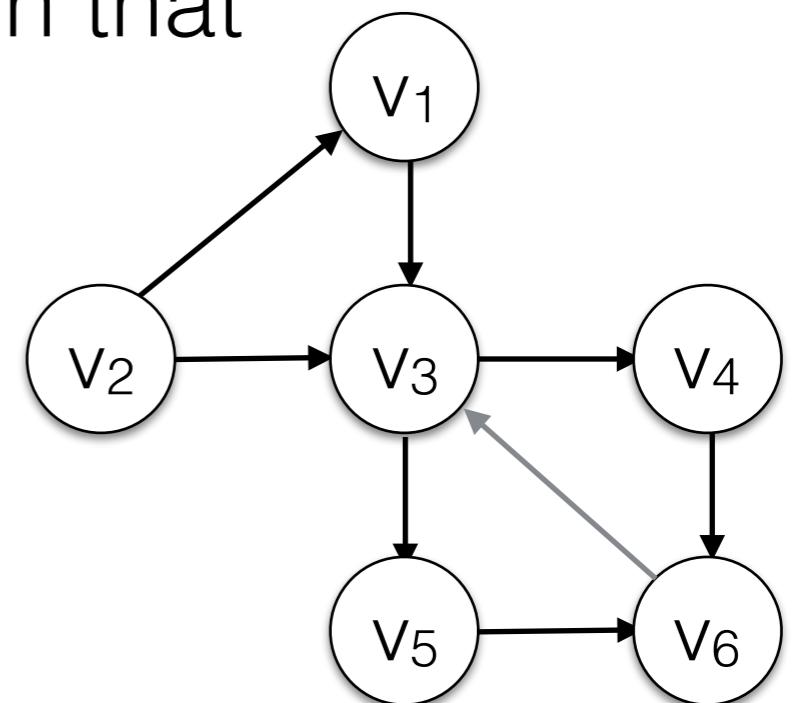
Cycles in Directed Graphs

- A **cycle** is a path (of length > 1) such that
 $w_1 = w_k$
- (v_3, v_4, v_6, v_3) is a cycle.



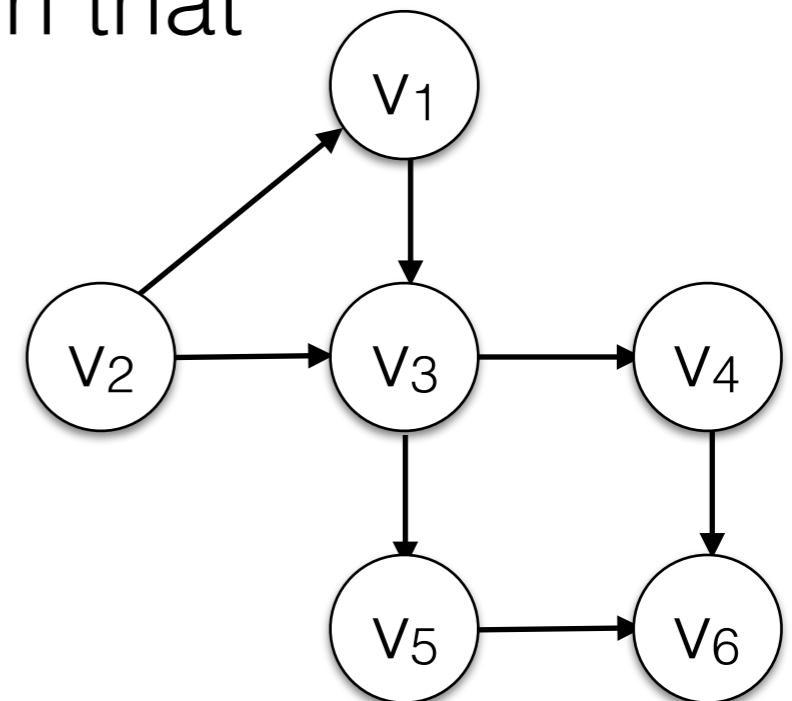
Cycles in Directed Graphs

- A **cycle** is a path (of length > 1) such that
 $w_1 = w_k$
- (v_3, v_4, v_6, v_3) is a cycle.
- A **Directed Acyclic Graph (DAG)** is a directed graph that contains no cycles.

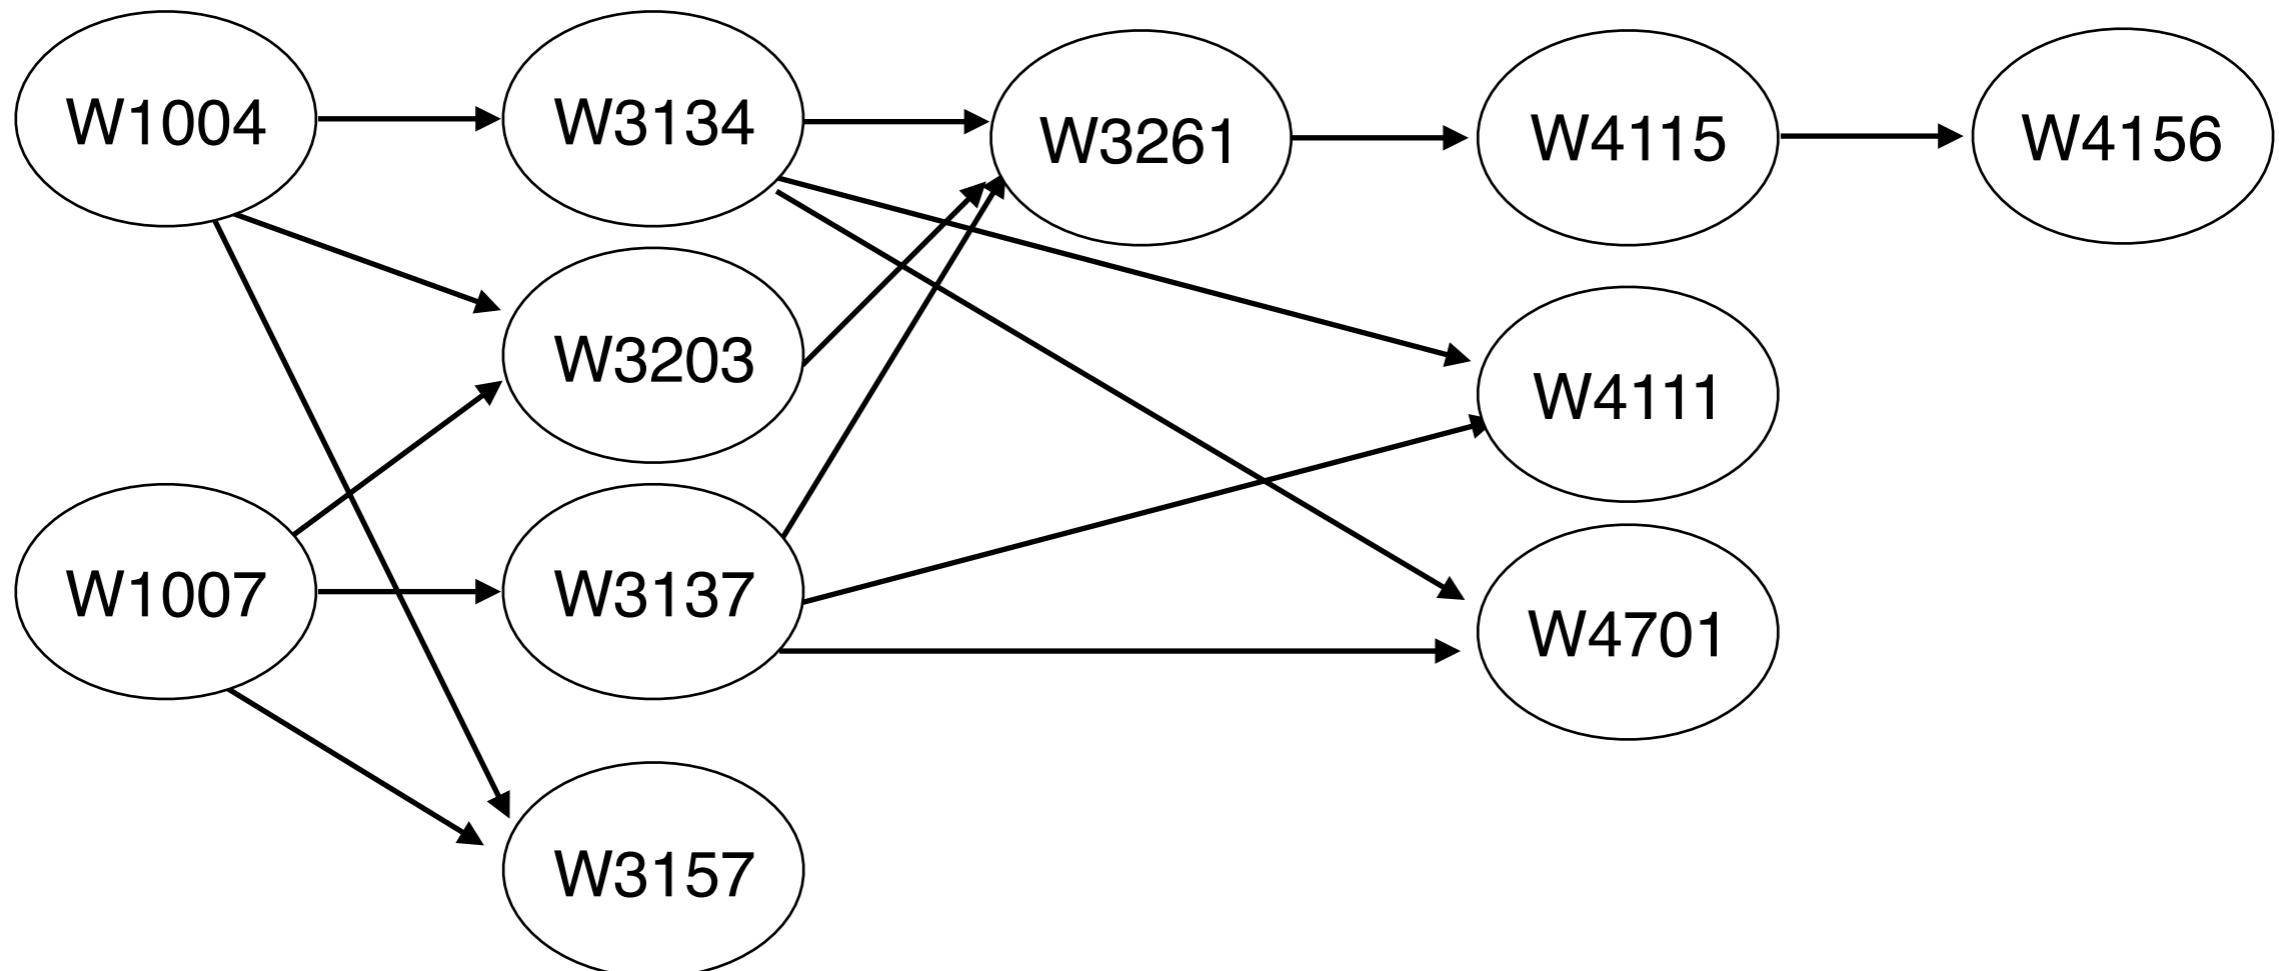


Cycles in Directed Graphs

- A **cycle** is a path (of length > 1) such that $w_1 = w_k$
- (v_3, v_4, v_6, v_3) is a cycle.
- A **Directed Acyclic Graph (DAG)** is a directed graph that contains no cycles.



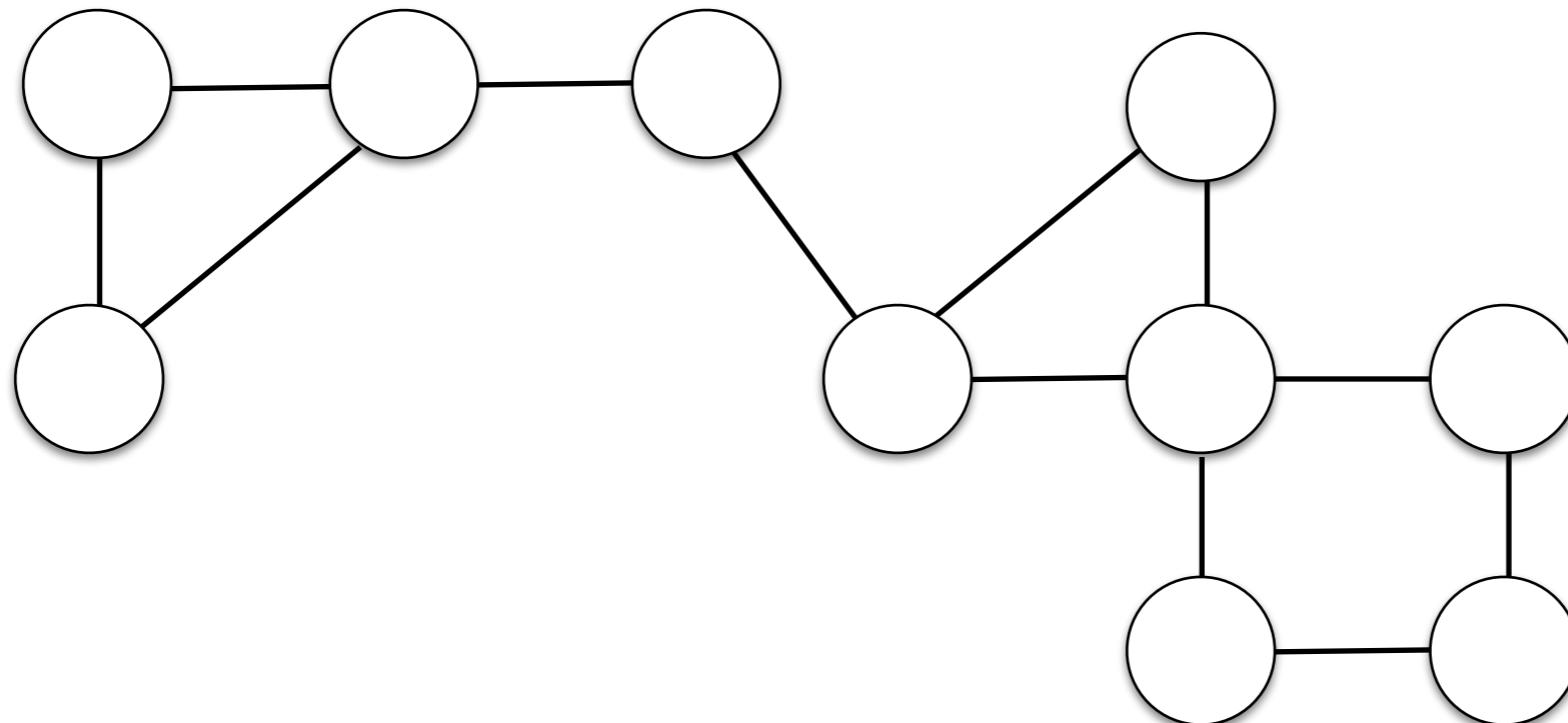
Columbia CS Course Prerequisites as a DAG



Please do not use this figure for program planning!
It's inaccurate!

Connectivity

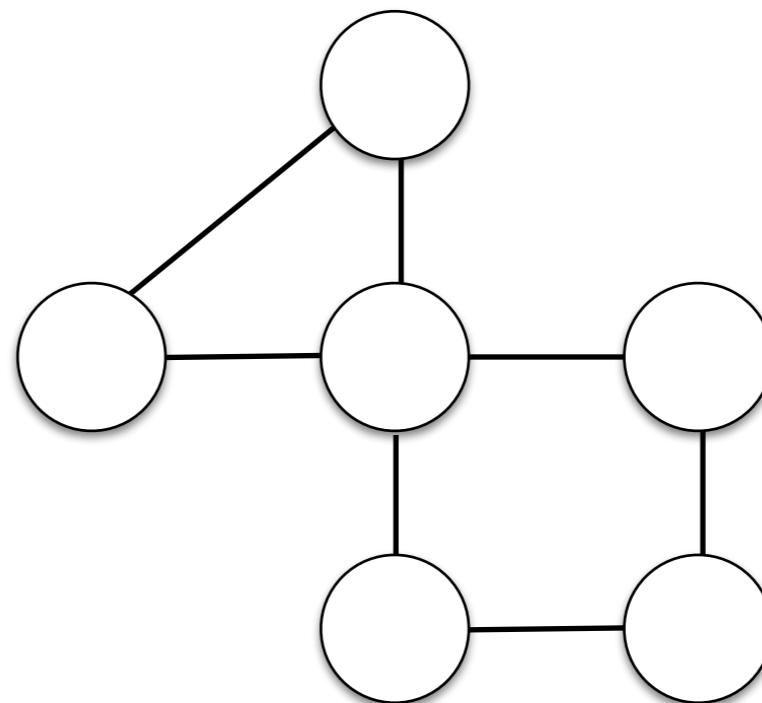
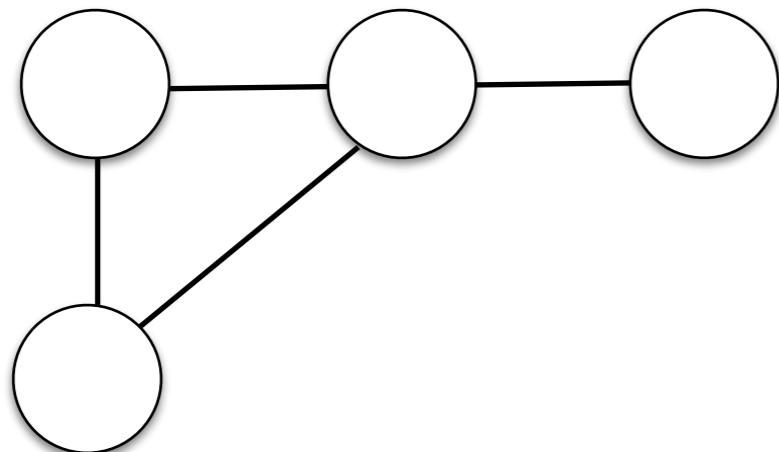
- An undirected graph is **connected** if there is a path from every vertex to every other vertex.



connected graph

Connectivity

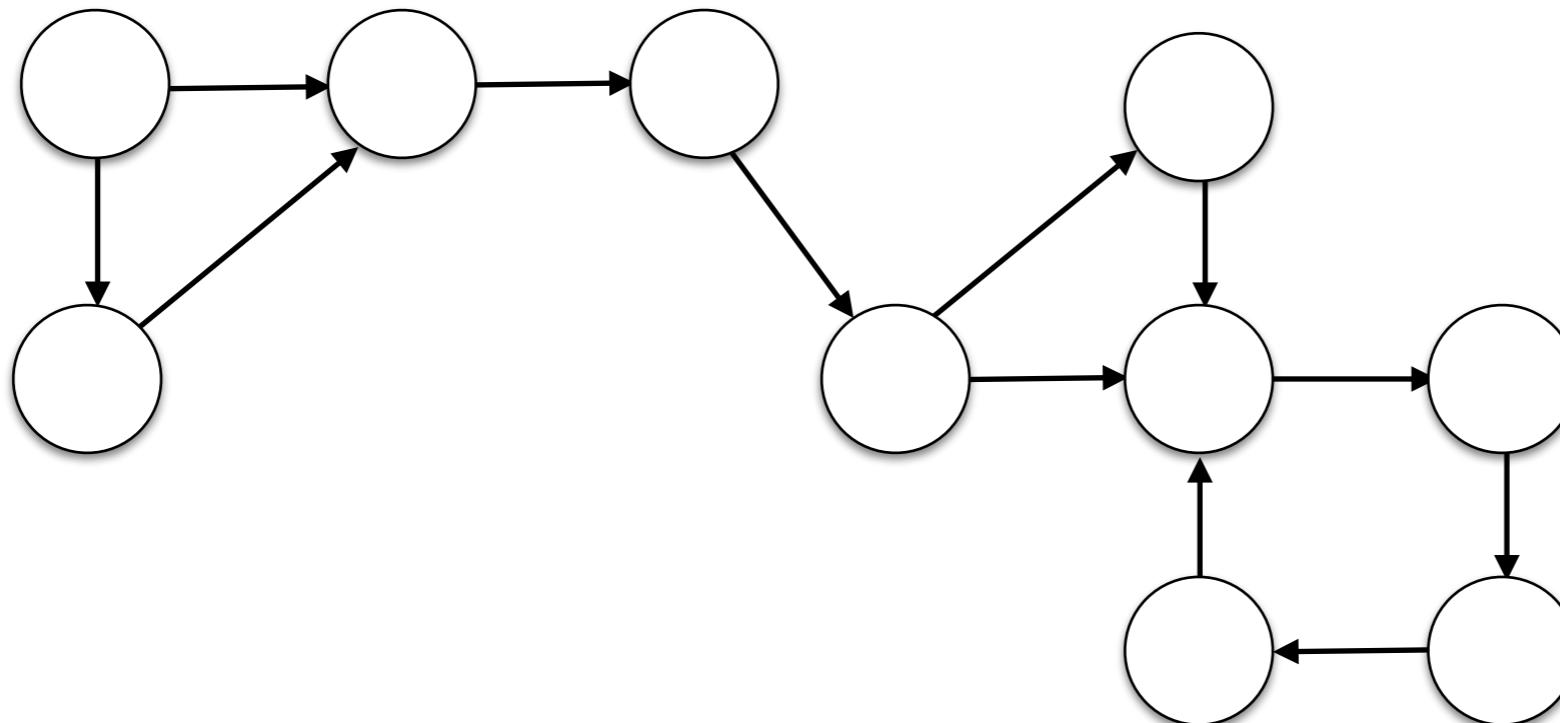
- An undirected graph is **connected** if there is a path from every vertex to every other vertex.



unconnected graph

Connectivity in Directed Graphs

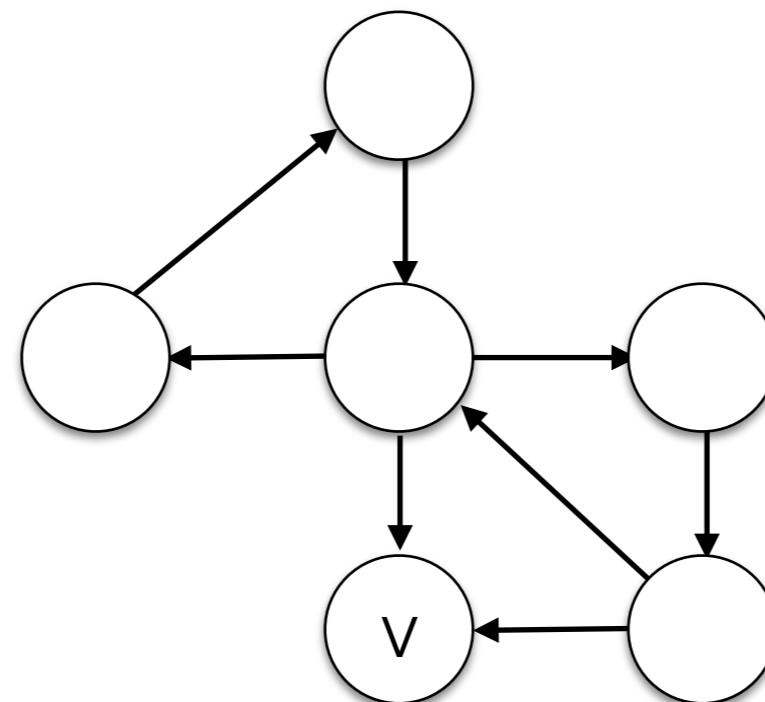
- A directed graph is **weakly connected** if there is an *undirected* path from every vertex to every other vertex.



weakly connected graph

Strongly Connected Graphs

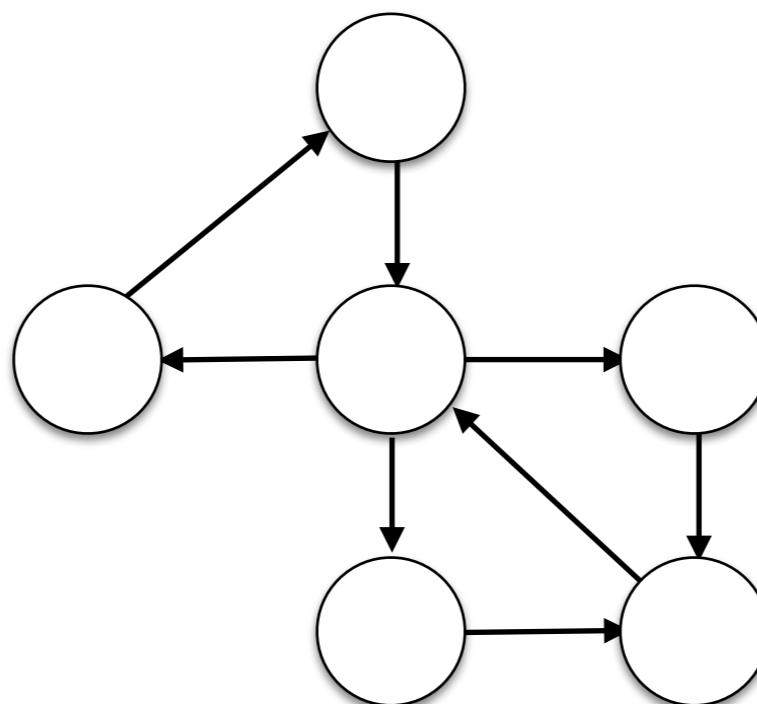
- A directed graph is **strongly connected** if there is a path from every vertex to every other vertex.



Weakly connected, but not strongly connected (no other vertex can be reached from v).

Strongly Connected Graphs

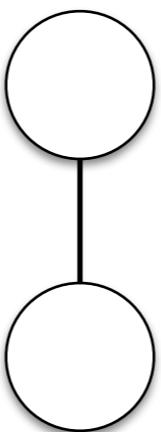
- A directed graph is **strongly connected** if there is a path from every vertex to every other vertex.



strongly connected

Complete Graphs

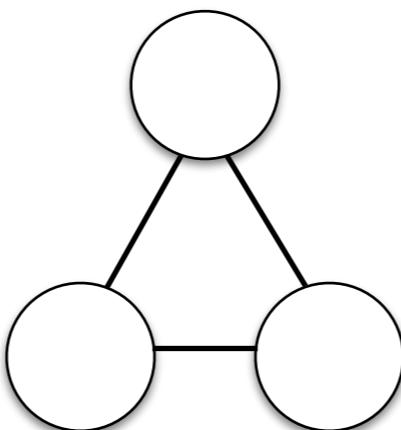
- A **complete graph** has edges between every pair of vertices.



N=2

Complete Graphs

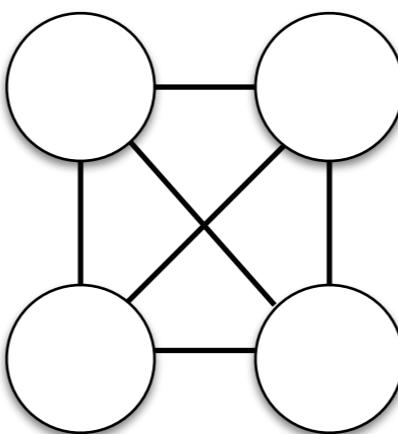
- A **complete graph** has edges between every pair of vertices.



$N=3$

Complete Graphs

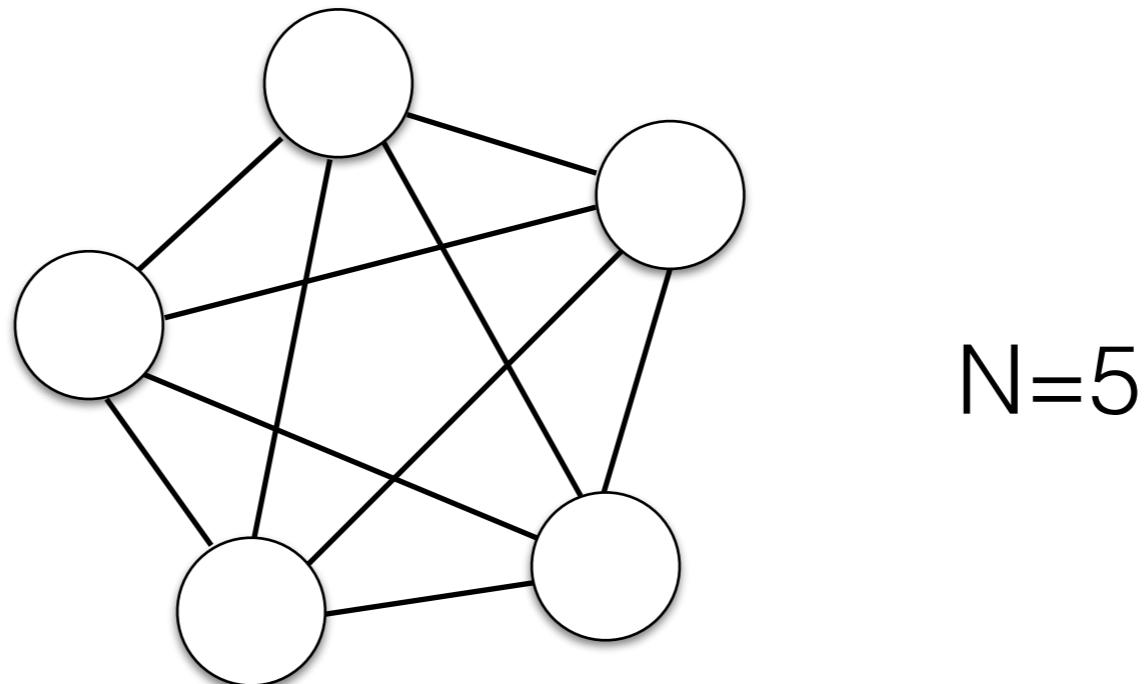
- A **complete graph** has edges between every pair of vertices.



$N=4$

Complete Graphs

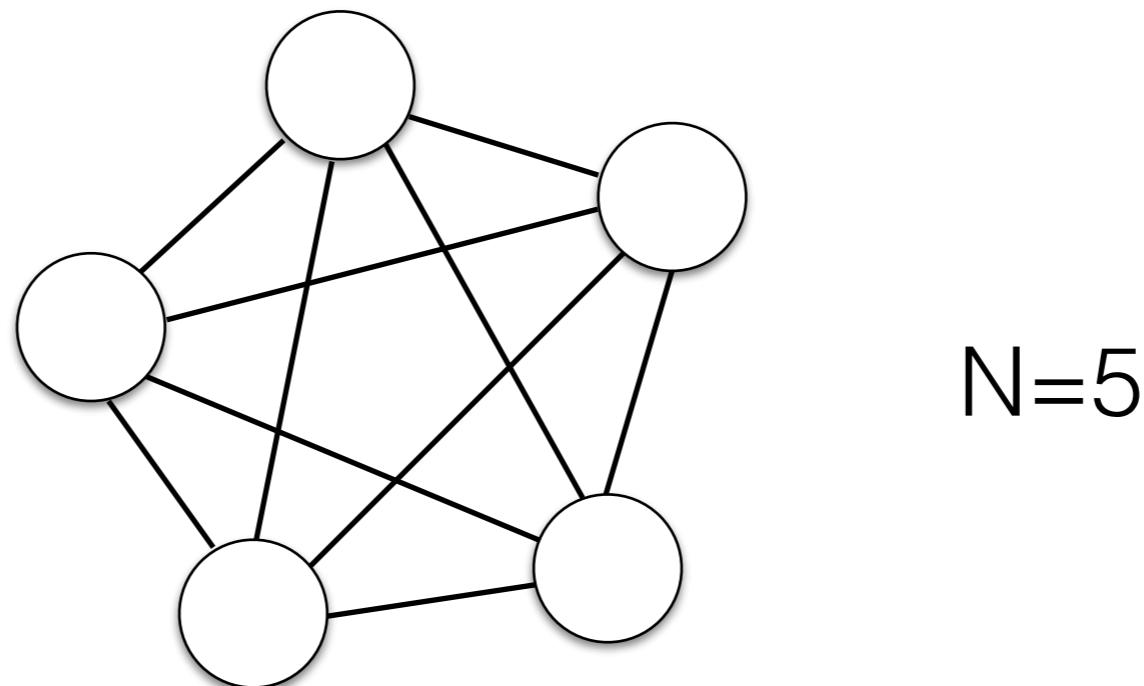
- A **complete graph** has edges between every pair of vertices.



How many edges are there in a complete graph of size N?

Complete Graphs

- A **complete graph** has edges between every pair of vertices.



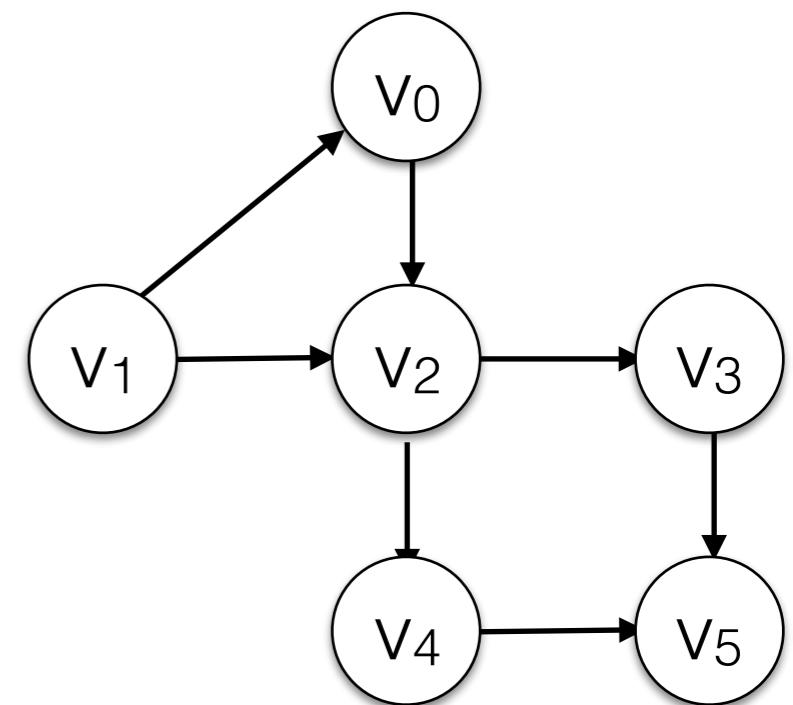
How many edges are there in a complete graph of size N?

$$\sum_{i=1}^{N-1} i = \frac{N \cdot (N - 1)}{2}$$

Representing Graphs

- Represent graph $G = (E, V)$, option 1:
 - $N \times N$ **Adjacency Matrix** represented as 2-dimensional Boolean[][].
 - $A[u][v] = \text{true}$ if $(u, v) \in E$, else false

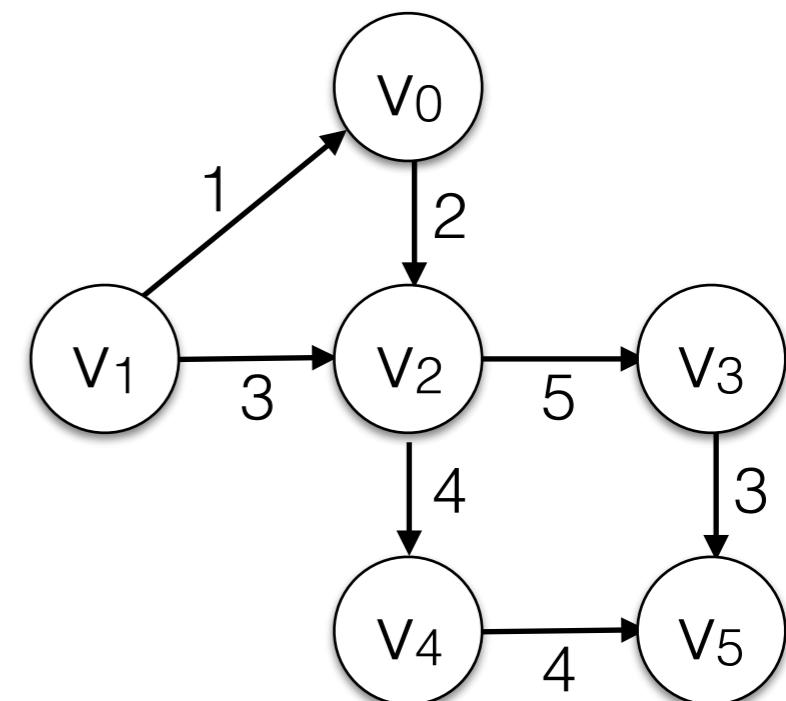
| | 0 | 1 | 2 | 3 | 4 | 5 |
|---|---|---|---|---|---|---|
| 0 | f | f | t | f | f | f |
| 1 | t | f | t | f | f | f |
| 2 | f | f | f | t | t | f |
| 3 | f | f | f | f | f | t |
| 4 | f | f | f | f | f | t |
| 5 | f | f | f | f | f | f |



Representing Graphs

- Represent graph $G = (E, V)$, option 1:
 - $N \times N$ **Adjacency Matrix** represented as 2-dimensional Integer[][].
 - $A[u][v] = \text{cost}(u, v)$ if $(u, v) \in E$, else ∞

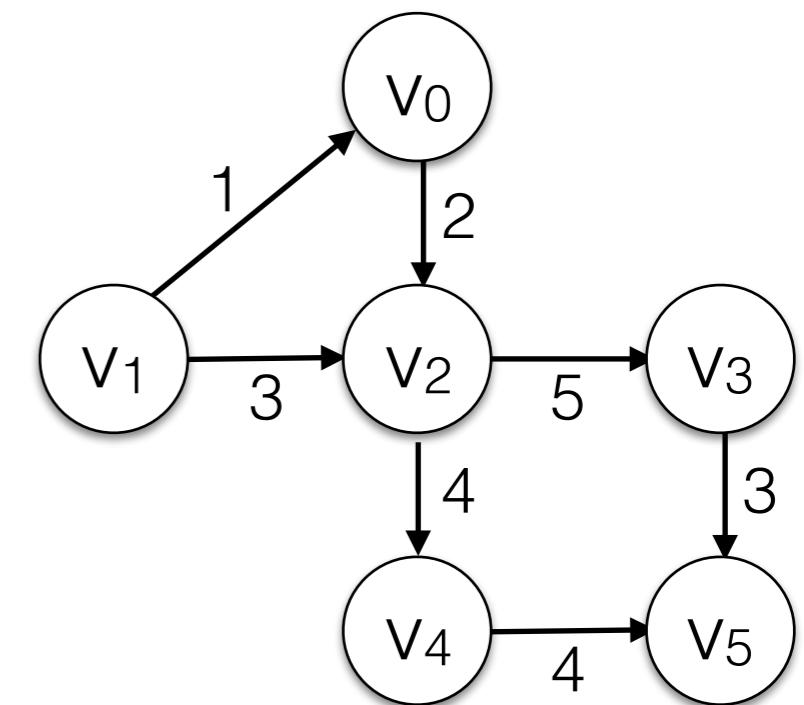
| | 0 | 1 | 2 | 3 | 4 | 5 |
|---|----------|----------|----------|----------|----------|----------|
| 0 | ∞ | ∞ | 2 | ∞ | ∞ | ∞ |
| 1 | 1 | ∞ | 3 | ∞ | ∞ | ∞ |
| 2 | ∞ | ∞ | ∞ | 5 | 4 | ∞ |
| 3 | ∞ | ∞ | ∞ | ∞ | ∞ | 3 |
| 4 | ∞ | ∞ | ∞ | ∞ | ∞ | 4 |
| 5 | ∞ | ∞ | ∞ | ∞ | ∞ | ∞ |



Representing Graphs

- Problem of Adjacency Matrix representation:
 - For **sparse** graphs (that contain much less than $|V|^2$ edges), a lot of array space is wasted.

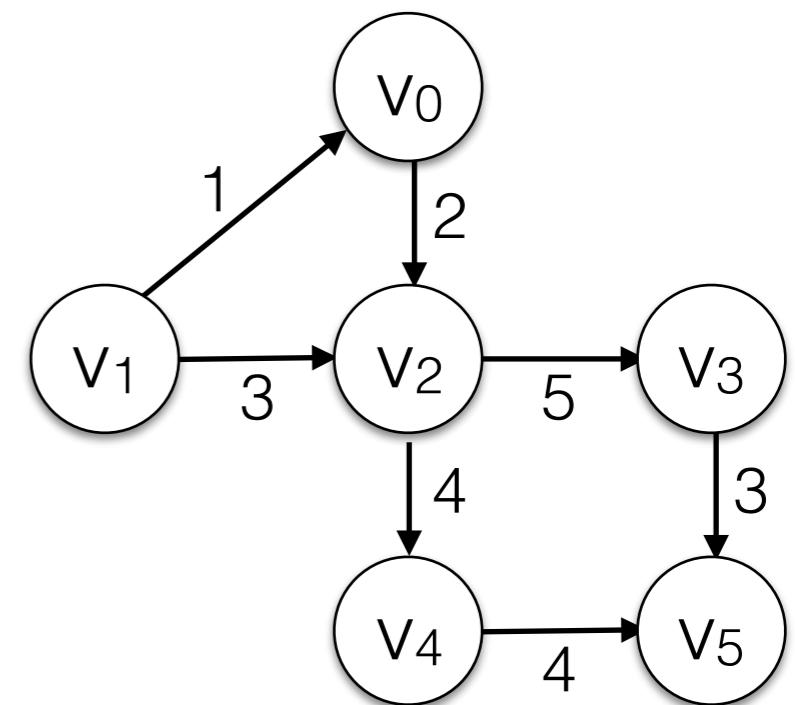
| | 0 | 1 | 2 | 3 | 4 | 5 |
|---|----------|----------|----------|----------|----------|----------|
| 0 | ∞ | ∞ | 2 | ∞ | ∞ | ∞ |
| 1 | 1 | ∞ | 3 | ∞ | ∞ | ∞ |
| 2 | ∞ | ∞ | ∞ | 5 | 4 | ∞ |
| 3 | ∞ | ∞ | ∞ | ∞ | ∞ | 3 |
| 4 | ∞ | ∞ | ∞ | ∞ | ∞ | 4 |
| 5 | ∞ | ∞ | ∞ | ∞ | ∞ | ∞ |



Representing Graphs

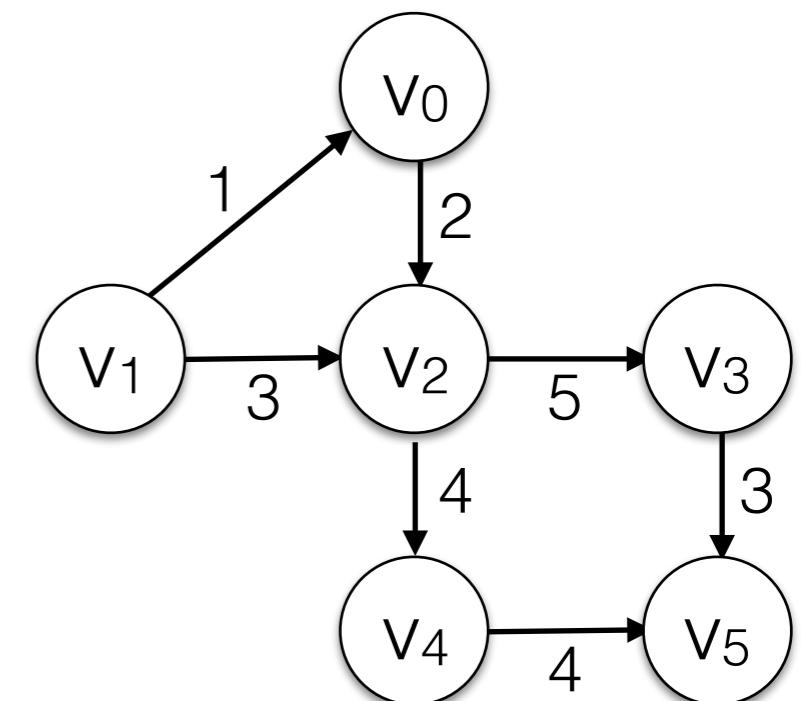
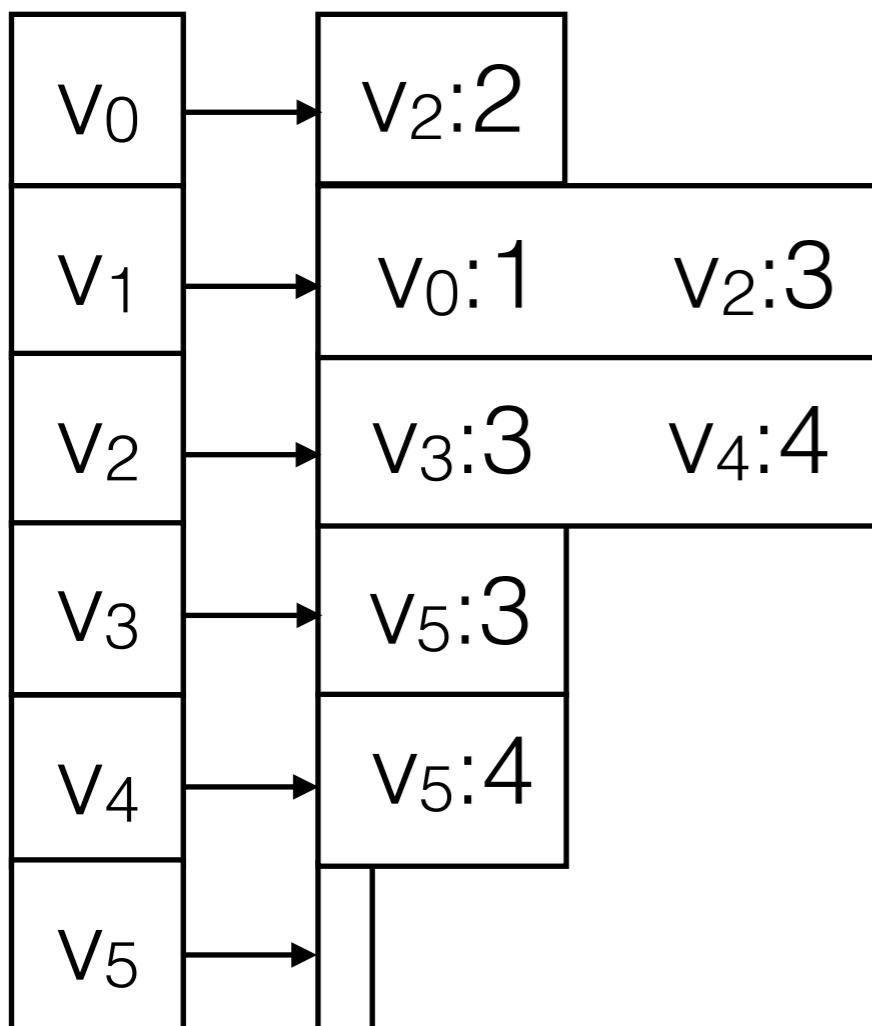
- Problem of Adjacency Matrix representation:
Space requirement: $\Theta(|V|^2)$
 - For **sparse** graphs (that contain much less than $|V|^2$ edges), a lot of array space is wasted.

| | 0 | 1 | 2 | 3 | 4 | 5 |
|---|----------|----------|----------|----------|----------|----------|
| 0 | ∞ | ∞ | 2 | ∞ | ∞ | ∞ |
| 1 | 1 | ∞ | 3 | ∞ | ∞ | ∞ |
| 2 | ∞ | ∞ | ∞ | 5 | 4 | ∞ |
| 3 | ∞ | ∞ | ∞ | ∞ | ∞ | 3 |
| 4 | ∞ | ∞ | ∞ | ∞ | ∞ | 4 |
| 5 | ∞ | ∞ | ∞ | ∞ | ∞ | ∞ |



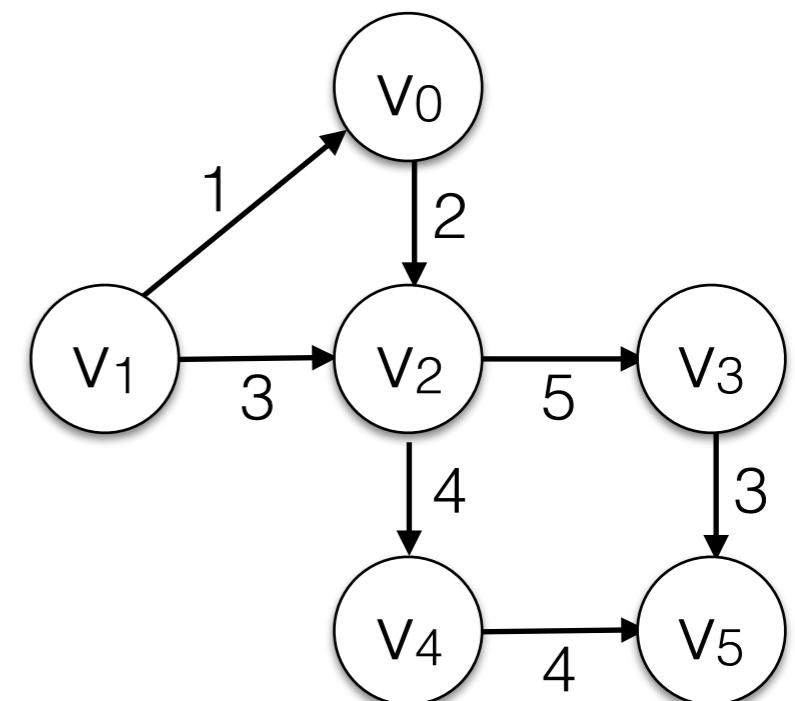
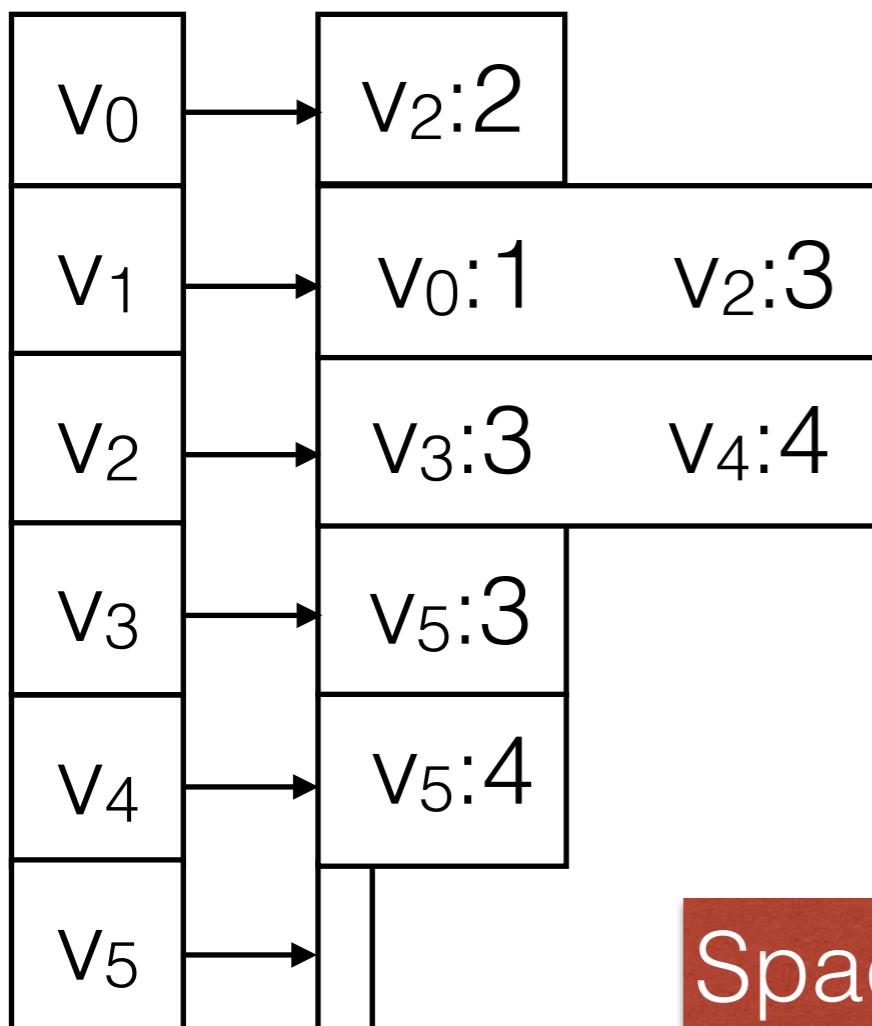
Representing Graphs

- Represent graph $G = (E, V)$, option 2: **Adjacency Lists**
 - For each vertex, keep a list of all adjacent vertices.



Representing Graphs

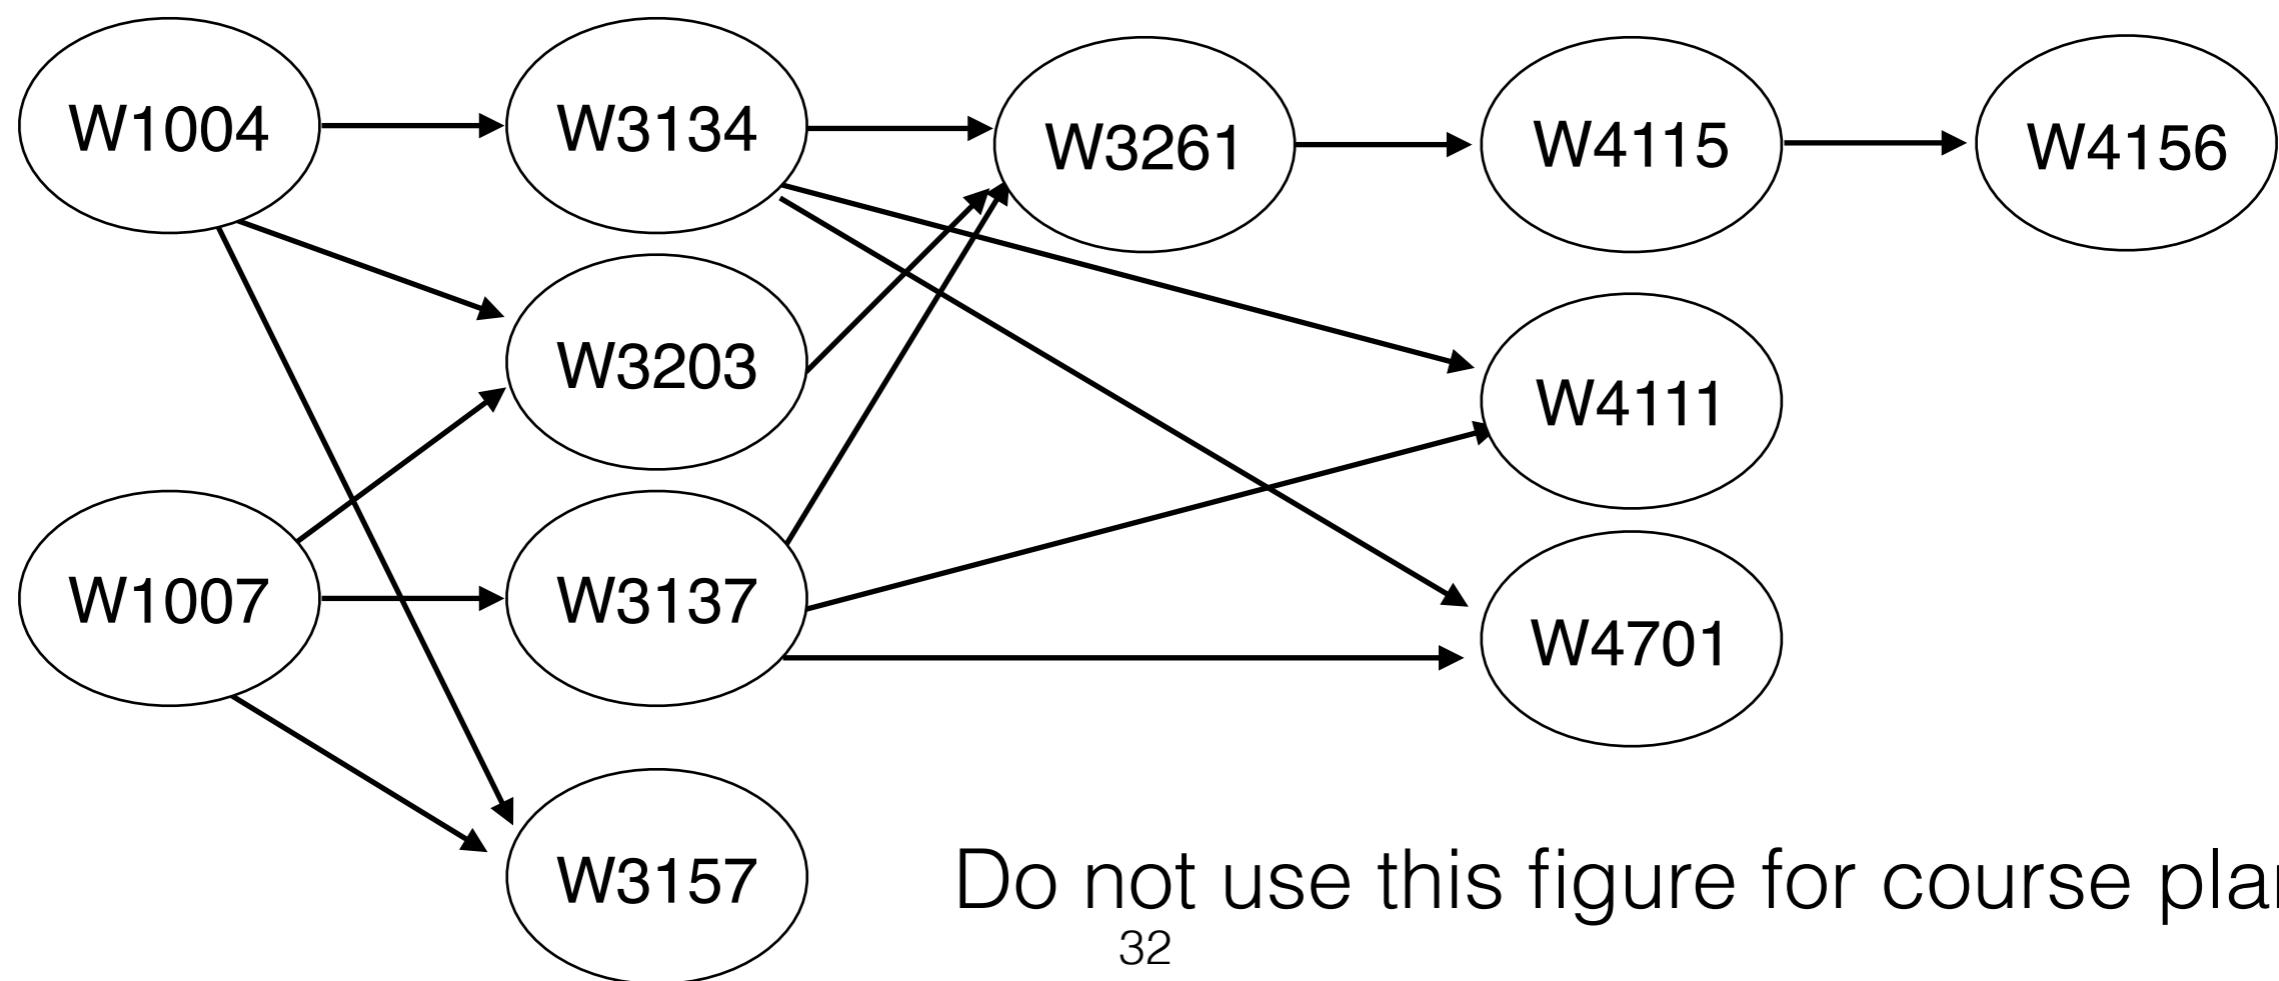
- Represent graph $G = (E, V)$, option 2: **Adjacency Lists**
 - For each vertex, keep a list of all adjacent vertices.



Space requirement: $\Theta(|V| + |E|)$

Topological Sort in DAGs

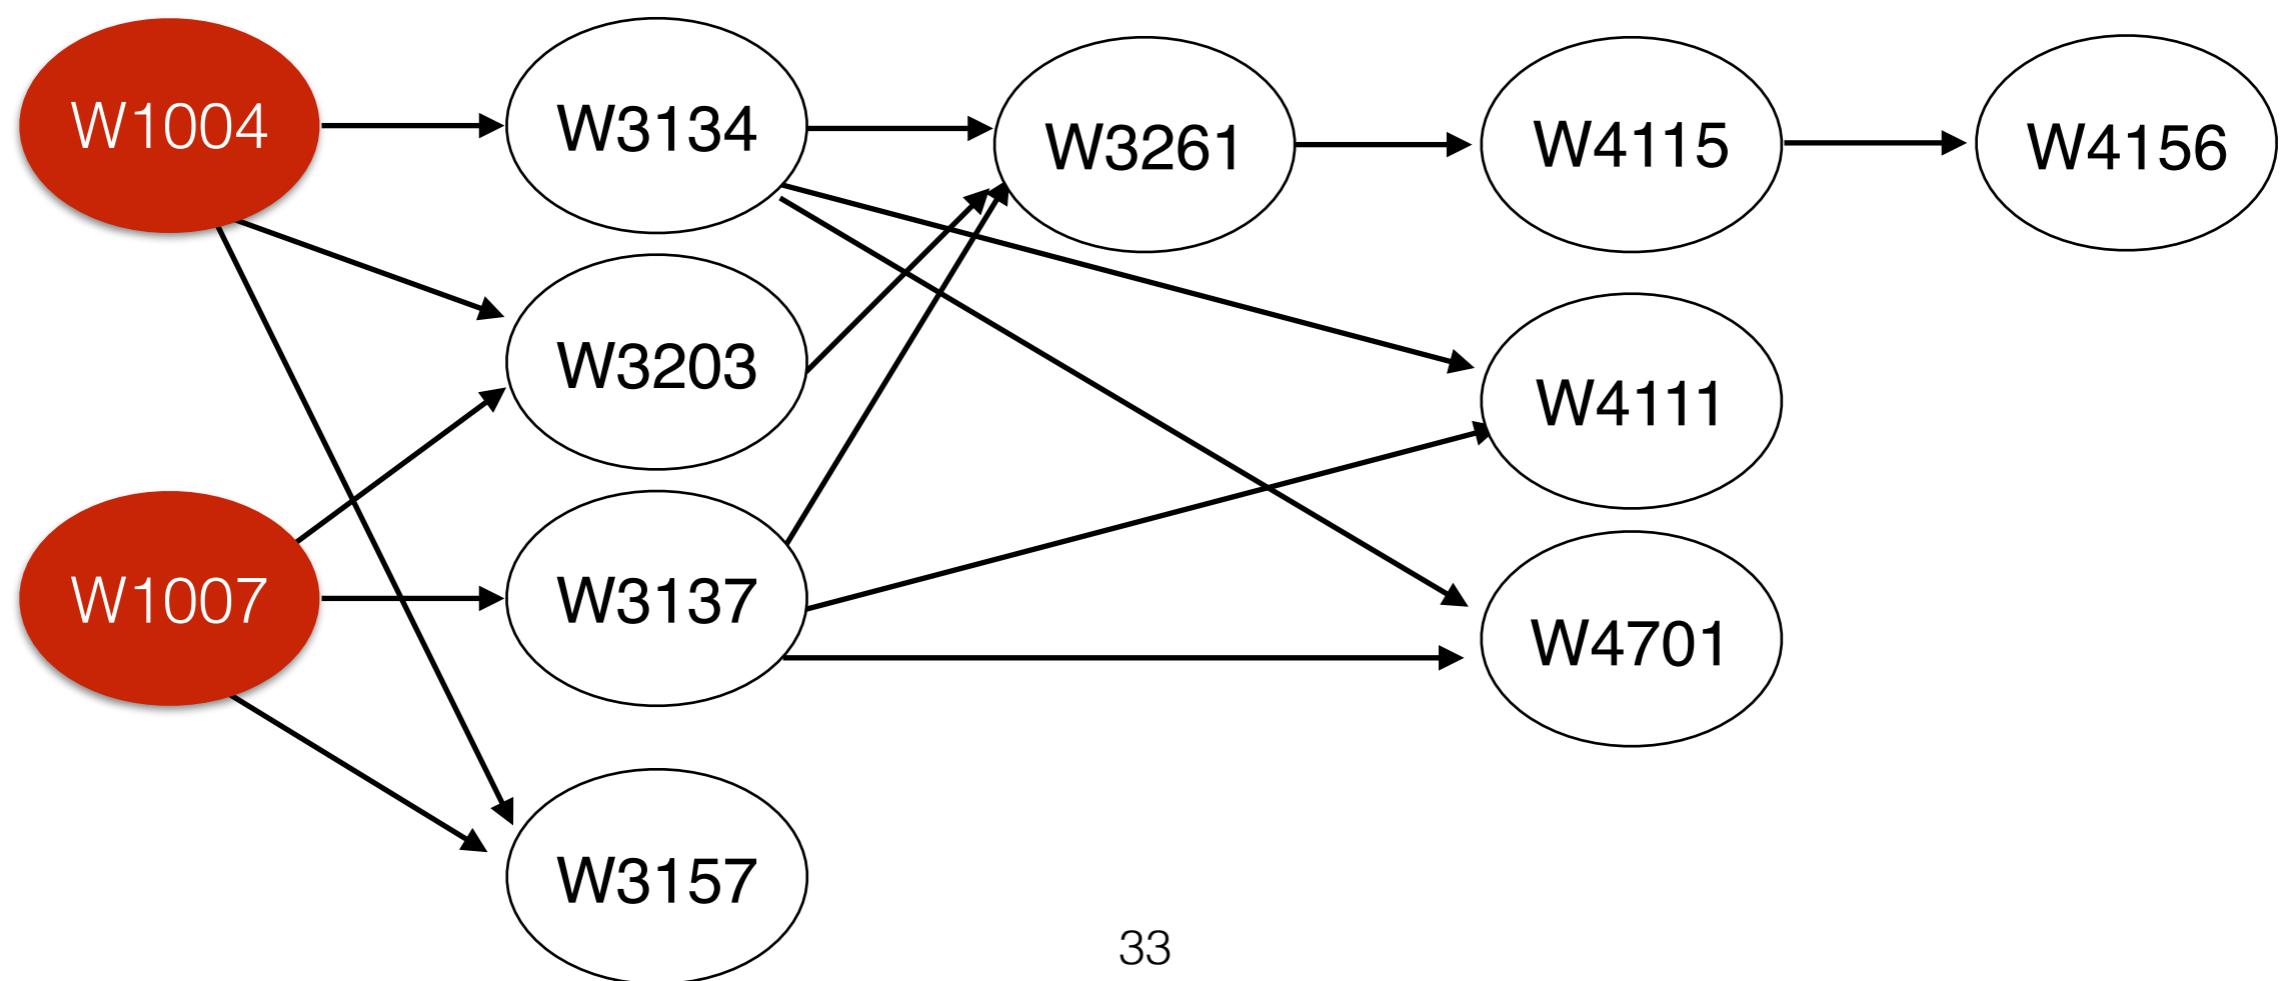
A **topological sort** of a DAG is an ordering of its vertices such that if there is a path from u to w , u appears before w in the ordering.



Topological Sort in DAGs

A topological sort of a DAG is an ordering of its vertices such that if there is a path from u to w , u appears before w in the ordering.

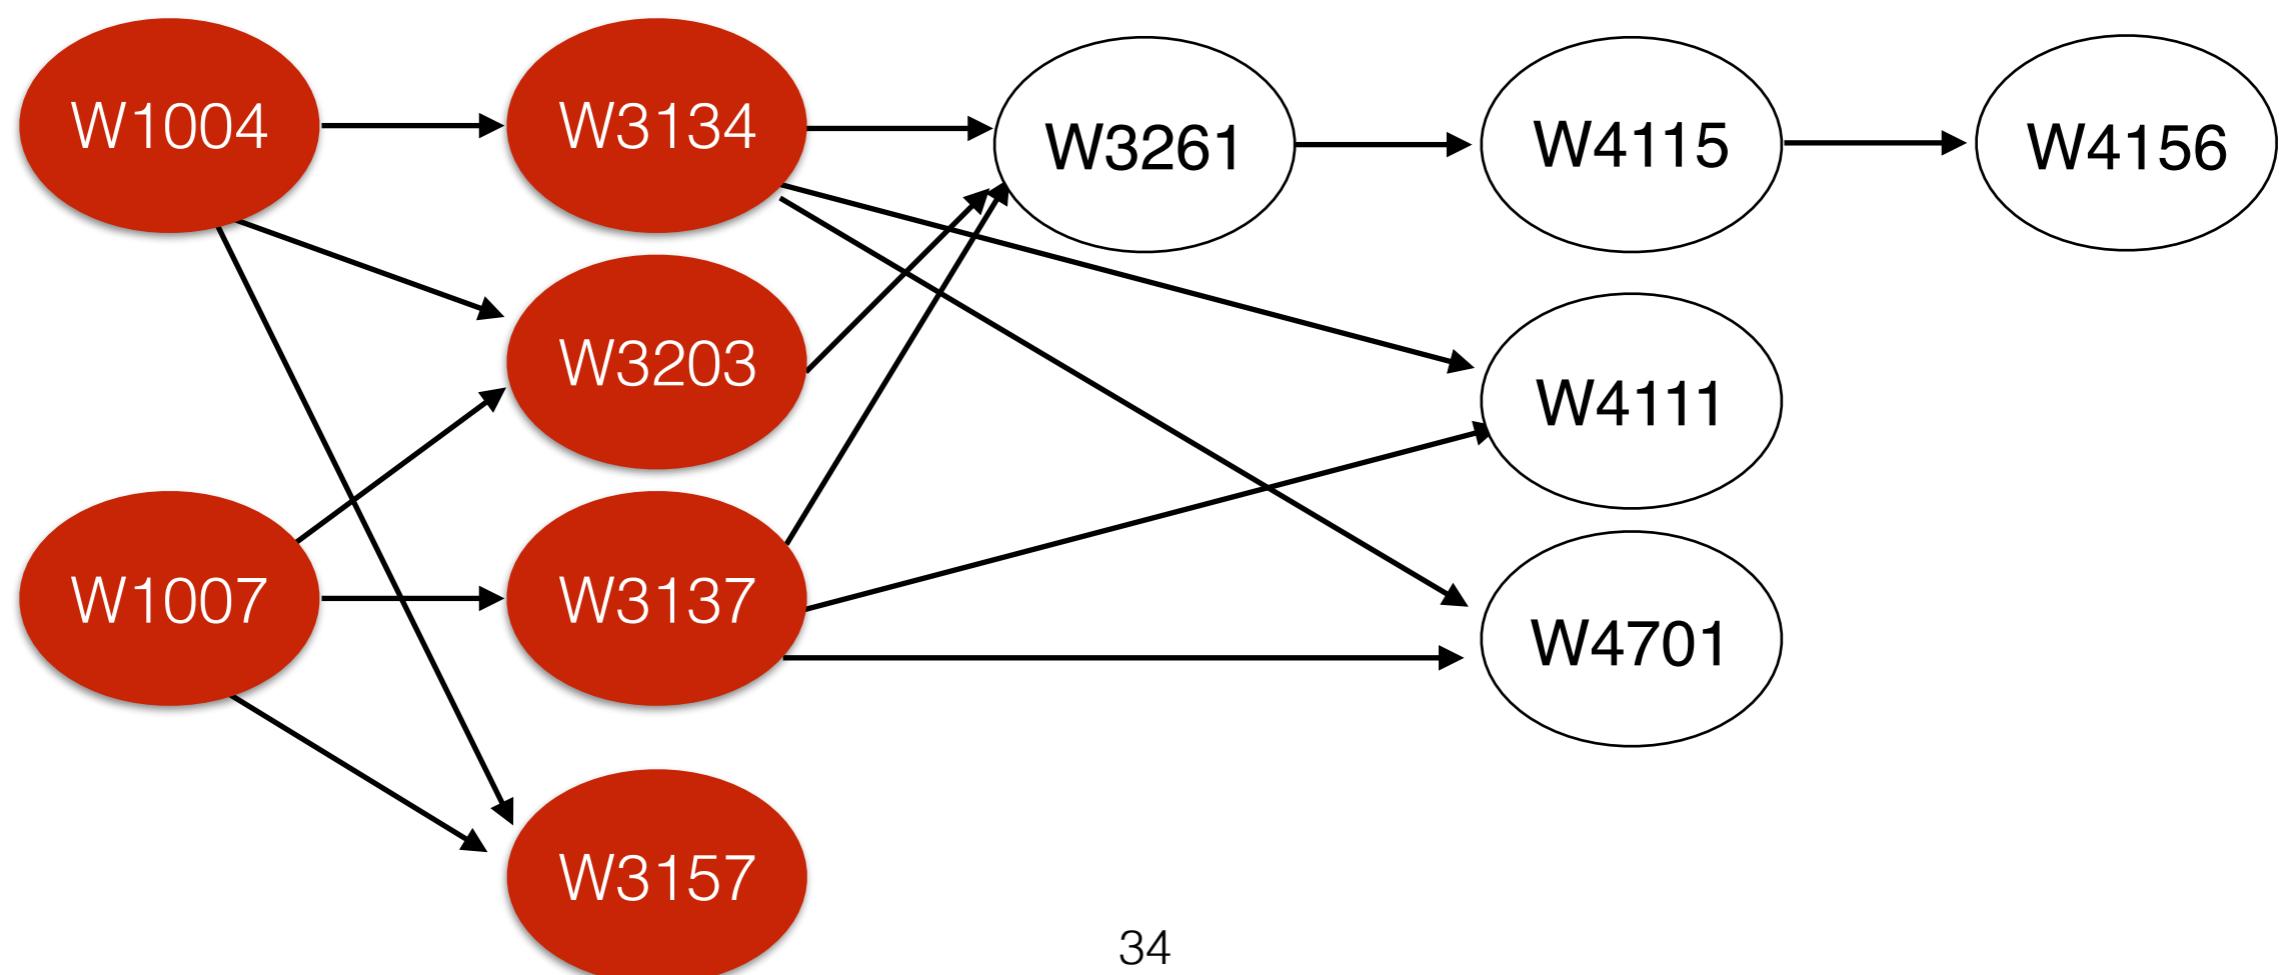
W1007 W1004



Topological Sort in DAGs

A topological sort of a DAG is an ordering of its vertices such that if there is a path from u to w , u appears before w in the ordering.

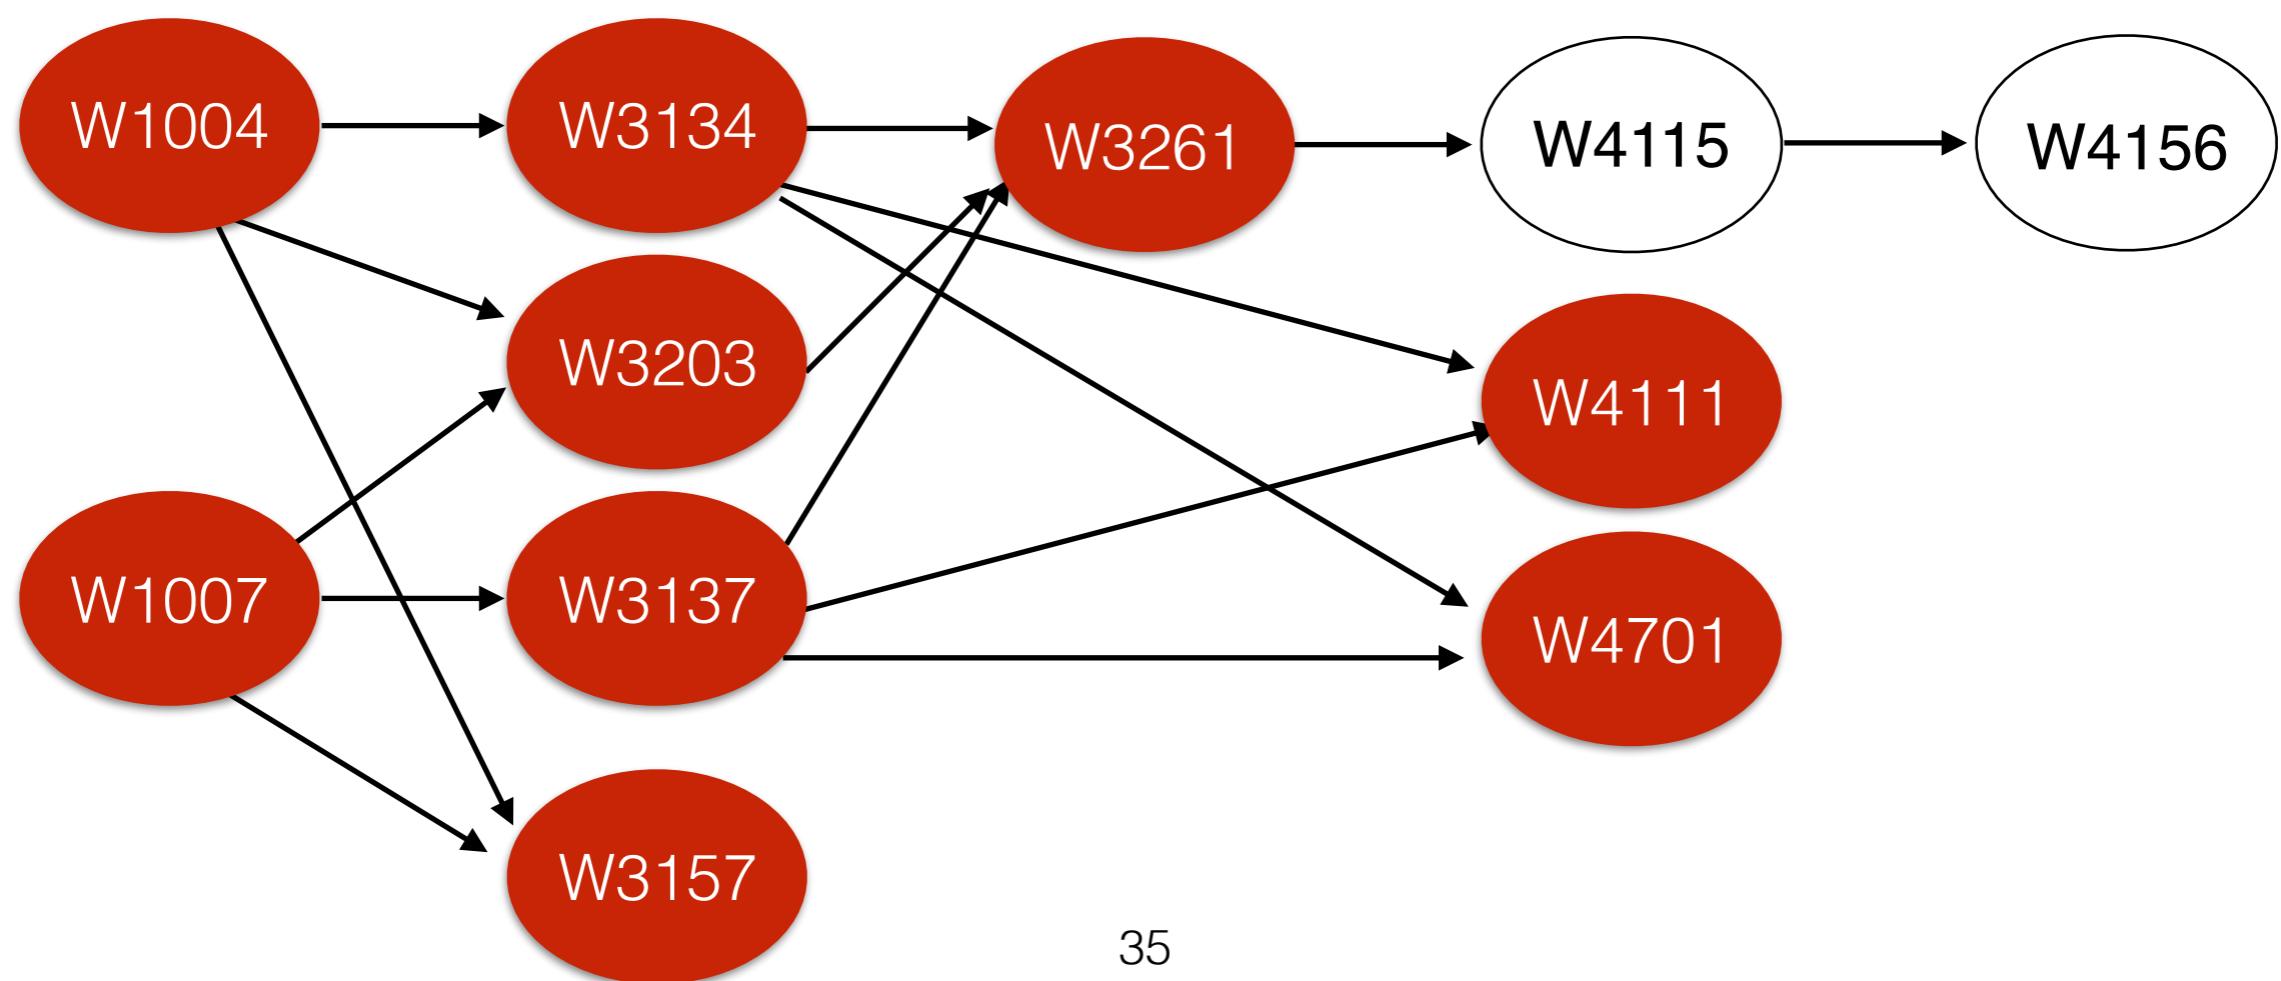
W1007 W1004 W3134 W3203 W3137 W3157



Topological Sort in DAGs

A topological sort of a DAG is an ordering of its vertices such that if there is a path from u to w , u appears before w in the ordering.

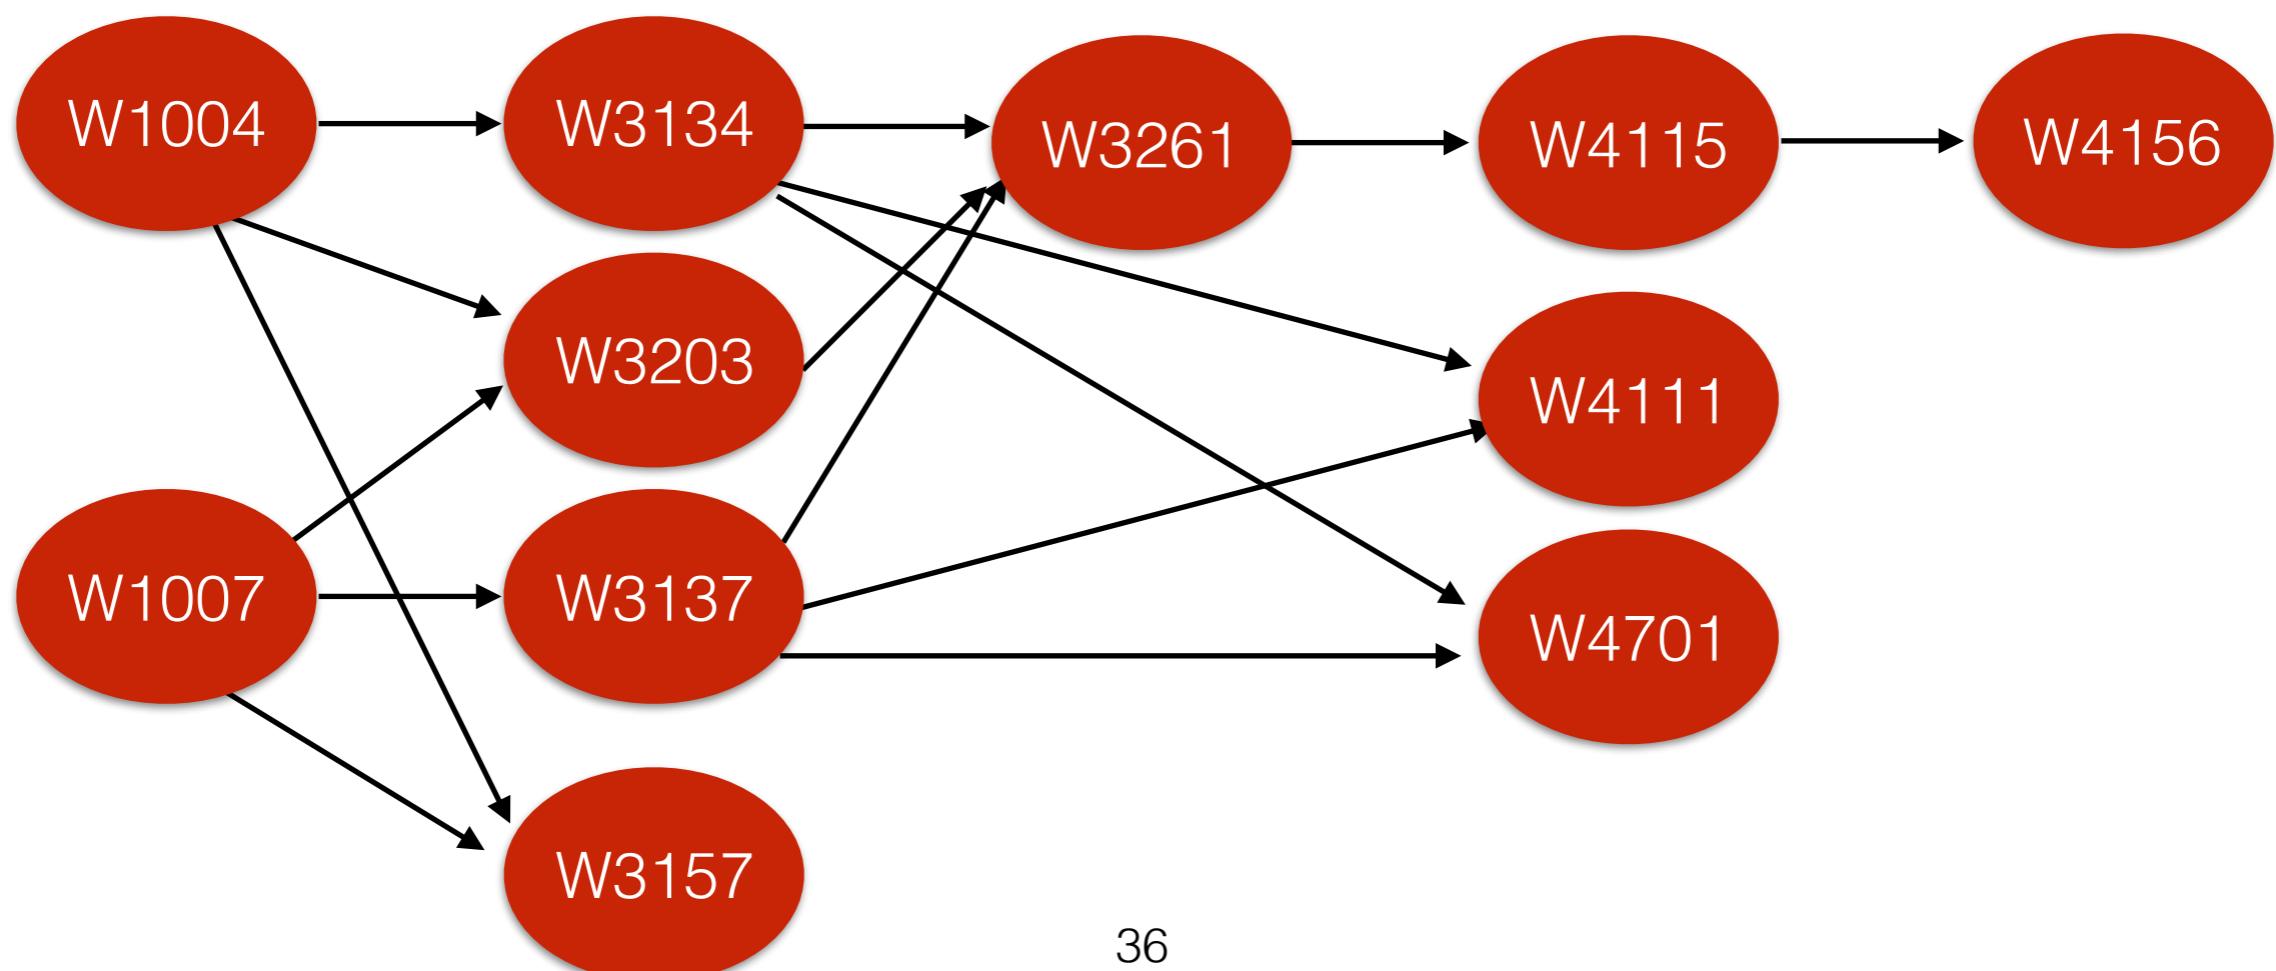
W1007 W1004 W3134 W3203 W3137 W3157 W3261 W4111 W4701



Topological Sort in DAGs

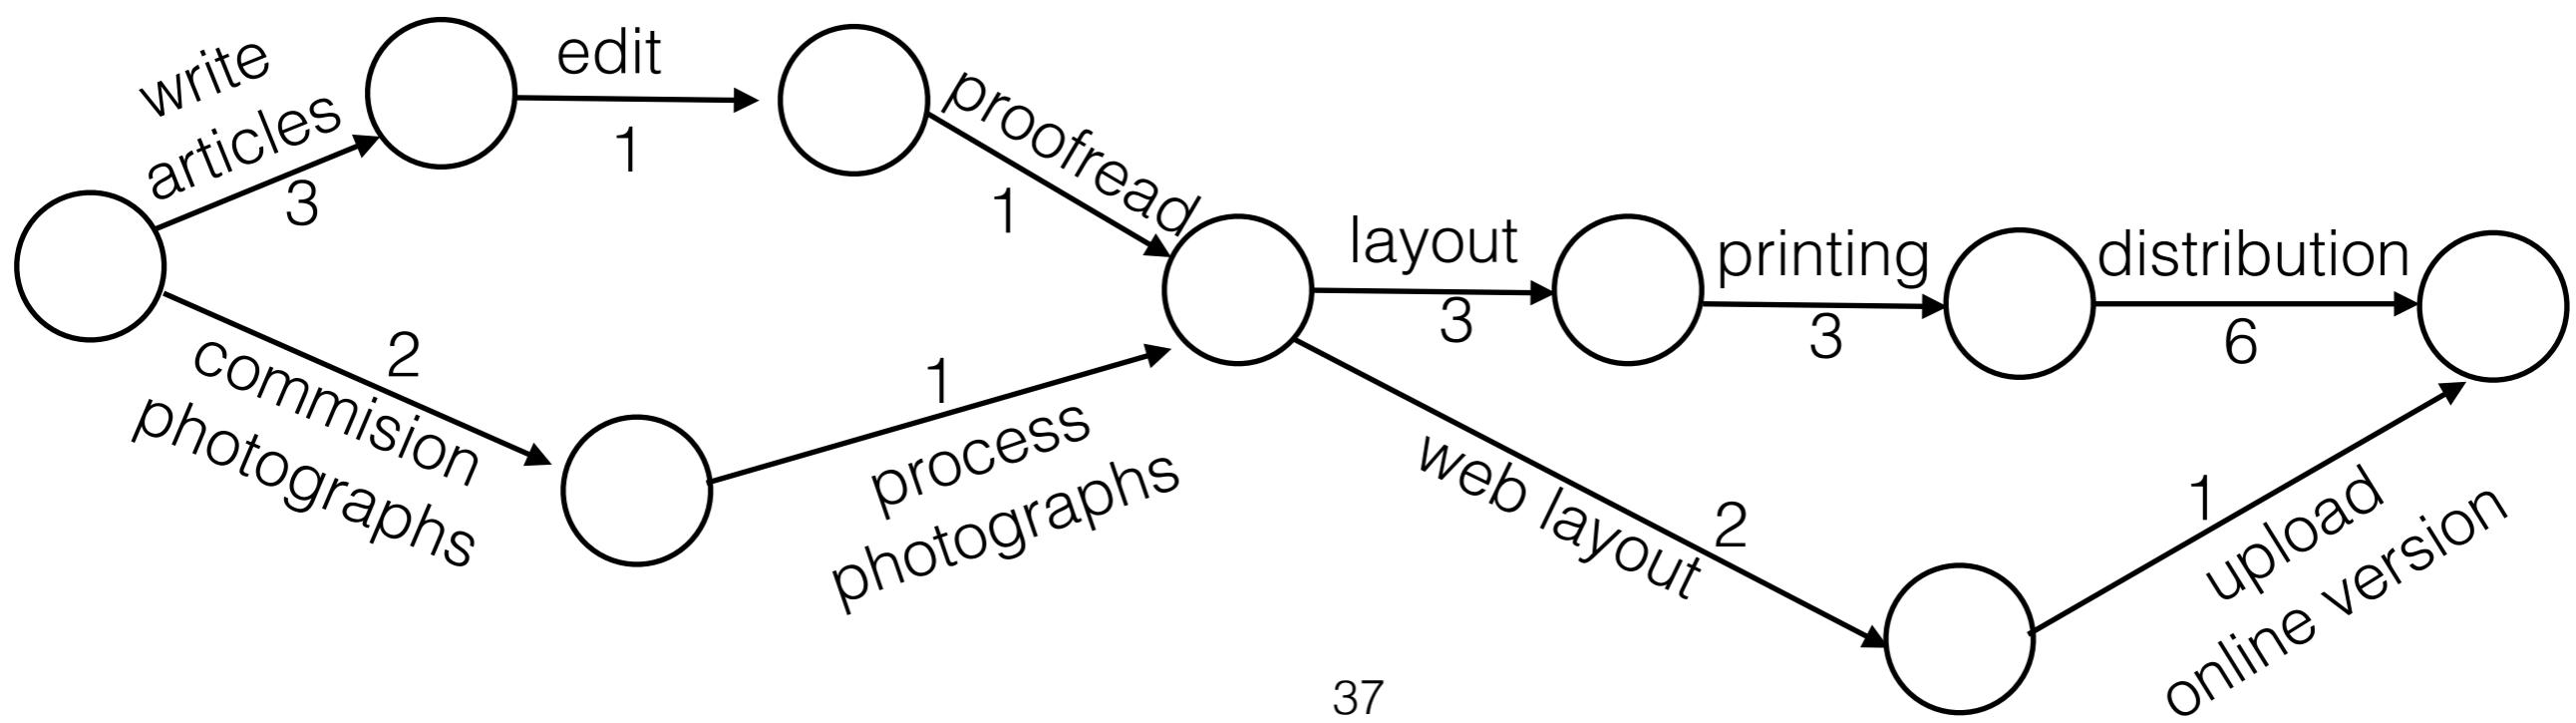
A topological sort of a DAG is an ordering of its vertices such that if there is a path from u to w , u appears before w in the ordering.

W1007 W1004 W3134 W3203 W3137 W3157 W3261 W4111 W4701 W4115 W4156



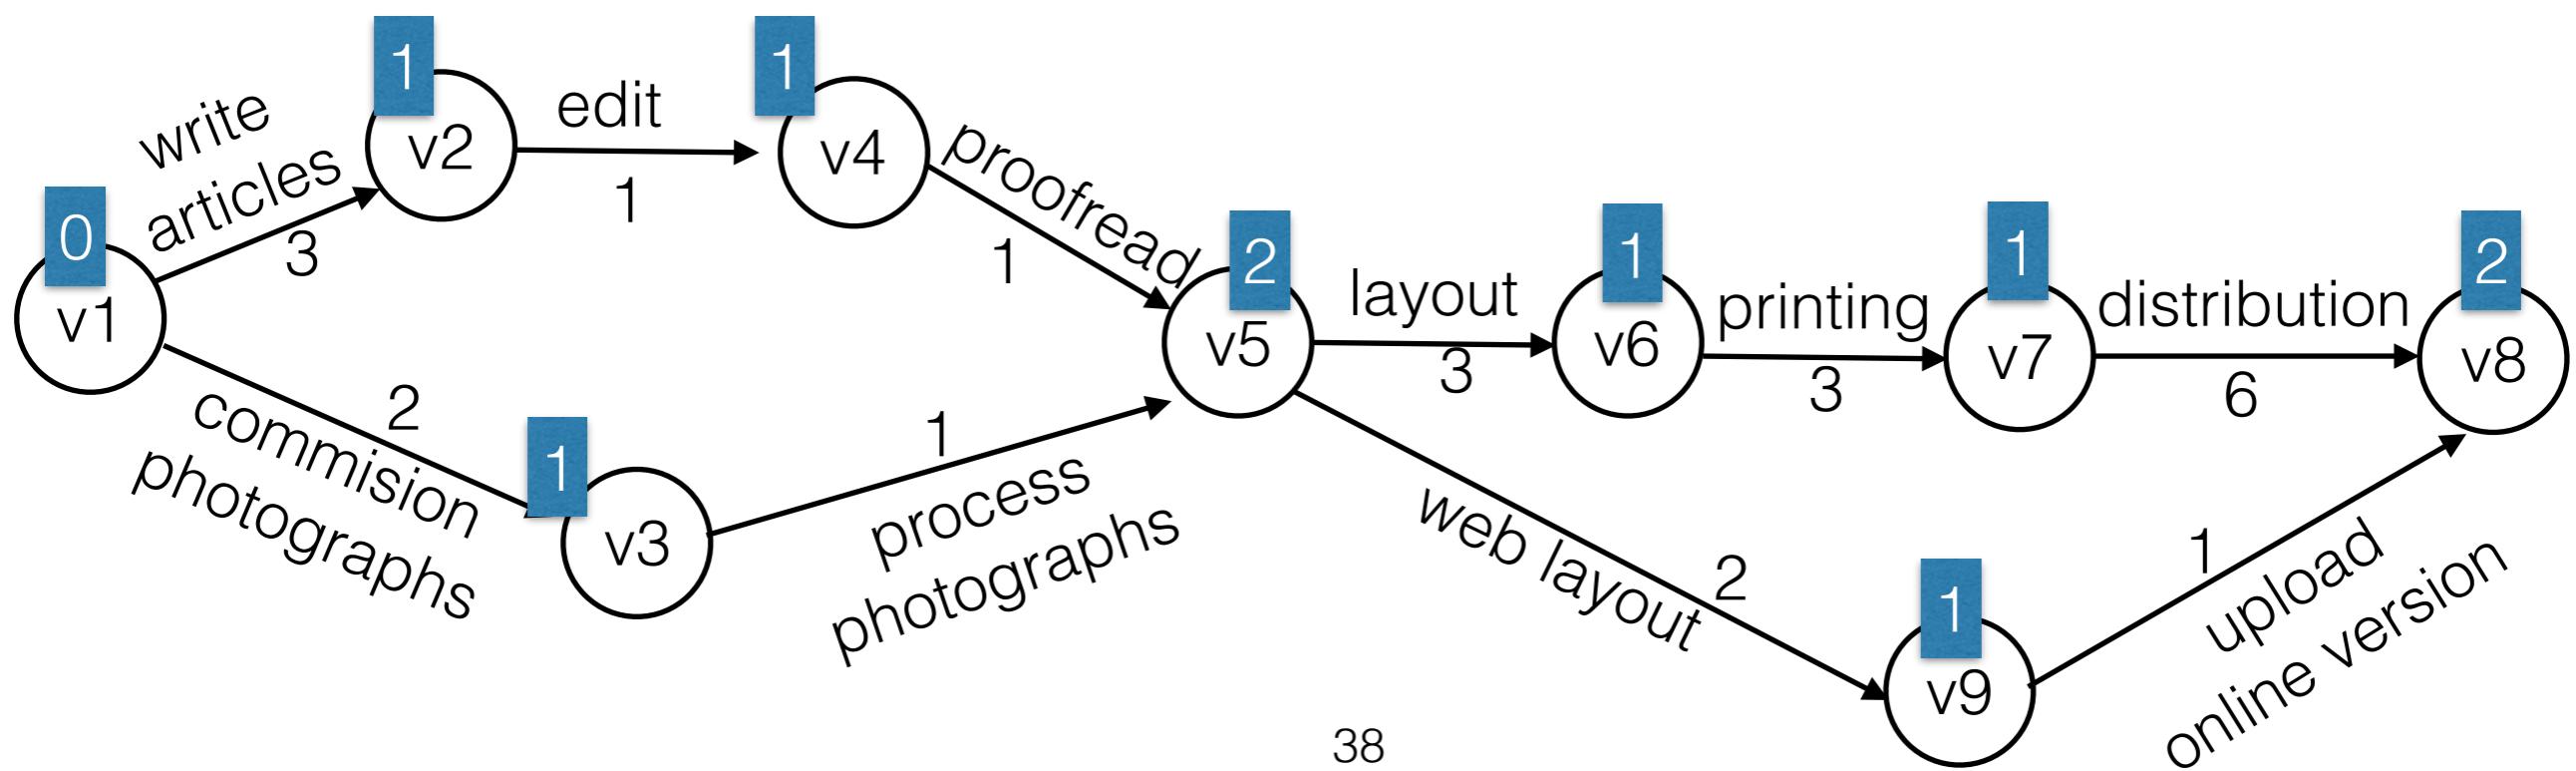
Application: Critical Path Analysis

- An **Event-Node Graph** is a DAG in which
 - Edges represent tasks, weights represents the time it takes to complete the task.
 - Vertices represent the event of completing a set of tasks.



Computing Topological Sort

- First annotate each vertex with the number of incoming edges (the **indegree**).

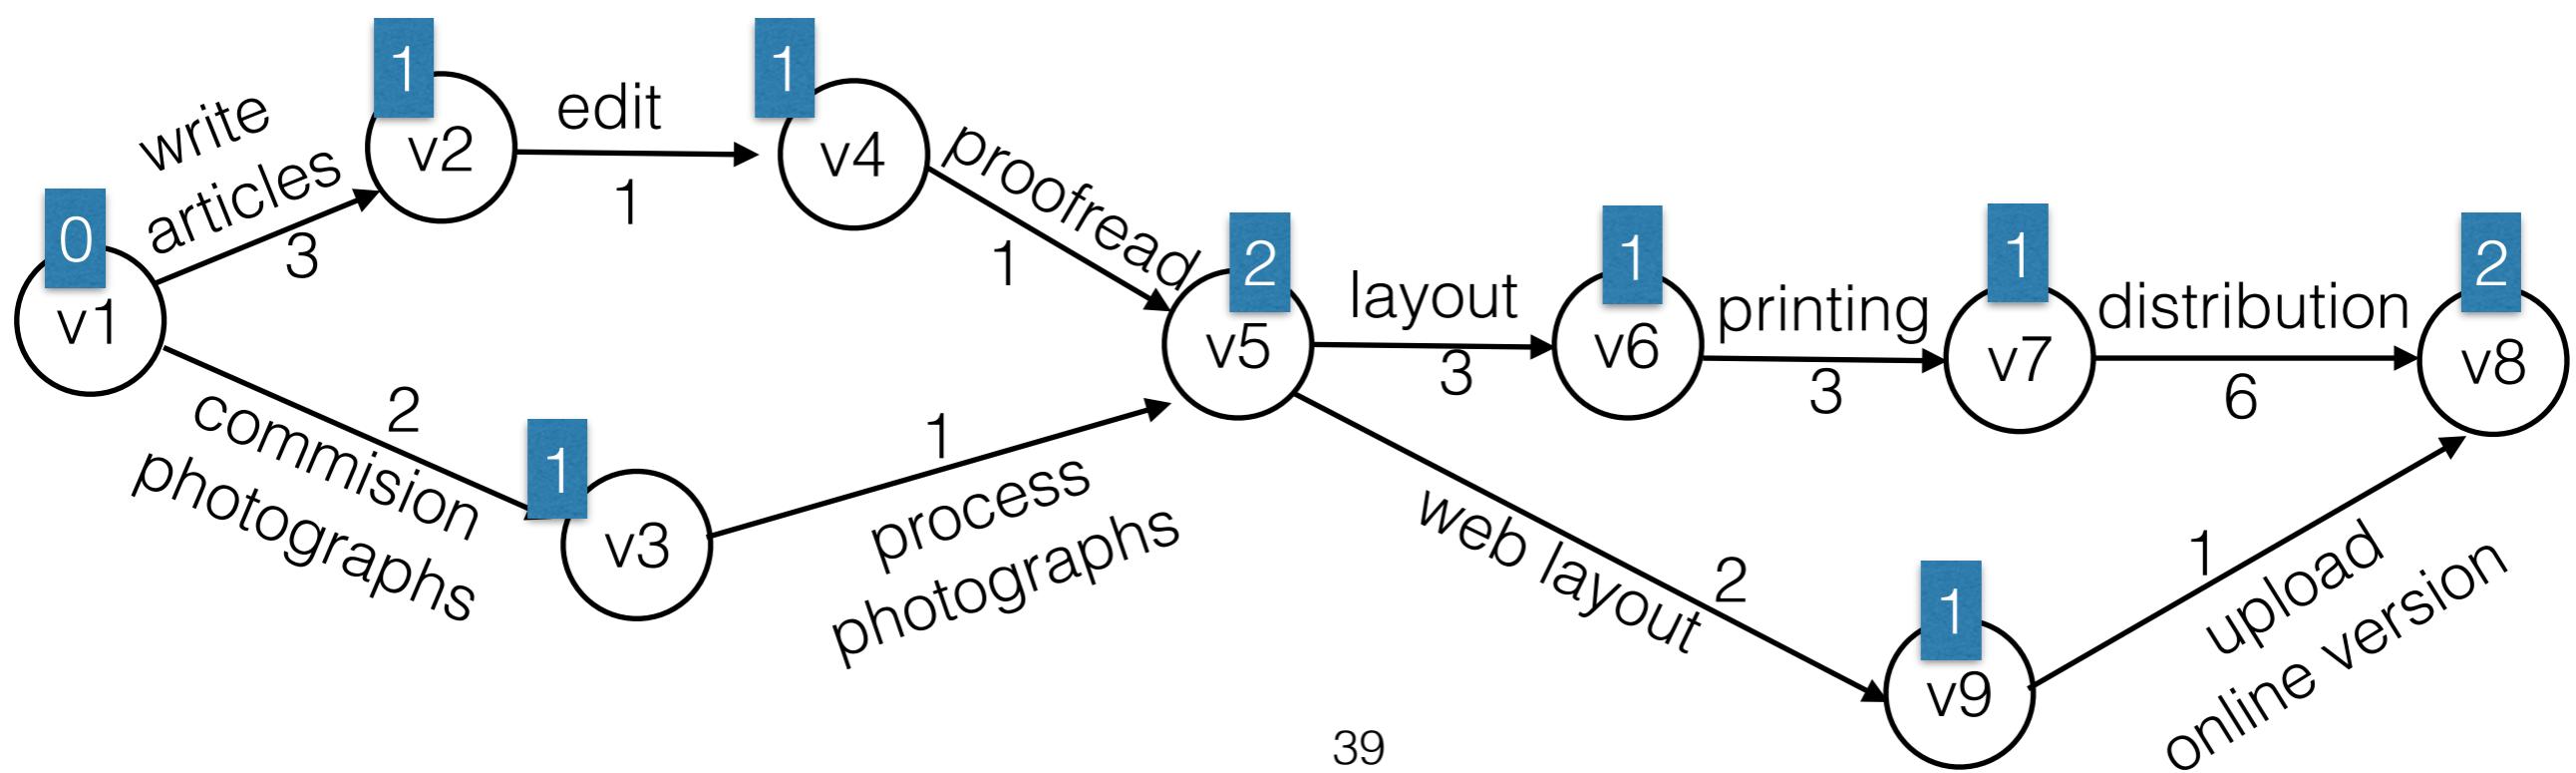


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v1

Output:

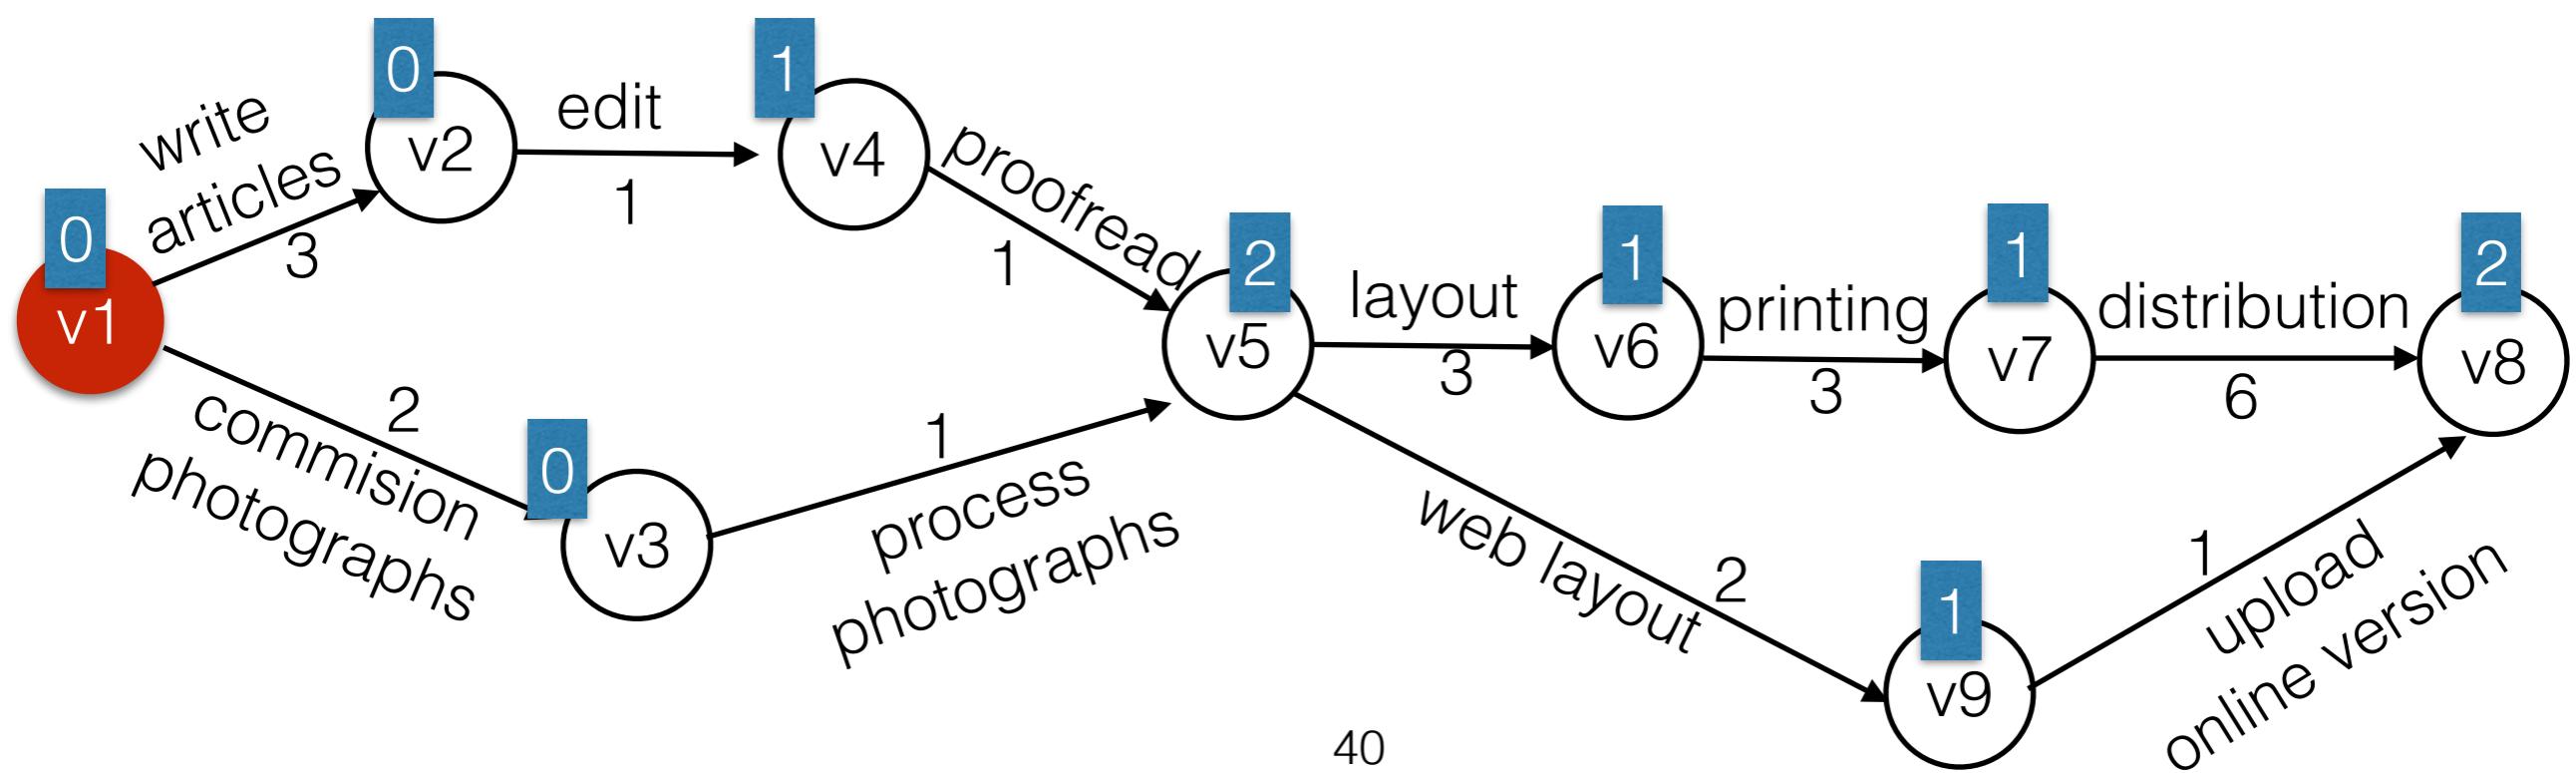


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v2 v3

Output: v1

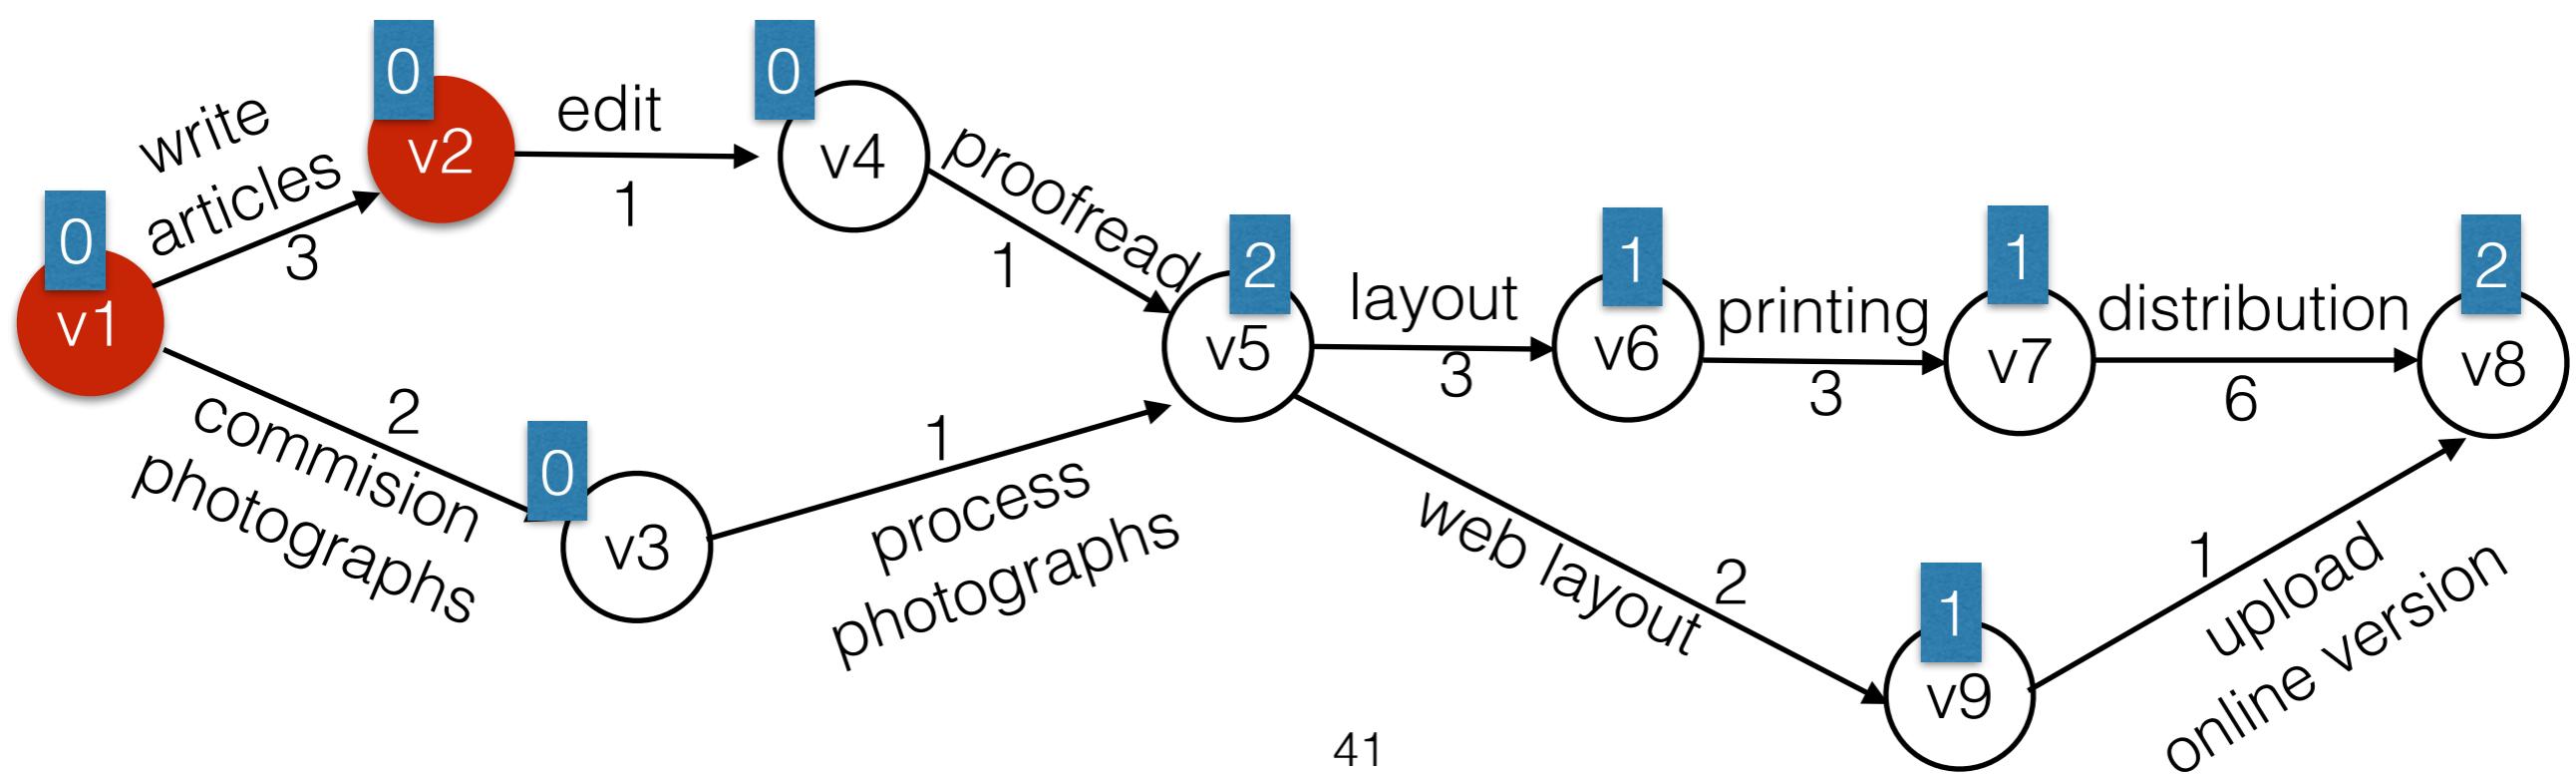


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v3 v4

Output: v1 v2

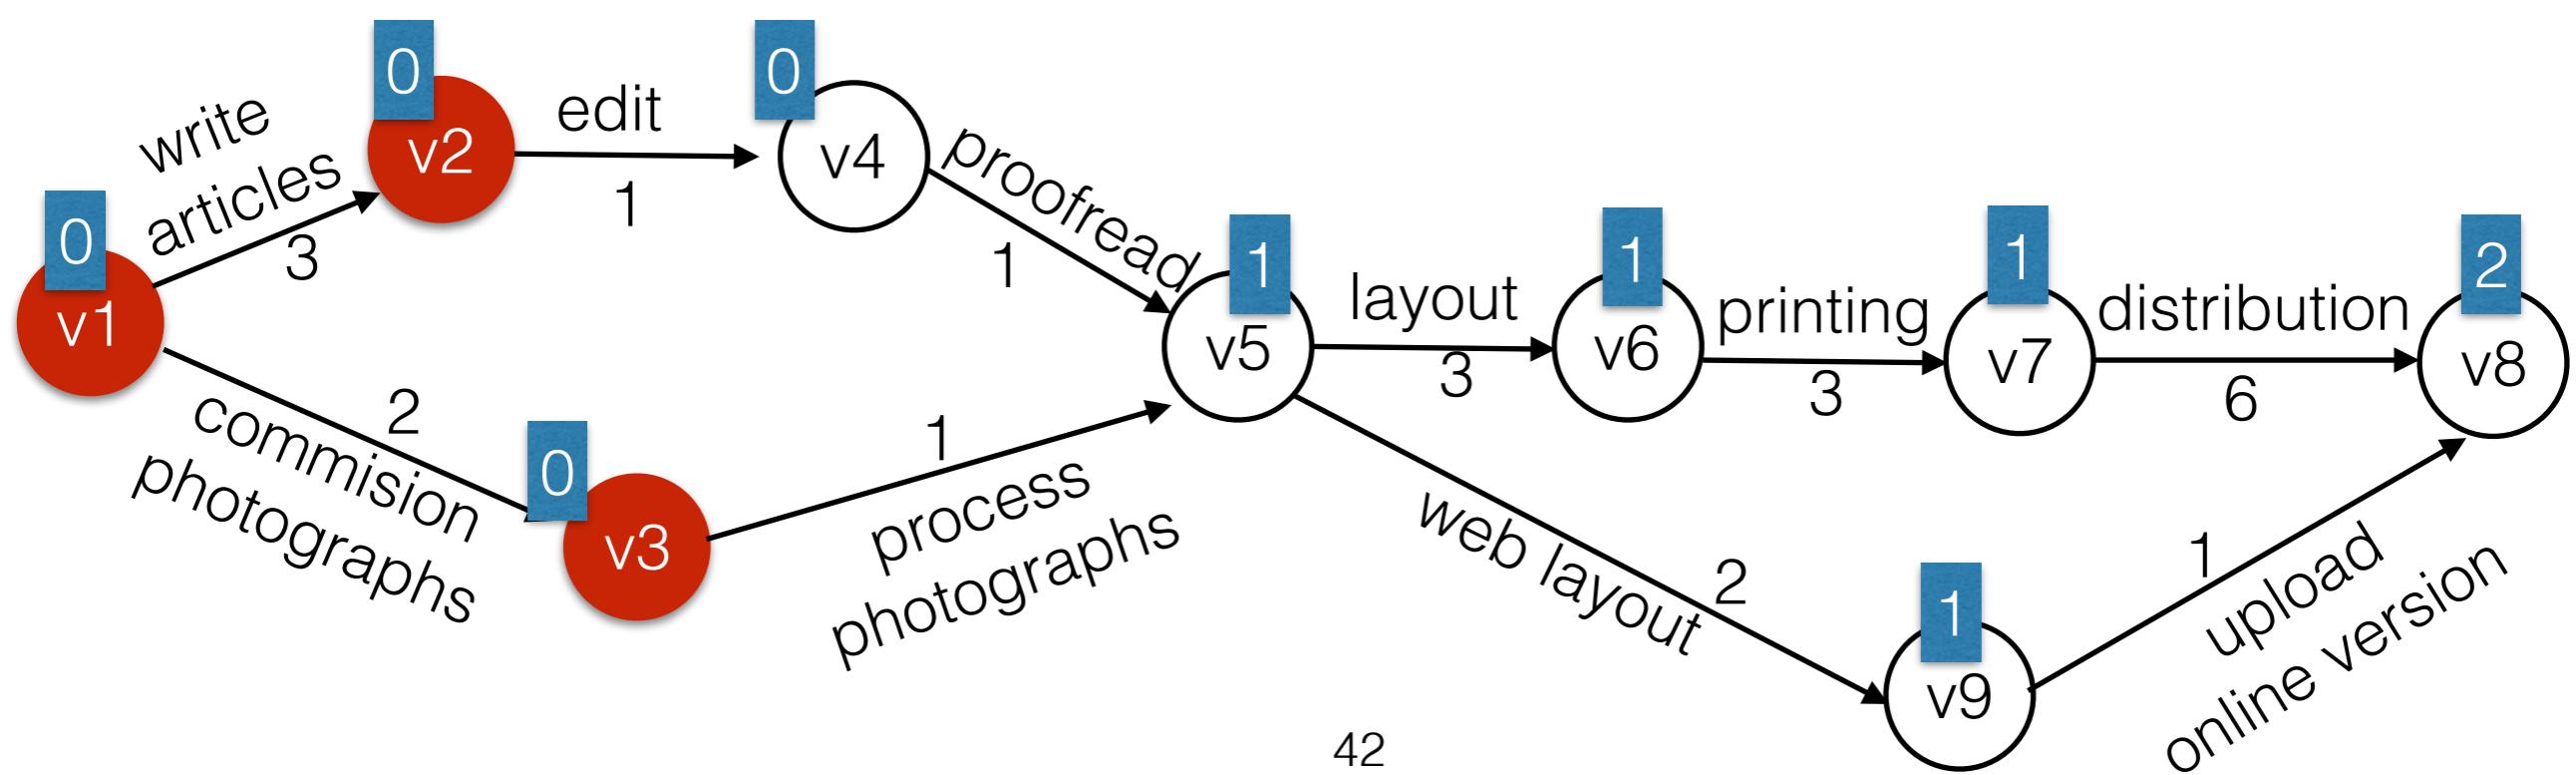


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v4

Output: v1 v2 v3

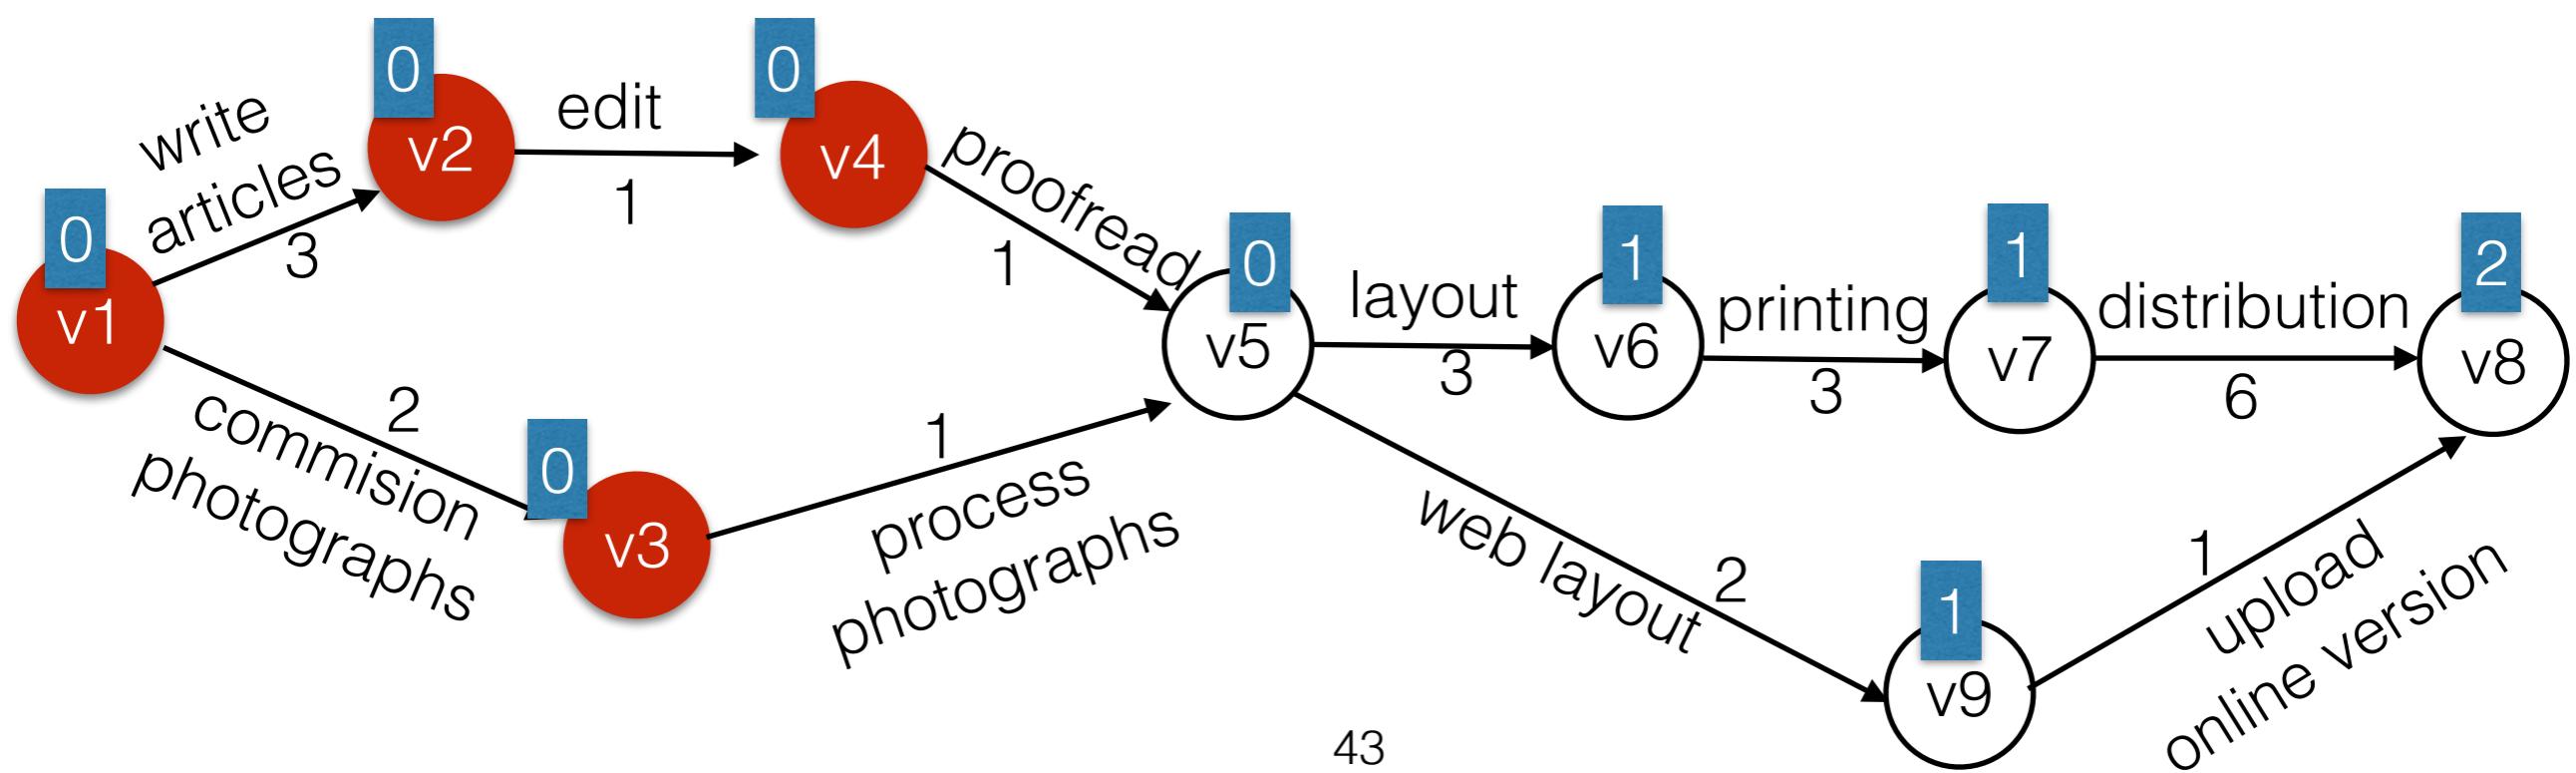


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v5

Output: v1 v2 v3 v4

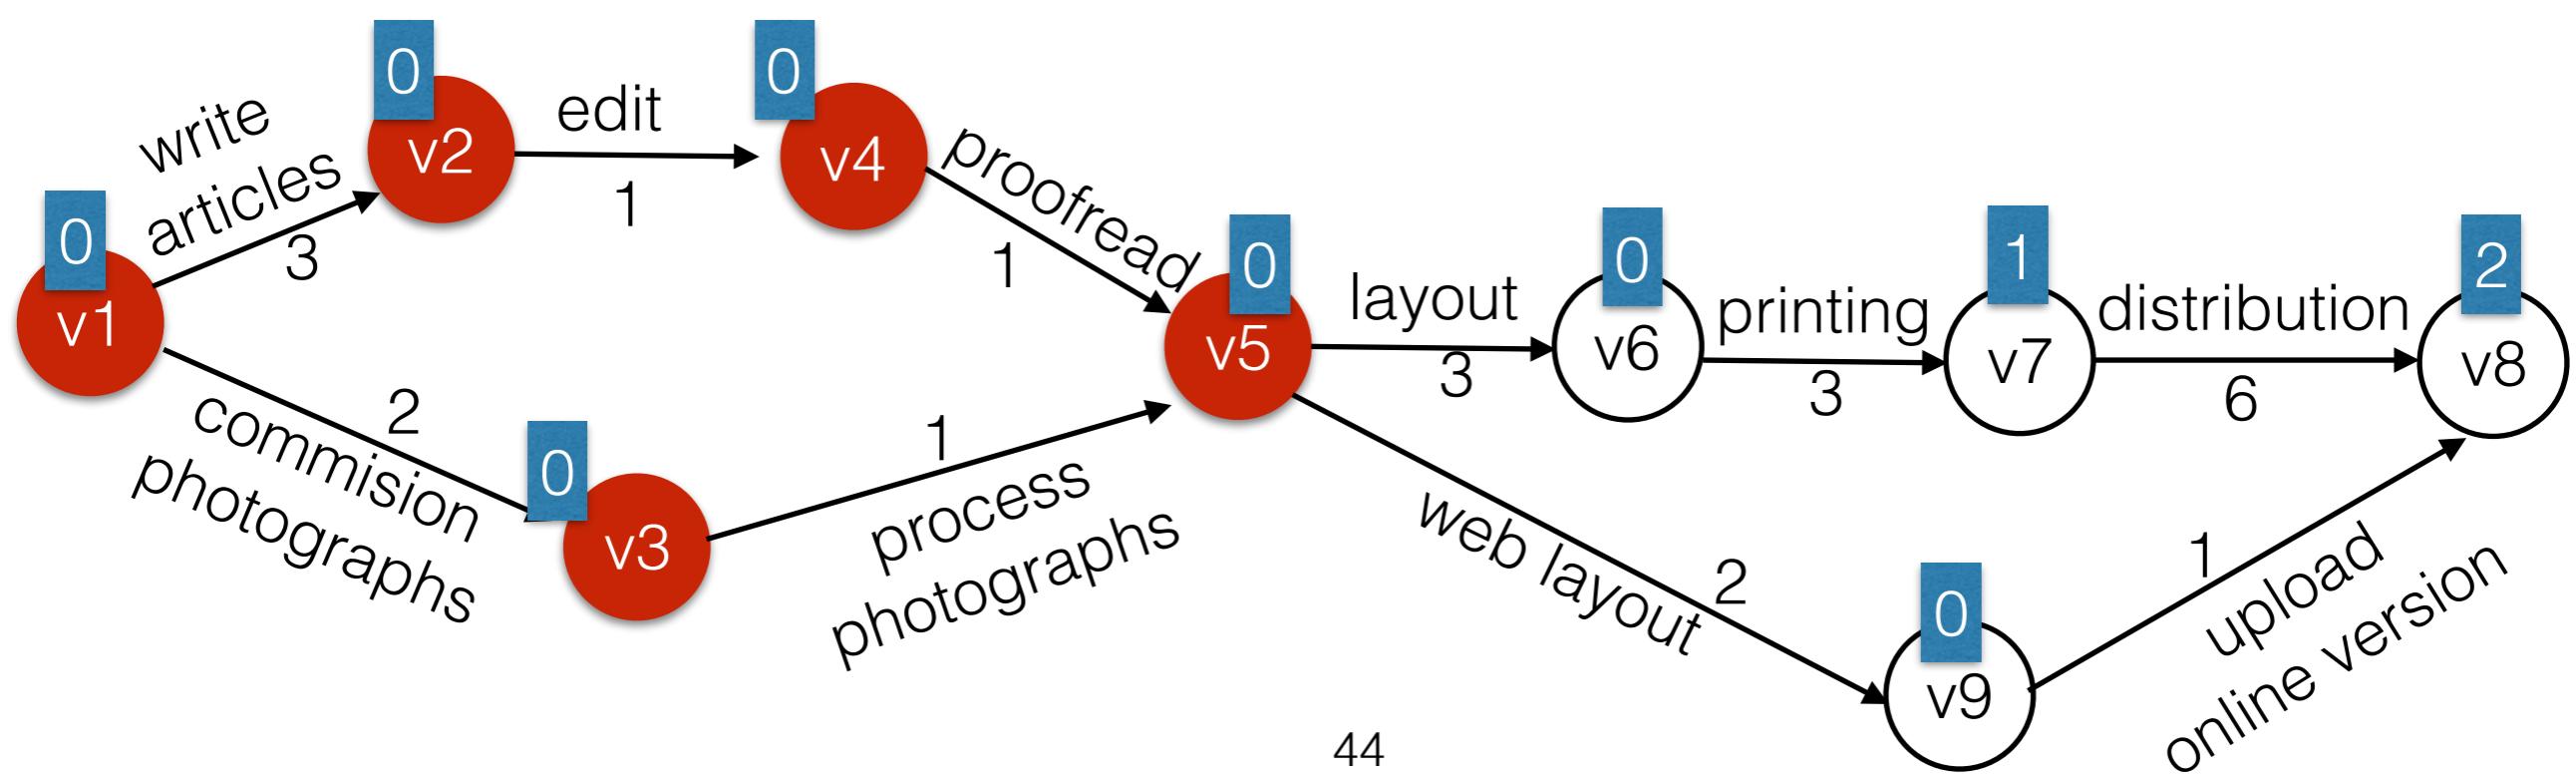


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v6 v9

Output: v1 v2 v3 v4 v5

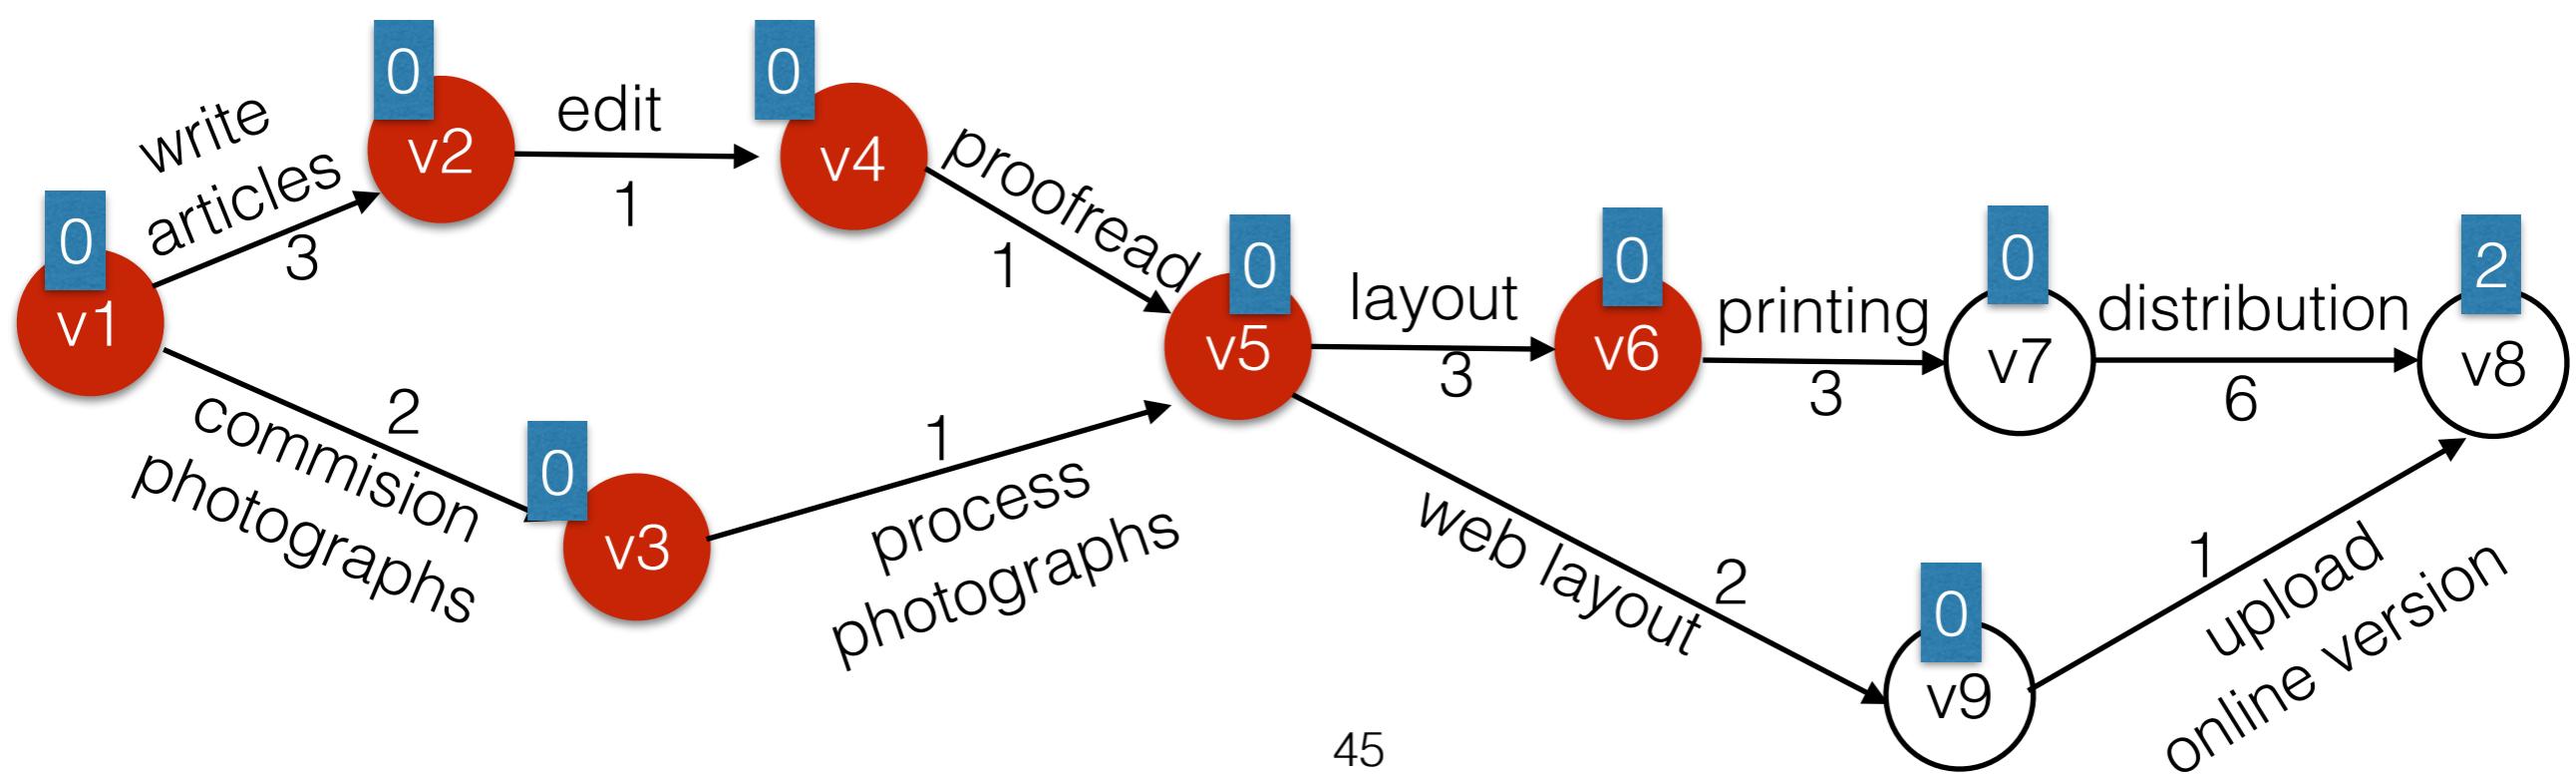


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v9 v7

Output: v1 v2 v3 v4 v5 v6

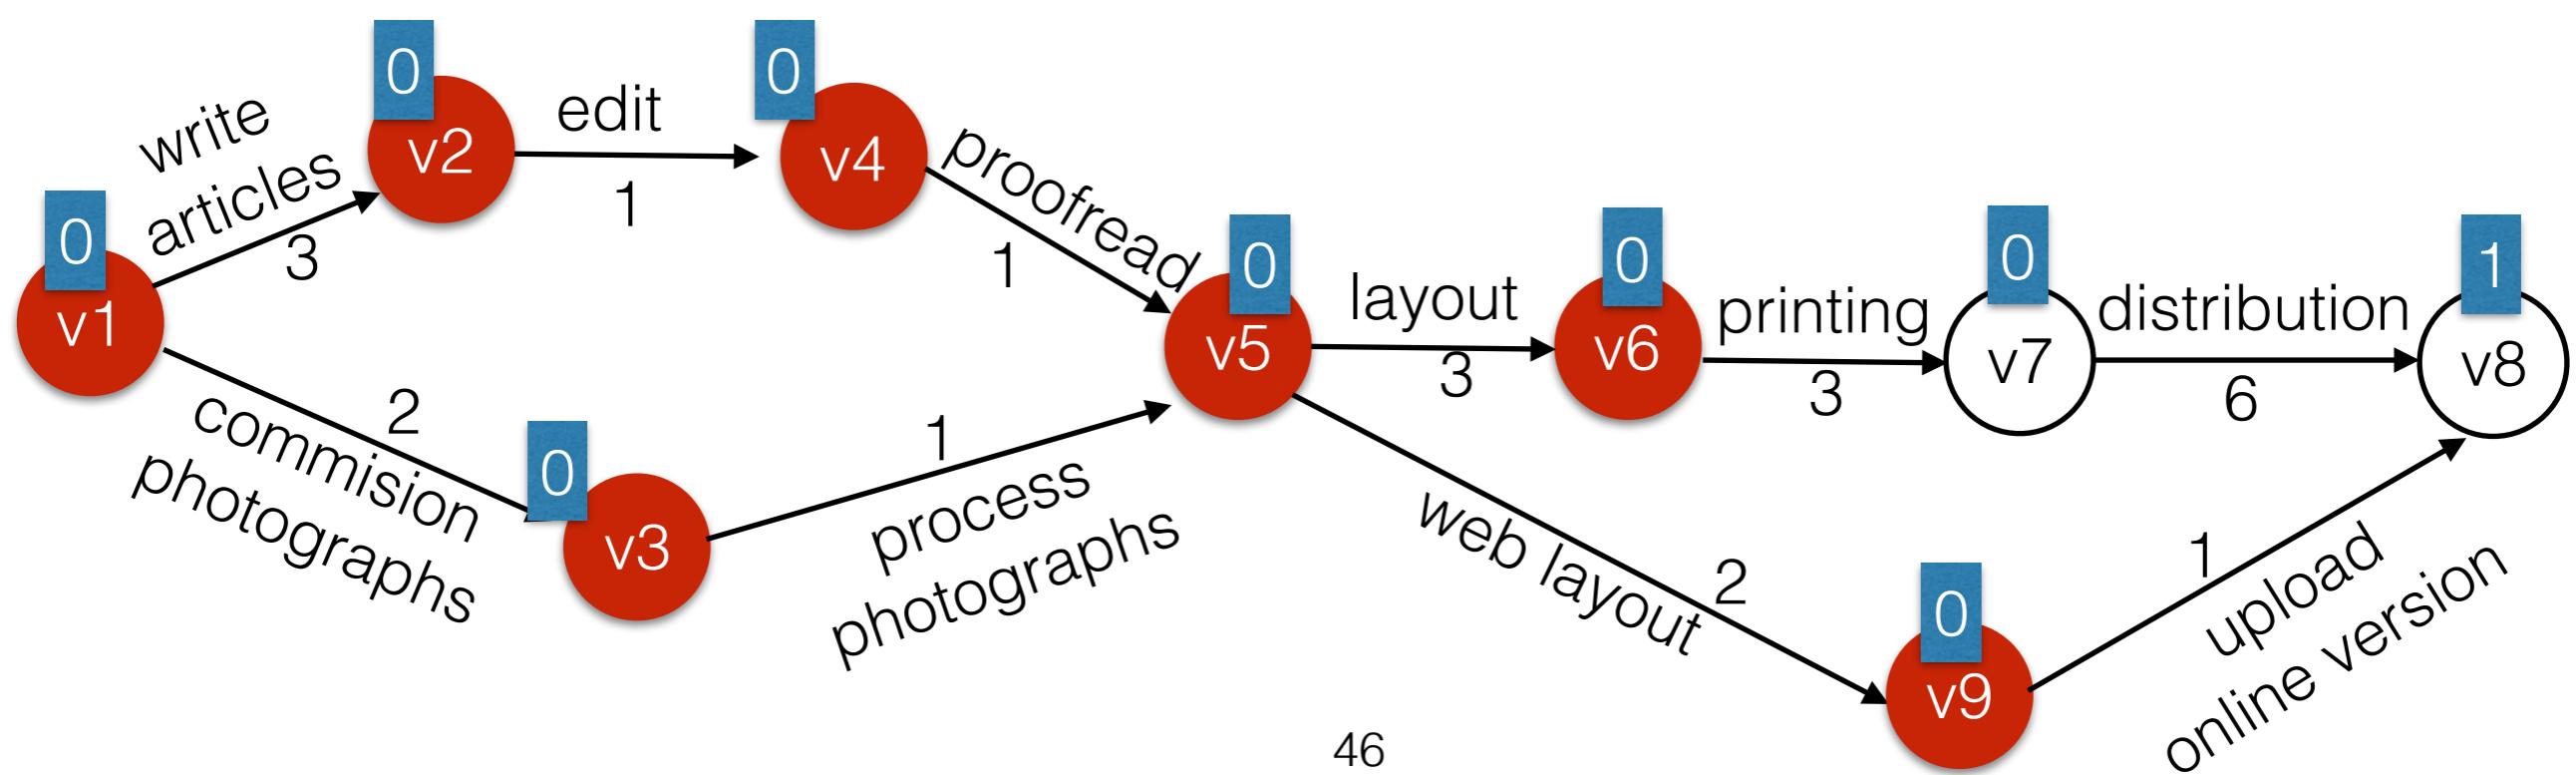


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v7

Output: v1 v2 v3 v4 v5 v6 v9

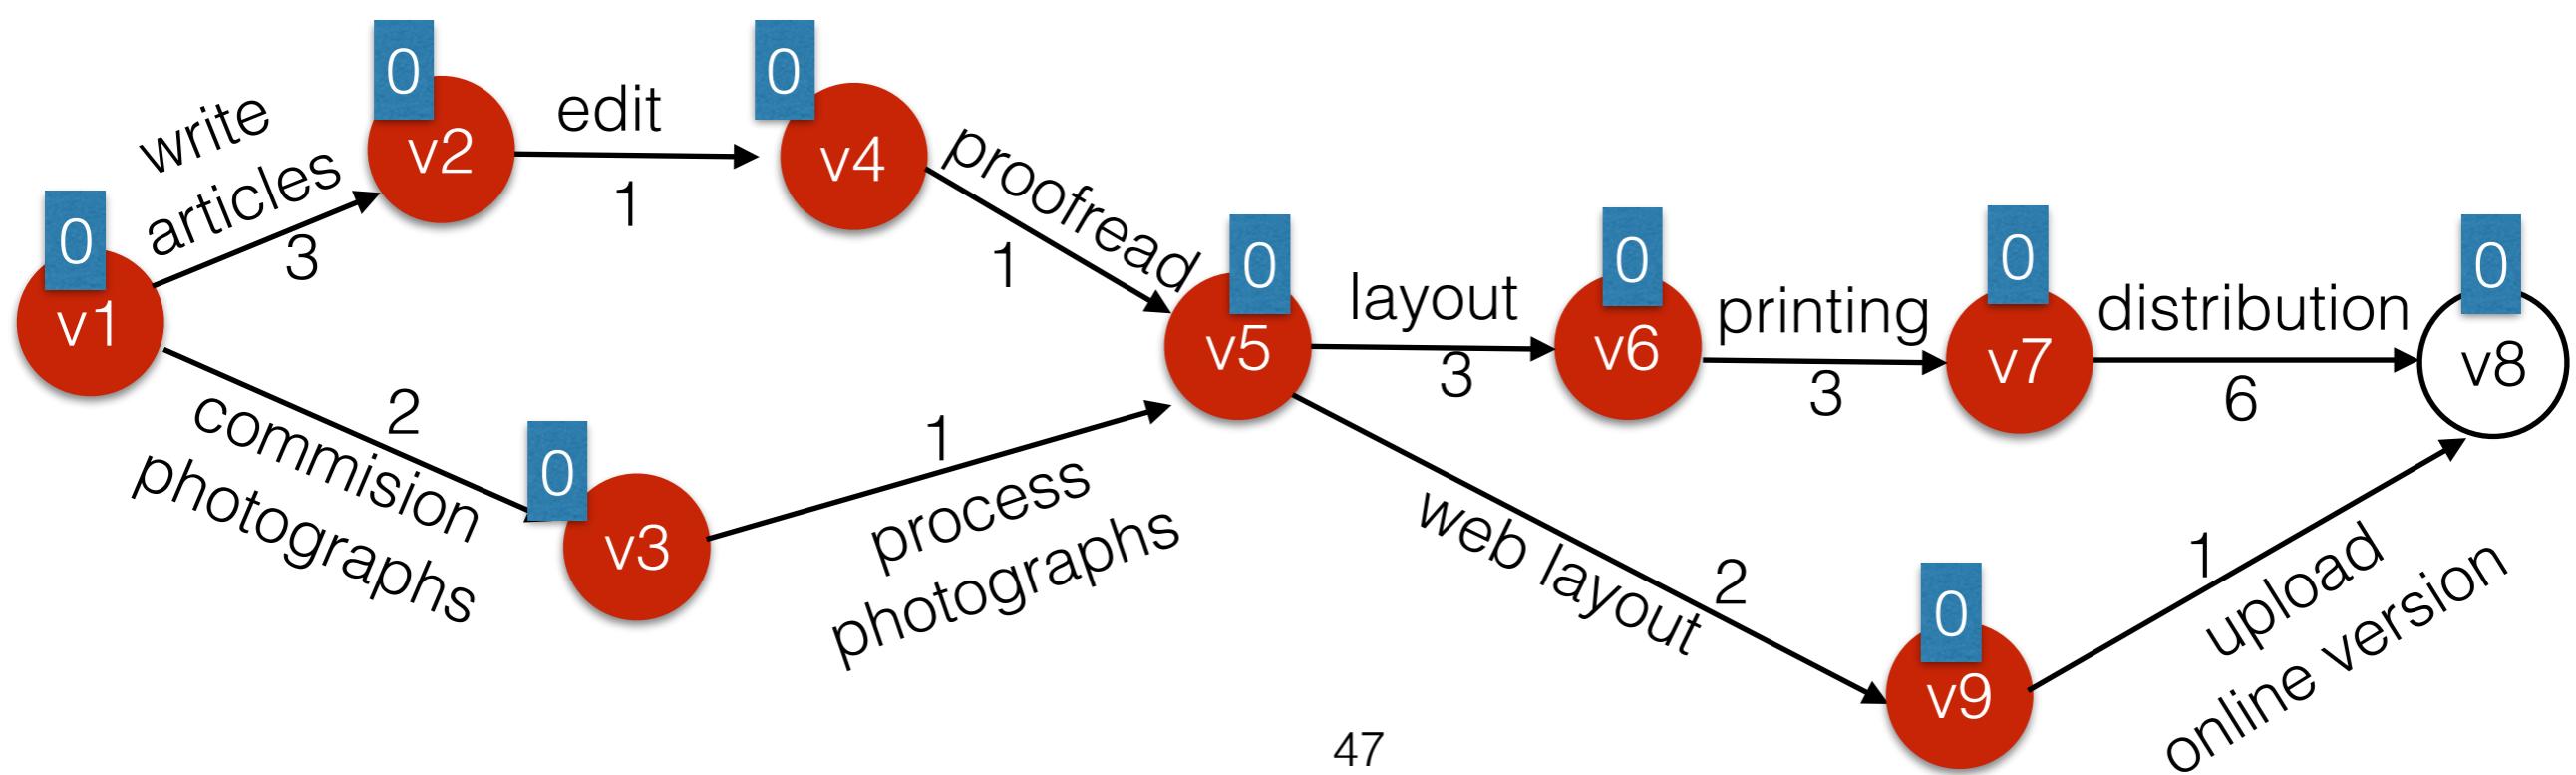


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v8

Output: v1 v2 v3 v4 v5 v6 v9 v7

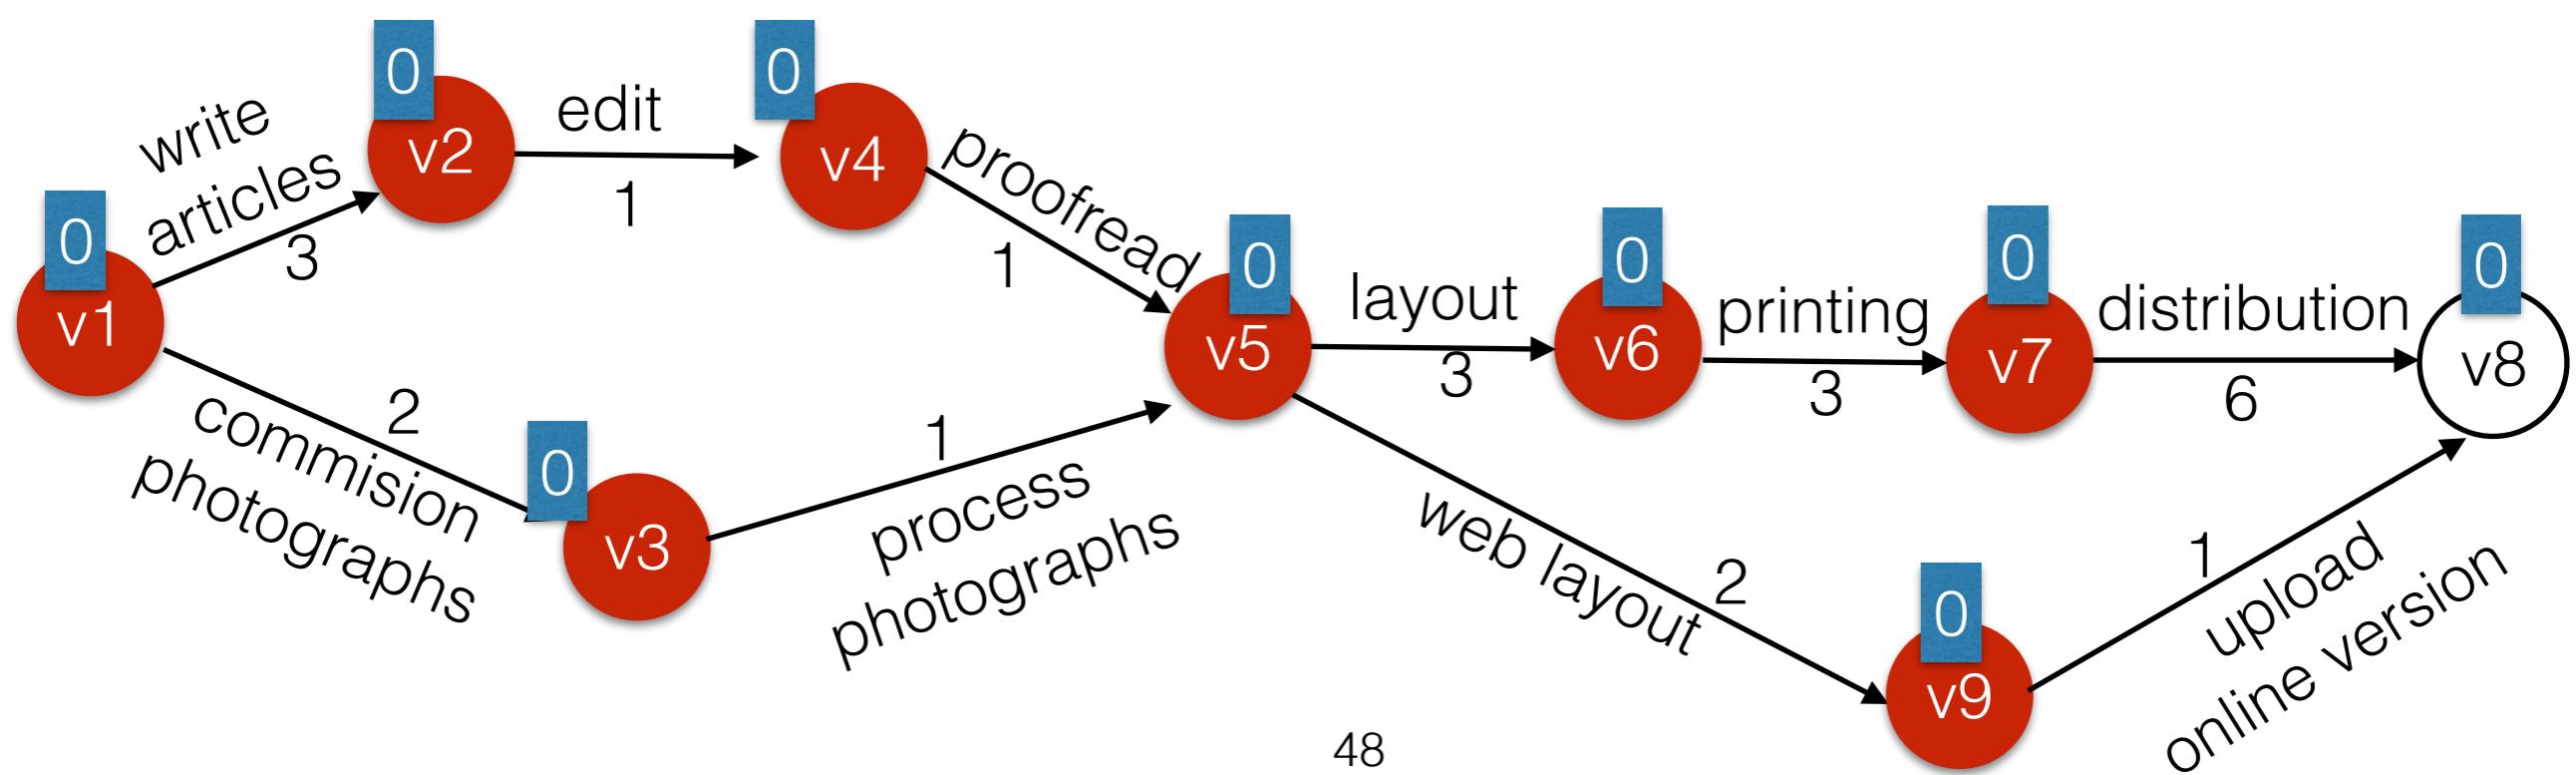


Computing Topological Sort

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes.
- If the indegree of any new vertex becomes 0, enqueue it.

Queue:

Output: v1 v2 v3 v4 v5 v6 v9 v7 v8



Topological Sort - Algorithm

- First annotate each vertex with its indegree.
- Enqueue all vertices with indegree 0.
- While the queue is not empty
 - dequeue a vertex.
 - add it to the topological sort.
 - decrement the indegree of its adjacent vertices.
 - if the indegree of any new vertex becomes 0, enqueue it.

Topological Sort - Algorithm

- First annotate each vertex with its indegree.
- Enqueue all vertices with indegree 0.
- While the queue is not empty
 - dequeue a vertex.
 - add it to the topological sort.
 - decrement the indegree of its adjacent vertices.
 - if the indegree of any new vertex becomes 0, enqueue it.

$$O(|V|+|E|)$$

Topological Sort - Algorithm

- First annotate each vertex with its indegree. $O(|V|+|E|)$
- Enqueue all vertices with indegree 0.
- While the queue is not empty $O(|E|)$ edges visited
 - dequeue a vertex.
 - add it to the topological sort.
 - decrement the indegree of its adjacent vertices.
 - if the indegree of any new vertex becomes 0, enqueue it.

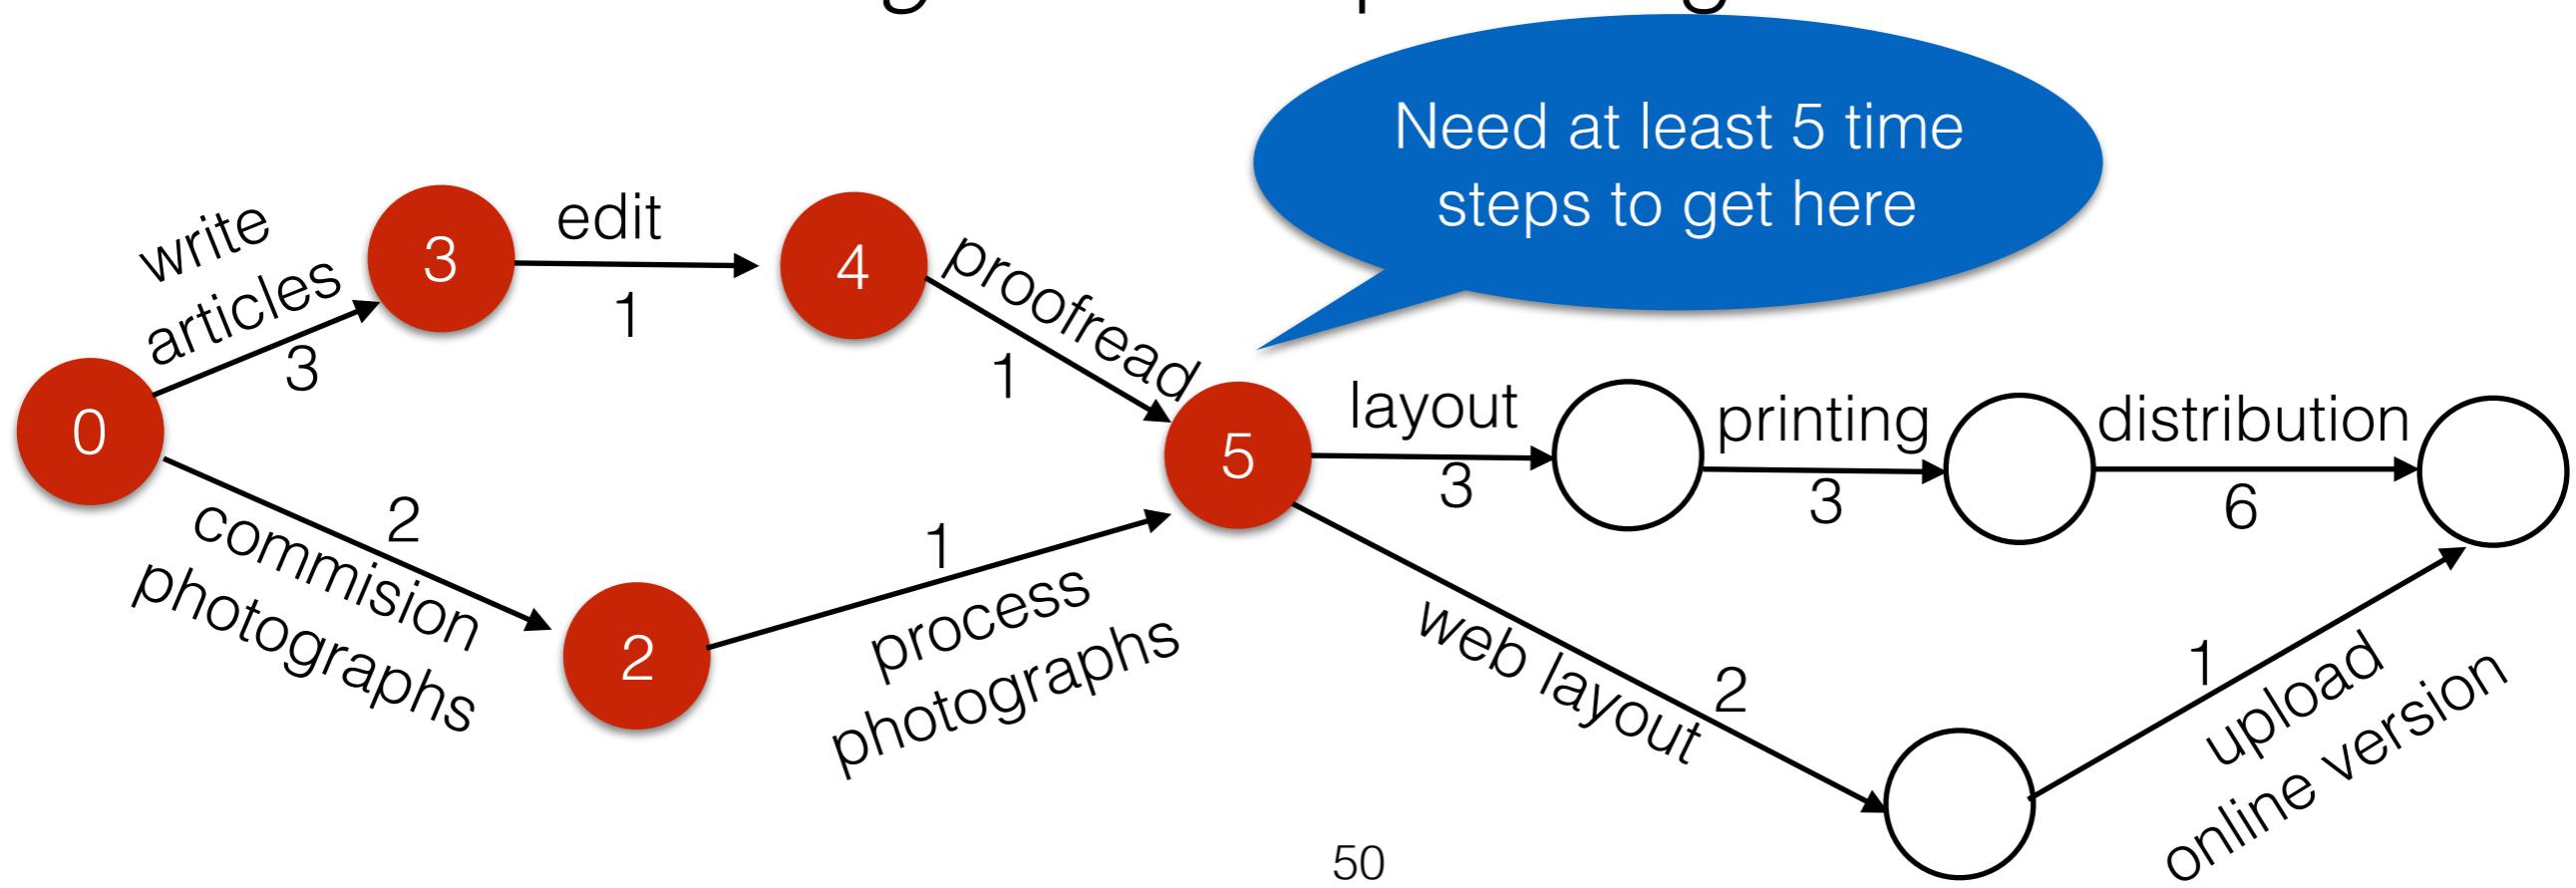
Topological Sort - Algorithm

- First annotate each vertex with its indegree. $O(|V|+|E|)$
- Enqueue all vertices with indegree 0.
- While the queue is not empty $O(|E|)$ edges visited
 - dequeue a vertex.
 - add it to the topological sort.
 - decrement the indegree of its adjacent vertices.
 - if the indegree of any new vertex becomes 0, enqueue it.

Total: $O(|V|+|E|)$

Earliest Completion Time

- Goal: Compute the earliest completion time for each vertex.
- Process vertices in topological order.
- What is the *highest cost* path to get to a vertex?

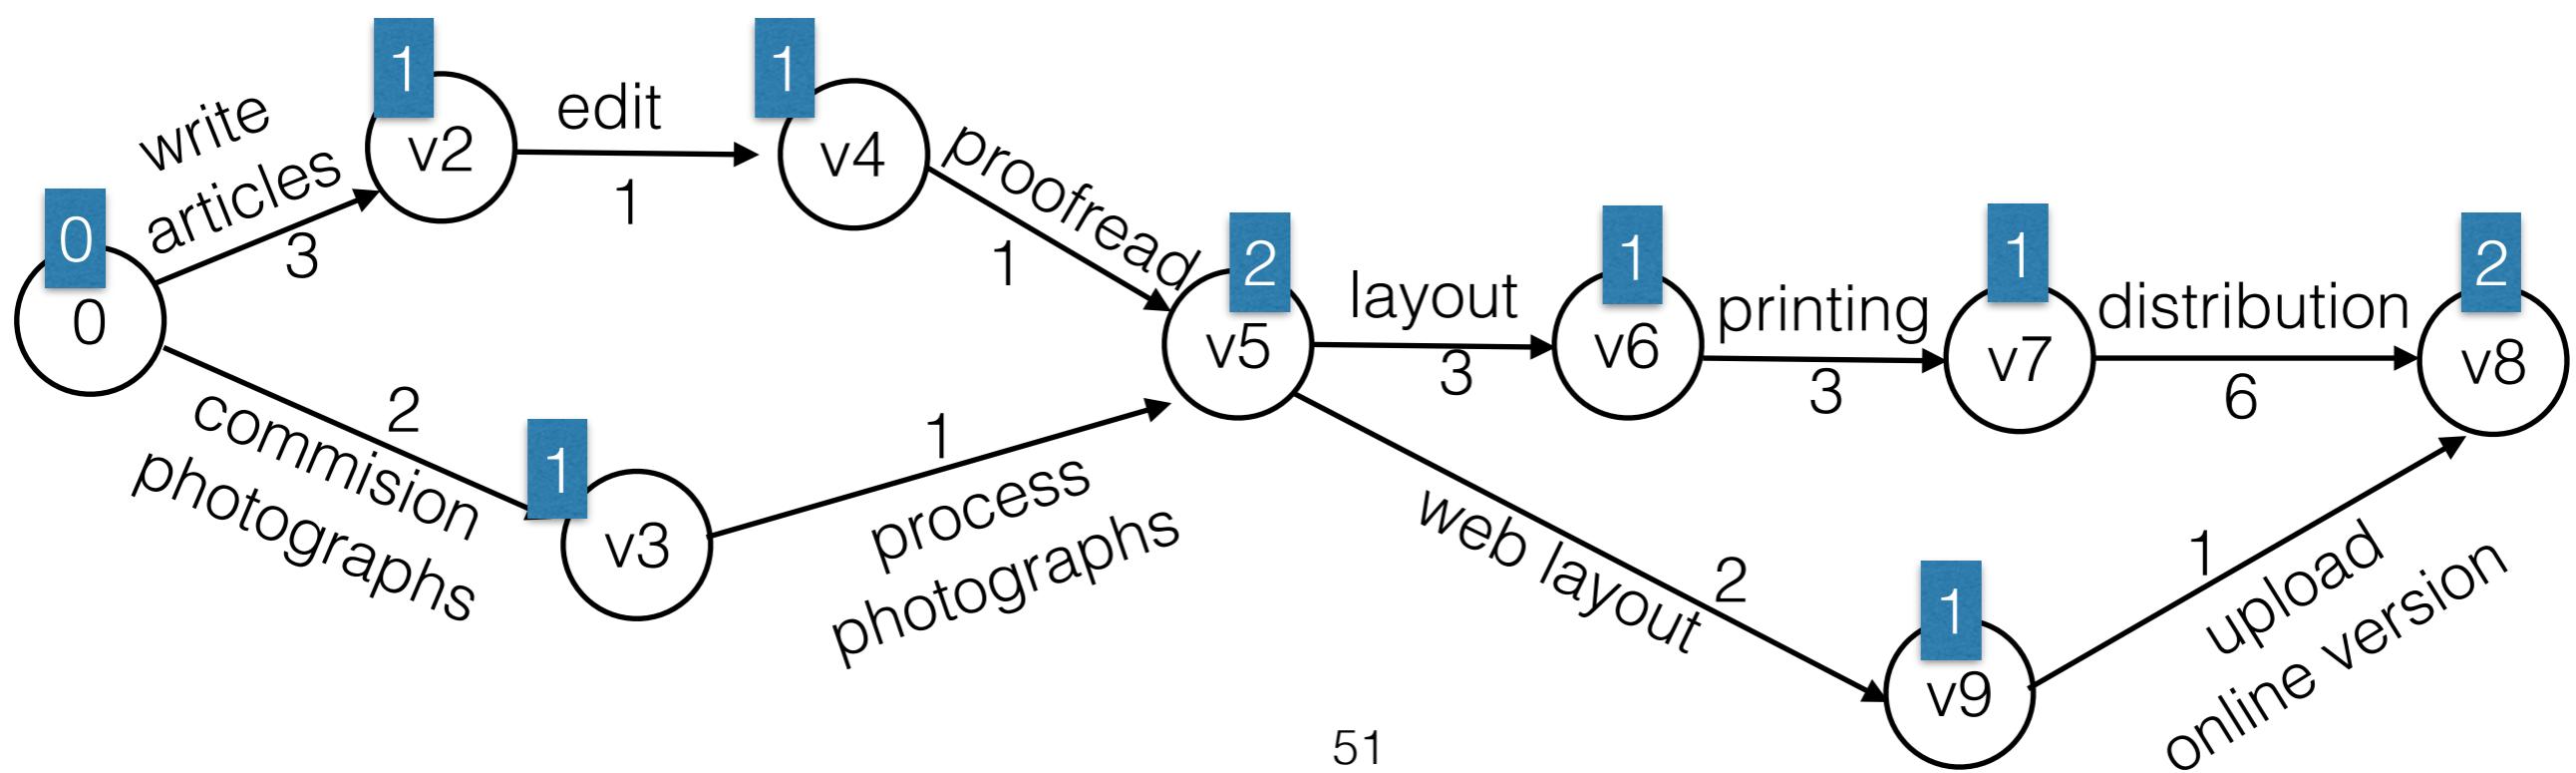


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v1

Output:

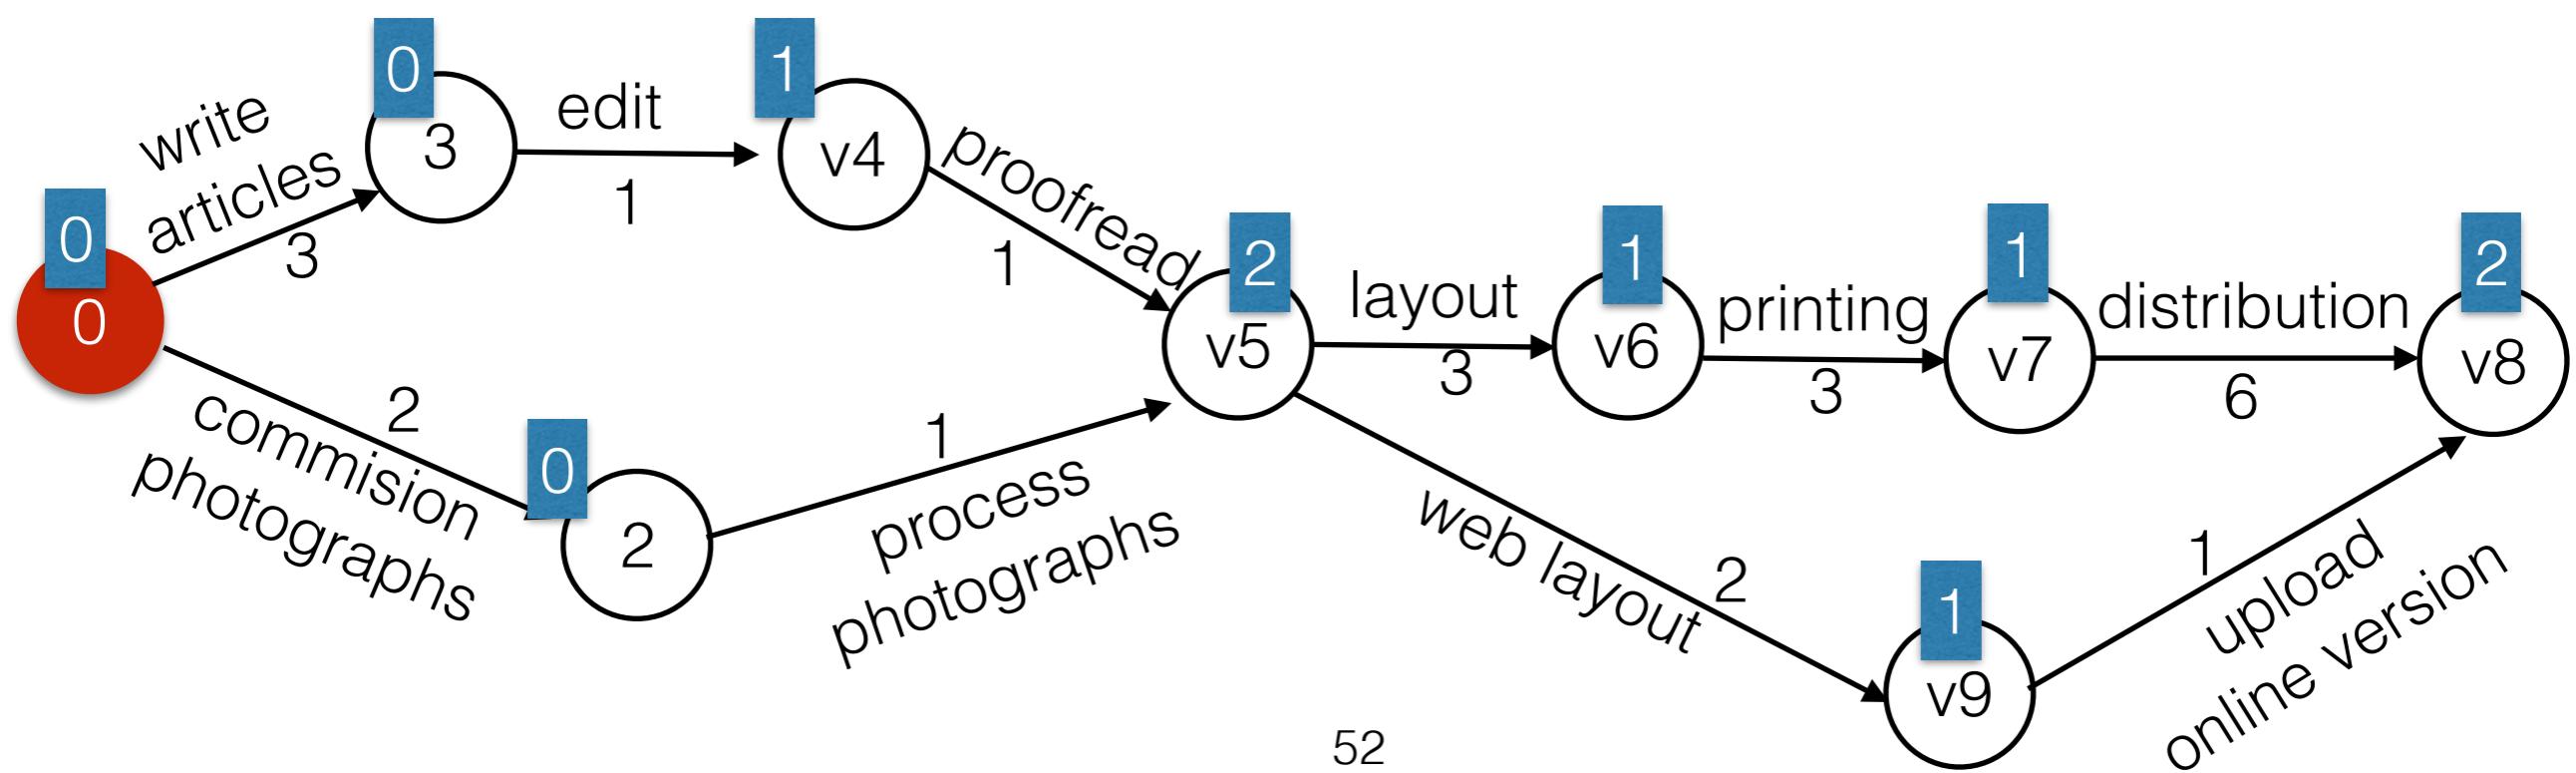


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v2 v3

Output: v1

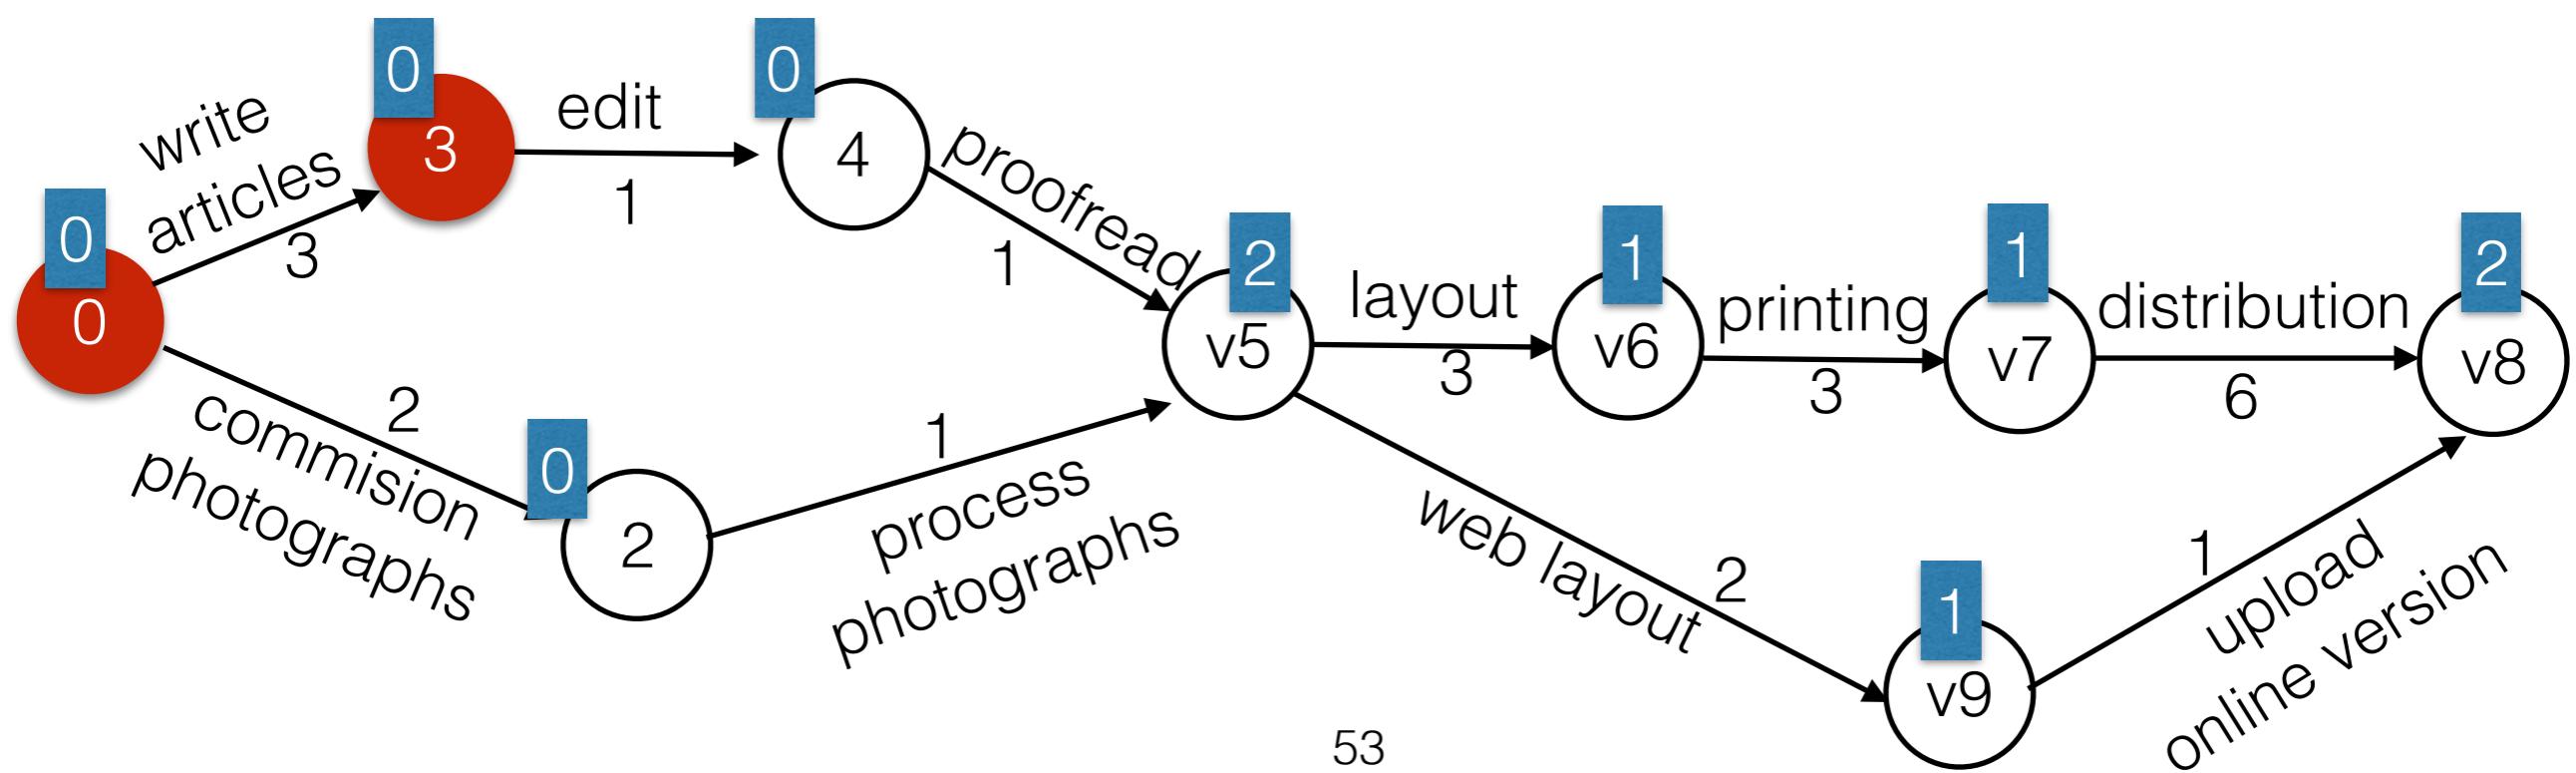


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v3 v4

Output: v1 v2

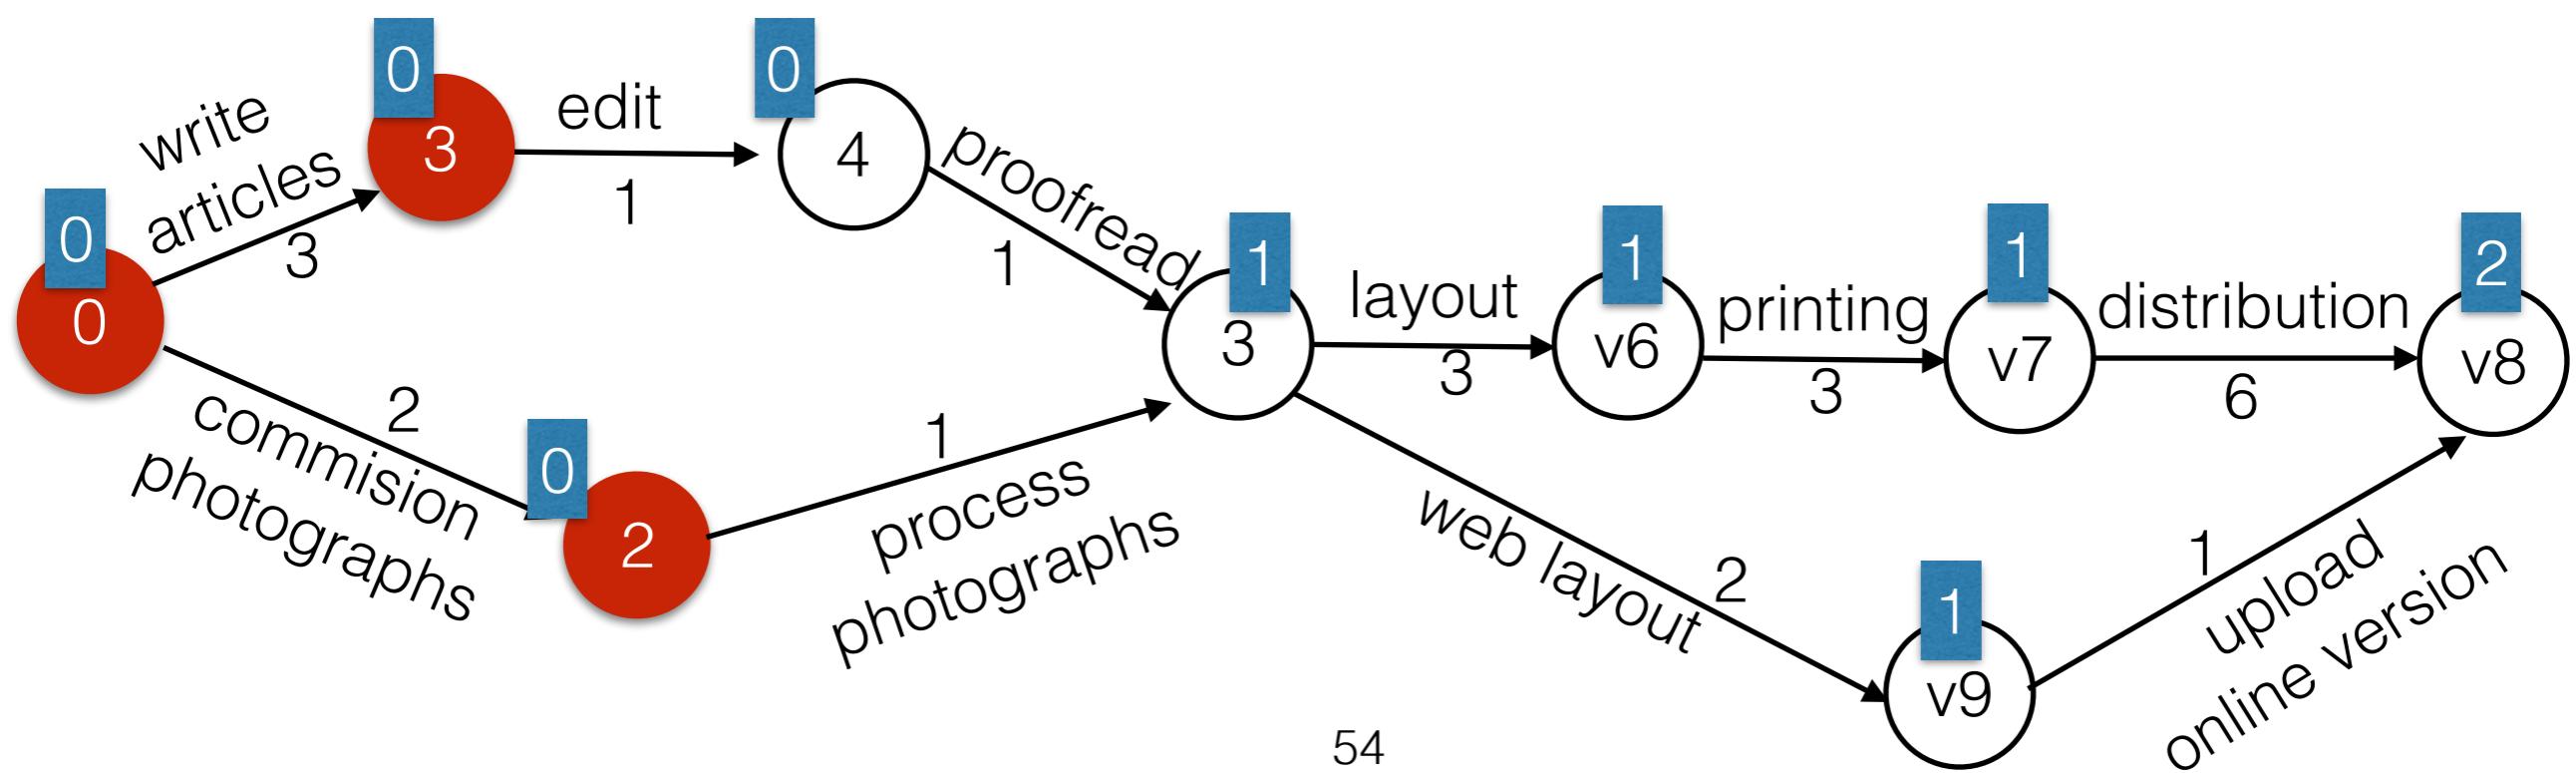


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v4

Output: v1 v2 v3

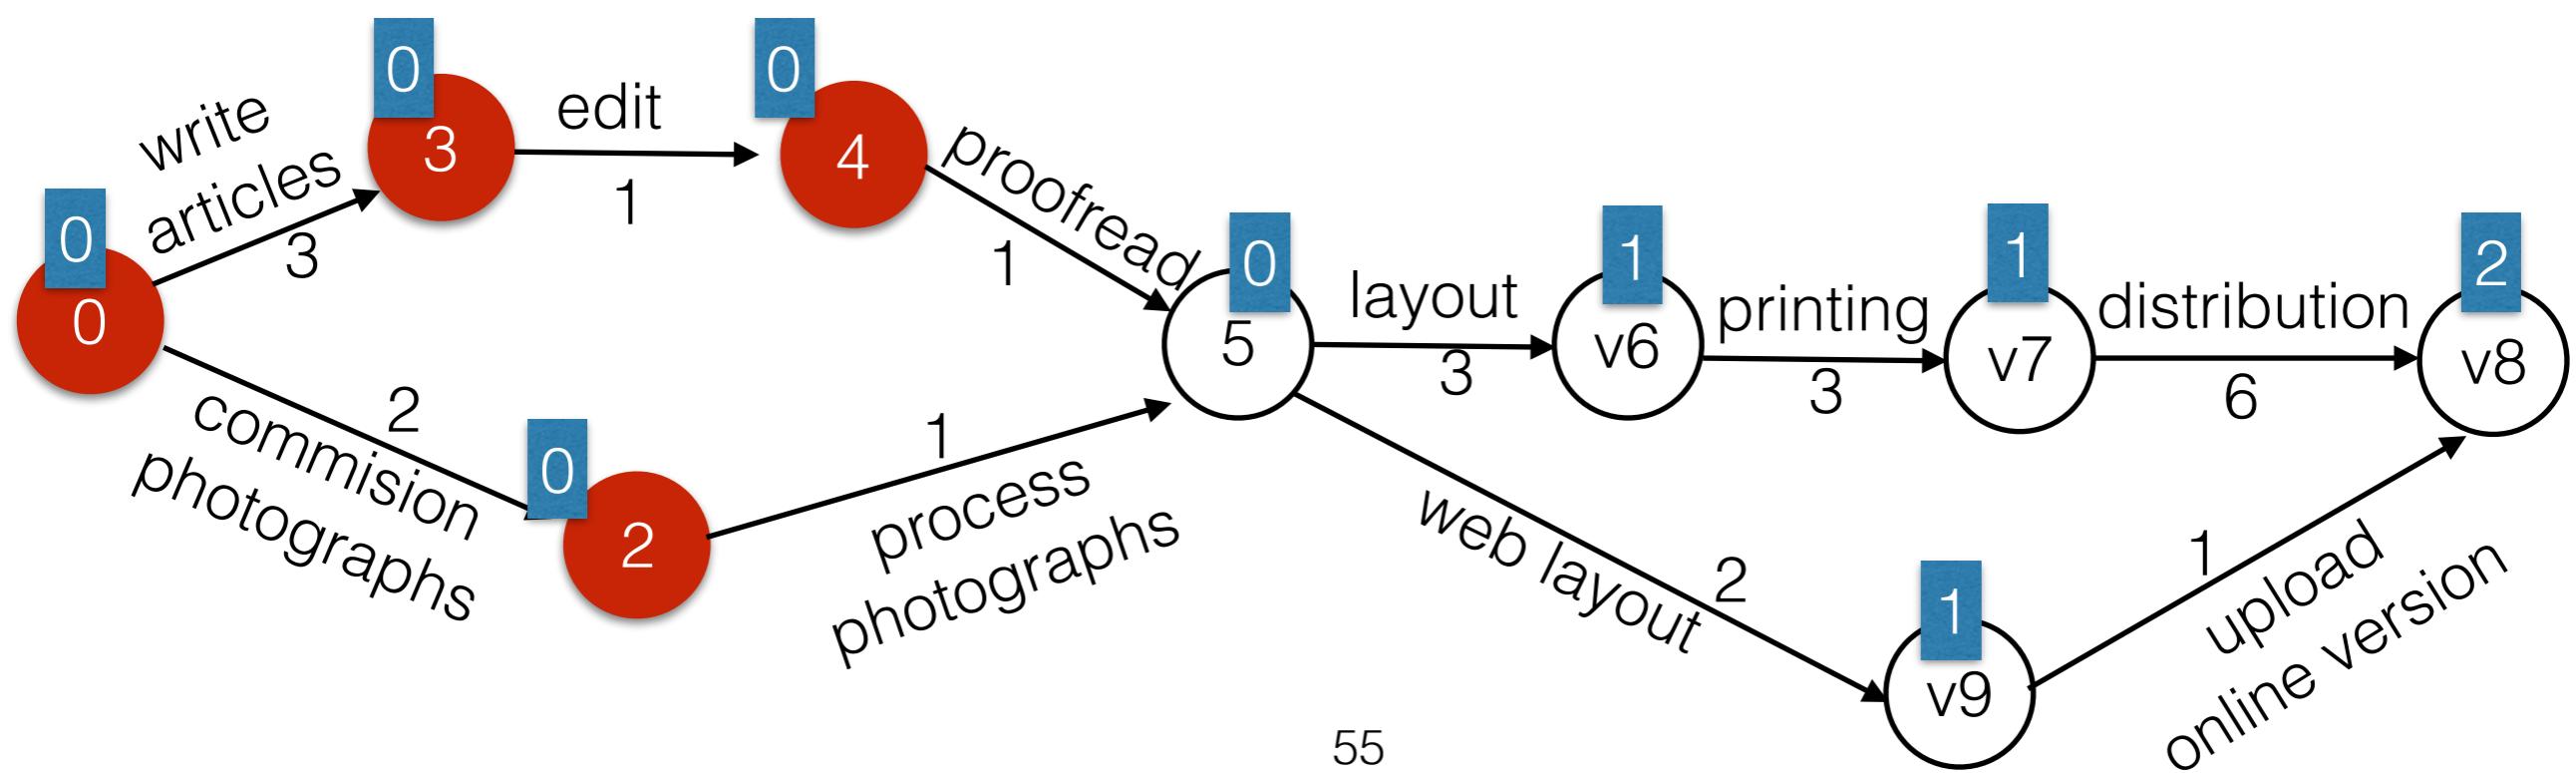


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v5

Output: v1 v2 v3 v4

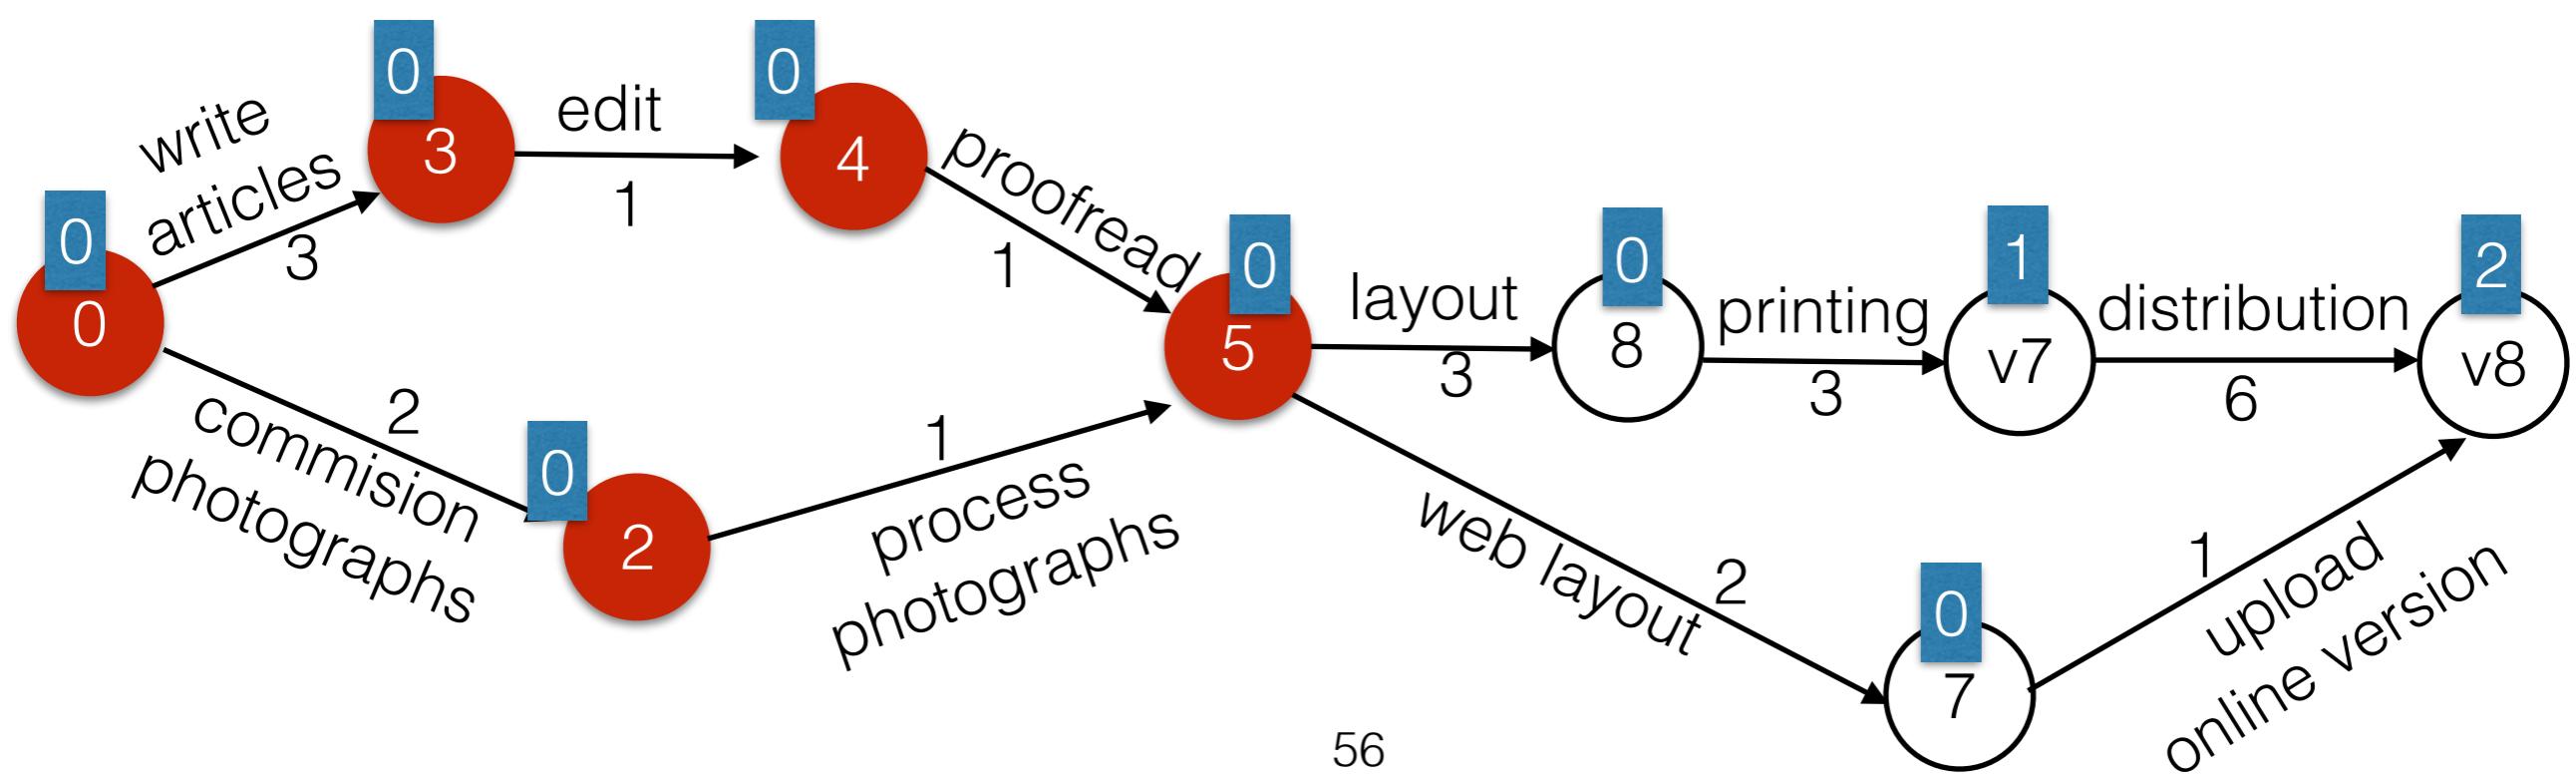


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v6 v9

Output: v1 v2 v3 v4 v5

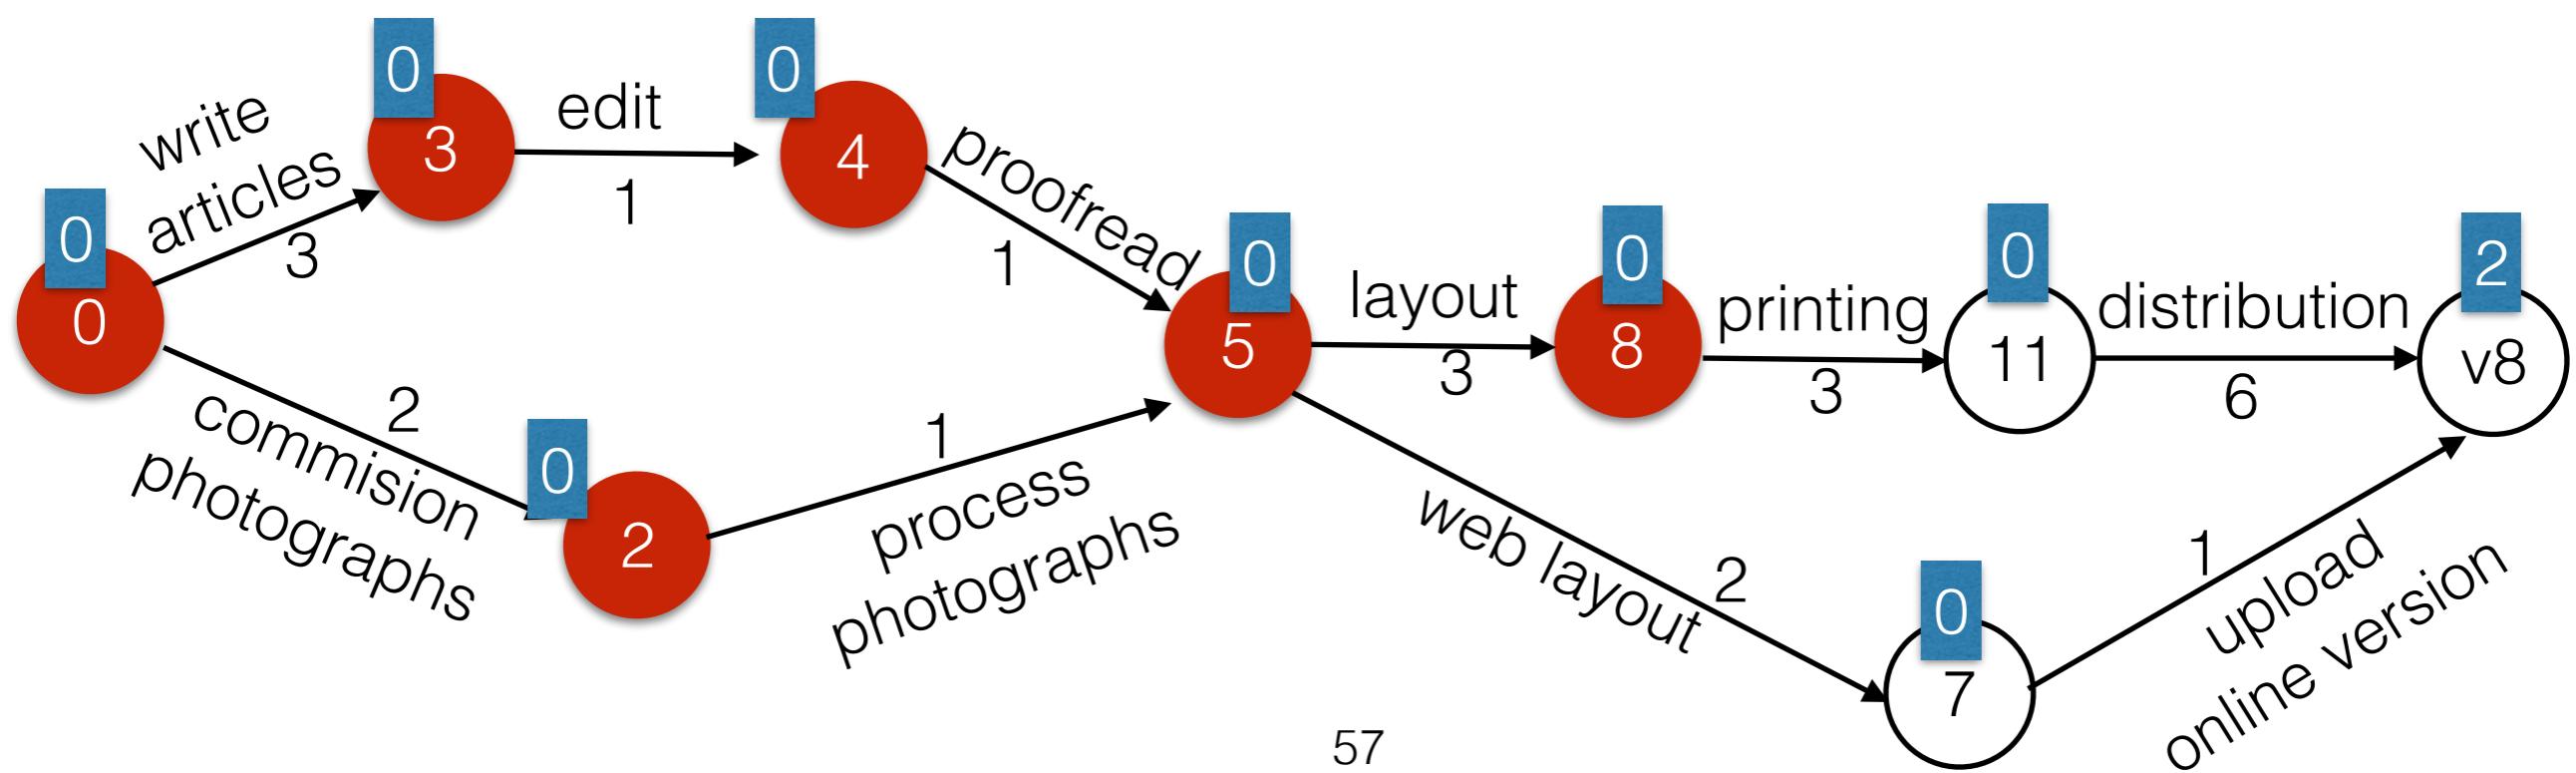


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v9 v7

Output: v1 v2 v3 v4 v5 v6

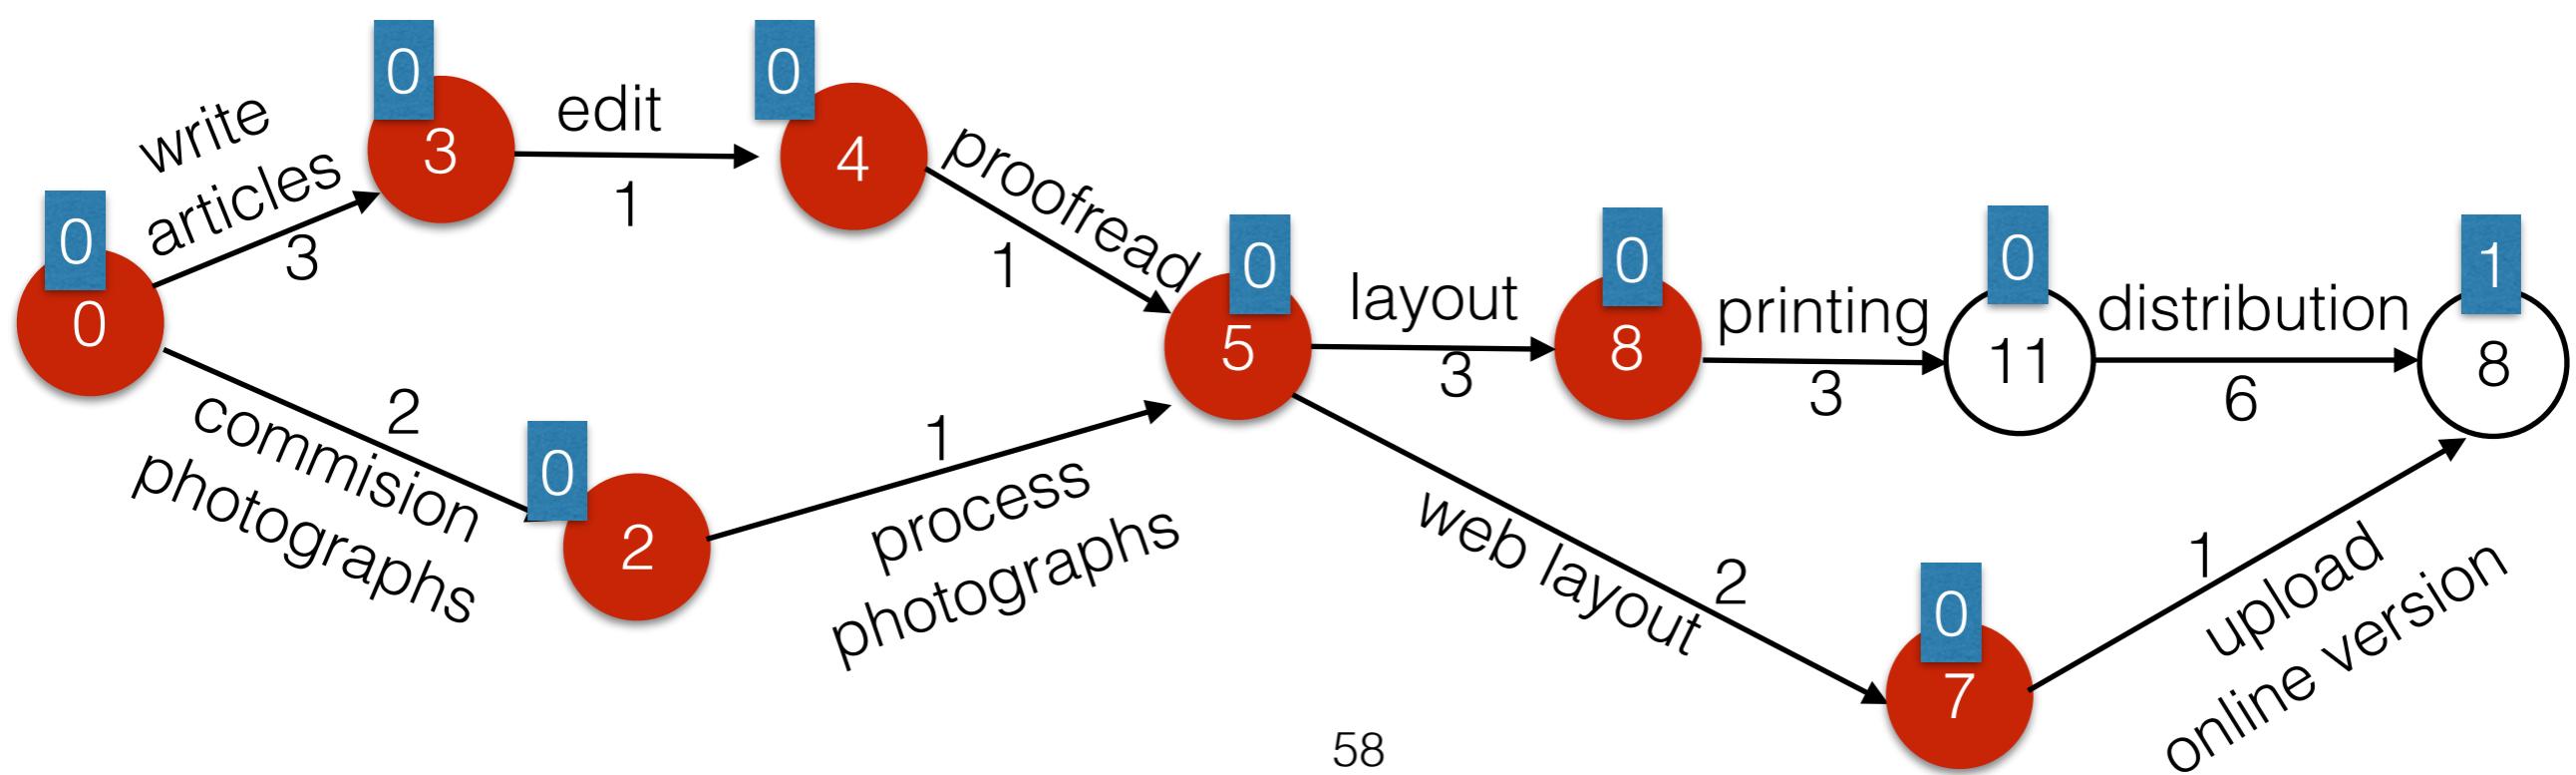


Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v7

Output: v1 v2 v3 v4 v5 v6 v9



Earliest Completion Time

- While the queue is not empty, dequeue a vertex, print it and decrement the indegree of its adjacent nodes. **Update earliest completion time for each adjacent node.**
- If the indegree of any new vertex becomes 0, enqueue it.

Queue: v8

Output: v1 v2 v3 v4 v5 v6 v9 v7

